

**PARALLEL COMPUTING
FOR
SCIENCE AND ENGINEERING**

VICTOR EIJKHOUT

0TH EDITION 2013

Public draft - open for comments

This book will be open source under CC-BY license.

This book covers OpenMP (ok, it will at some point) and MPI. For both systems it covers some of the latest features. There is a healthy emphasis on practical examples.

Contents

1	MPI	6
1.1	<i>Distributed memory and message passing</i>	6
1.2	<i>Basic concepts</i>	8
1.3	<i>Point-to-point communication</i>	10
1.4	<i>Collectives</i>	23
1.5	<i>Data types</i>	28
1.6	<i>Communicators</i>	29
1.7	<i>Hybrid programming: MPI and threads</i>	32
1.8	<i>Leftover topics</i>	33
1.9	<i>Programming projects</i>	38
1.10	<i>Reference to the routines</i>	50
1.11	<i>Review questions</i>	63
1.12	<i>Literature</i>	63
2	Ruminations on parallelism	65
2.1	<i>Fully parallel work</i>	65
2.2	<i>Introduction: Quicksort</i>	65
2.3	<i>Data partitioning</i>	67
3	Hybrid computing	69
4	Support libraries	70
	Appendices	71
A	Practical tutorials	71
A.1	<i>Managing projects with Make</i>	72
A.2	<i>Debugging</i>	81
A.3	<i>Tracing</i>	88
B	Index and list of acronyms	90

Chapter 1

MPI

In this chapter you will learn the use of the main tool for distributed memory programming: the Message Passing Interface (MPI) library. The MPI library has about 250 routines, many of which you may never need. Since this is a textbook, not a reference manual, we will focus on the important concepts and give the important routines for each concept. What you learn here should be enough for most common purposes. You are advised to keep a reference document handy, in case there is a specialized routine, or to look up subtleties about the routines you use.

1.1 Distributed memory and message passing

In its simplest form, a distributed memory machine is a bunch of single computers hooked up with network cables. In fact, this has a name: a *Beowulf cluster*. As you recognize from that setup, each processor will run an independent program, and has its own memory without direct access to other processors' memory. MPI is the magic that makes multiple instantiations of the same executable run so that they know about each other and can exchange data through the network.

One of the reasons that MPI is so successful as a tool for high performance on clusters is that it is very explicit: the programmer controls many details of the data motion between the processors. Consequently, a capable programmer can write very efficient code with MPI. Unfortunately, that programmer will have to spell things out in considerable detail. For this reason, people sometimes call MPI ‘the assembly language of parallel programming’. If that sounds scary, be assured that things are not that bad. You can get started fairly quickly with MPI, using just the basics, and coming to the more sophisticated tools only when necessary.

Another reason that MPI was a big hit with programmers is that it does not ask you to learn a new language: it is a library that can be interface to C/C++ or Fortran; there are even bindings to Python. A related point is that it is easy to install: there are free implementations that you can download and install on any computer that has a Unix-like operating system, even if that is not a parallel machine.

1.1.1 History

Mid 1990s, many parties involved, big concensus. Many competing packages before, few after.

1.1.2 Basic model

Here we sketch the two most common scenarios for using MPI. In the first, the user is working on an interactive machine, which has network access to a number of hosts, typically a network of workstations; see figure 1.1. The user types the command `mpirun` and supplies

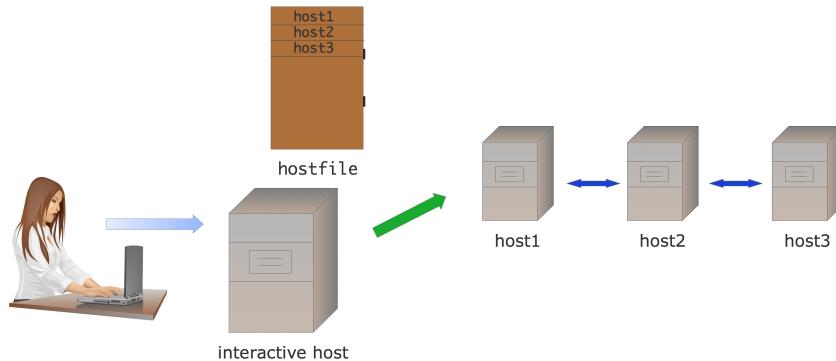


Figure 1.1: Interactive MPI setup

- The number of hosts involved,
- their names, possibly in a hostfile,
- and other parameters, such as whether to include the interactive host; followed by
- the name of the program and its parameters.

The `mpirun` program then makes an `ssh` connection to each of the hosts, giving them sufficient information that they can find each other. All the output of the processors is piped through the `mpirun` program, and appears on the interactive console.

In the second scenario (figure 1.2) the user prepares a *batch job* script with commands, and these will be run when the *batch scheduler* gives a number of hosts to the job. Now the batch script contains the `mpirun`

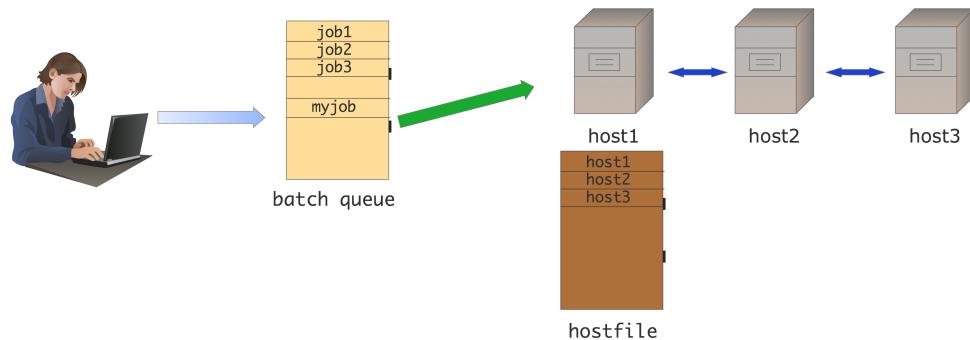


Figure 1.2: Batch MPI setup

command, or some variant such as `ibrun`, and the hostfile is dynamically generated when the job starts. Since the job now runs at a time when the user may not be logged in, any screen output goes into an output file.

1. MPI

You see that in both scenarios the parallel program is started by the `mpirun` command, which only supports an Single Program Multiple Data (SPMD) mode of execution: all hosts execute the same program.

To first order, the network is symmetric. However, the truth is more complicated ('topology-aware communication').

1.1.3 Making and running an MPI program

MPI is a library, called from programs in ordinary programming languages such as C/C++ or Fortran. To compile such a program you use your regular compiler:

```
gcc -c my_mpi_prog.c -I/path/to/mpi.h  
gcc -o my_mpi_prog my_mpi_prog.o -L/path/to/mpi -lmpich
```

However, MPI libraries may have different names between different architectures, making it hard to have a portable makefile. Therefore, MPI typically has shell scripts around your compiler call:

```
mpicc -c my_mpi_prog.c  
mpicc -o my_mpi_prog my_mpi_prog.o
```

MPI programs can be run on many different architectures. Obviously it is your ambition (or at least your dream) to run your code on a cluster with a hundred thousand processors and a fast network. But maybe you only have a small cluster with plain *ethernet*. Or maybe you're sitting in a plane, with just your laptop. An MPI program can be run in all these circumstances – within the limits of your available memory of course.

The way this works is that you do not start your executable directly, but you use a program, typically called `mpirun` or something similar, which makes a connection to all available processors and starts a run of your executable there. So if you have a thousand nodes in your cluster, `mpirun` can start your program once on each, and if you only have your laptop it can start a few instances there. In the latter case you will of course not get great performance, but at least you can test your code for correctness.

1.2 Basic concepts

1.2.1 Initialization / finalization

The reference for the commands introduced here can be found in section [1.10.1.1](#).

Every program that uses MPI needs to initialize and finalize exactly once. In C, the calls are

```
ierr = MPI_Init (&argc, &argv);  
// your code  
ierr = MPI_Finalize();
```

where `argc` and `argv` are the arguments of the main program. The corresponding Fortran calls are

```
call MPI_Init(ierr)
// your code
call MPI_Finalize()
```

(There is a call `MPI_Abort` if you want to abort execution completely.)

We make a few observations.

- MPI routines return an error code. In C, this is a function result; in Fortran it is the final parameter in the calling sequence.
- For most routines, this parameter is the only difference between the C and Fortran calling sequence, but some routines differ in some respect related to the languages. In this case, C has a mechanism for dealing with commandline arguments that Fortran lacks.
- This error parameter is zero if the routine completes successfully, and nonzero otherwise. You can write code to deal with the case of failure, but by default your program will simply abort on any MPI error. See section ?? for more details.

The commandline arguments `argc` and `argv` are only guaranteed to be passed to process zero, so the best way to pass commandline information is by a broadcast (section 1.4.1).

1.2.2 Communicators

Before we can discuss any further MPI routines we need to take a first look at an important concept: that of the *communicator*. A communicator stands for a group of processes. Thus, almost all MPI routines have a communicator argument: you can never ask how many processes there are, or send data to process 5, you have to ask how many processes in a given communicator, or send data to process 5 in a communicator.

There are several reasons for having communicators. One is that sometimes you may want to divide your processes in two groups, each of which engages in a different activity. Another reason is for ease of writing software libraries. If a software library that uses MPI starts by creating its own communicator, even if it contains the same processes as the communicator in the calling program, it can safely use that and never run the risk of confusion between messages in the library code and messages in the user code. We will discuss this in more detail later on.

For most of your MPI programming, the only communicator you will use is `MPI_COMM_WORLD` which contains all your processes. In section 1.6 you will learn of the mechanisms for creating other communicators from it.

1.2.3 Distinguishing between processes

The reference for the commands introduced here can be found in section 1.10.1.2.

In the SPMD model you run the same executable on each of a set of processors; see section HPSC-2.2.2. So how can you do anything useful if all processors run the same code? Here is where your first two MPI routines come in, which query the `MPI_COMM_WORLD` communicator and each process' place in it.

With `MPI_Comm_size` a processor can query how many processes there are in total, and with `MPI_Comm_rank` it can find out what its number is. This rank is a number from zero to the comm size minus one. (Zero-based indexing is used even if you program in *Fortran*.)

Using these calls, the simplest MPI programs does this:

```
// helloworld.c
MPI_Init(&argc,&argv);
MPI_Comm_size(MPI_COMM_WORLD,&ntids);
MPI_Comm_rank(MPI_COMM_WORLD,&mytid);
printf("Hello, this is processor %d out of %d\n",mytid,ntids);
MPI_Finalize();
```

1.3 Point-to-point communication

MPI has two types of message passing routines: point-to-point and collectives. In this section we will discuss point-to-point communication, which involves the interaction of a unique sender and a unique receiver. Collectives, which involve all processes in some joint fashion, will be discussed in the next section.

There is a lot to be said about simple sending and receiving of data. We will go into three broad categories of operations: blocking and non-blocking two-sided communication, and the somewhat more tricky one-sided communication.

Two-sided communication is a little like email: one party send data, which needs to be specified, to another party. The other party can then be expecting a message from a specified sender or it can be open to receiving from any source, but in either case the receiver indicates that something is to be expected. One-sided communication is very different in nature. Compare it to leaving your front-door open and people can bring things to your house, or take them, without you noticing.

1.3.1 Blocking communication

The reference for the commands introduced here can be found in section 1.10.2.

In two-sided communication, one process issues a send call and the other a receive call. Life would be easy if the send call put the data somewhere in the network for the receiving process to find whenever it gets around to its receive call. This ideal scenario is pictured figure 1.3. Of course, if the receiving process gets

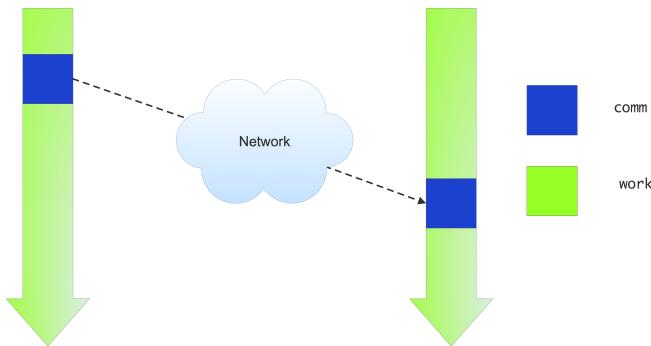


Figure 1.3: Illustration of an ideal send-receive interaction

to the receive call before the send call has been issued, it will be idle until that happens.

Even if processes are optimally synchronized communication introduces some overhead: there is an initial latency connected with every message, and the network also has a limited bandwidth which leads to a transfer time per byte; see HPSC-[2.6.7](#).

The above ideal scenario is not realistic: it assumes that somewhere in the network there is buffer capacity for all messages that are in transit. Since this message volume can be large, we have to worry explicitly about management of send and receive *buffers*.

The easiest scenario is that the sending process keeps the message data in its address space until the receiving process has indicated that it is ready to receive it. This is pictured in figure 1.4. This is known as

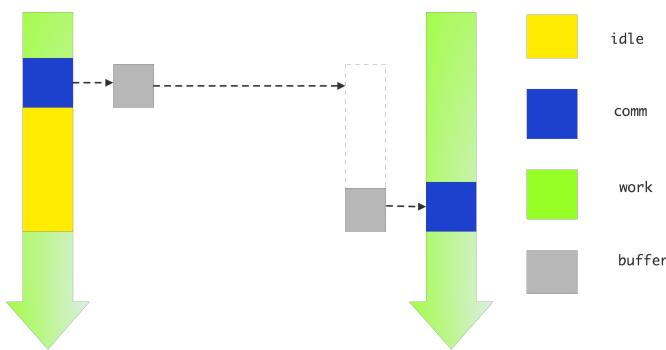


Figure 1.4: Illustration of a blocking communication: the sending processor is idle until the receiving processor issues and finishes the receive call

blocking communication: a process that issues a send or receive call will then block until the corresponding receive or send call is successfully concluded.

It is clear what the first problem with this scenario is: if your processes are not perfectly synchronized your performance may degrade because processes spend time waiting for each other; see HPSC-[2.10.1](#).

But there is a more insidious, more serious problem. Suppose two process need to exchange data, and consider the following pseudo-code, which purports to exchange data between processes 0 and 1:

```
other = 1-mytid; /* if I am 0, other is 1; and vice versa */
send(target=other);
receive(source=other);
```

Imagine that the two processes execute this code. They both issue the send call... and then can't go on, because they are both waiting for the other to issue a receive call. This is known as *deadlock*.

Formally you can describe deadlock as follows. Draw up a graph where every process is a node, and draw a directed arc from process A to B if A is waiting for B. There is deadlock if this directed graph has a loop.

The solution to the deadlock in the above example is to first do the send from 0 to 1, and then from 1 to 0 (or the other way around). So the code would look like:

```

if ( /* I am processor 0 */ ) {
    send(target=other);
    receive(source=other);
} else {
    receive(source=other);
    send(target=other);
}

```

There is even a third, even more subtle problem with blocking communication. Consider the scenario where every processor needs to pass data to its successor, that is, the processor with the next higher rank. The basic idea would be to first send to your successor, then receive from your predecessor. Since the last processor does not have a successor it skips the send, and likewise the first processor skips the receive. The pseudo-code looks like:

```

successor = mytid+1; predecessor = mytid-1;
if ( /* I am not the last processor */ )
    send(target=successor);
if ( /* I am not the first processor */ )
    receive(source=predecessor)

```

This code does not deadlock. All processors but the last one block on the send call, but the last processor executes the receive call. Thus, the processor before the last one can do its send, and subsequently continue to its receive, which enables another send, et cetera.

In one way this code does what you intended to do: it will terminate (instead of hanging forever on a deadlock) and exchange data the right way. However, the execution now suffers from *unexpected serialization*: only one processor is active at any time, so what should have been a parallel operation becomes a sequential

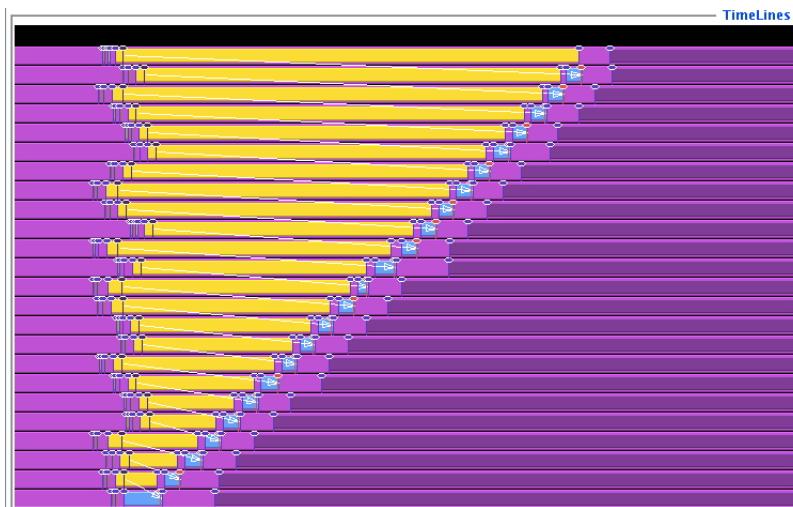


Figure 1.5: Trace of a simple send-recv code

one. This is illustrated in figure 1.5.

Exercise 1.1. Modify your earlier code, run it and reproduce the trace output of figure 1.5.

It is possible to orchestrate your processes to get an efficient and deadlock-free execution, but doing so is a bit cumbersome. There are better solutions which we will explore next.

Exercise 1.2. Give pseudo-code for a solution that uses blocking operations, and is parallel, although maybe not optimally so.

Above, you saw a code fragment with a conditional send:

```
MPI_Comm_rank( .... &mytid );
successor = mytid+1
if ( /* I am not the last processor */ )
    send(target=successor);
```

MPI allows for the following variant which makes the code slightly more homogeneous:

```
MPI_Comm_rank( .... &mytid );
if ( /* I am not the last processor */ )
    successor = mytid+1
else
    successor = MPI_PROC_NULL;
send(target=successor);
```

where the send call is executed by all processors; the target value of `MPI_PROC_NULL` means that no actual send is done. The null processor value is also of use with the `MPI_Sendrecv` call.

1.3.1.1 Deadlock-free blocking communication

The reference for the commands introduced here can be found in section 1.10.3.

Above you saw that with blocking sends the precise ordering of the send and receive calls is crucial. Use the wrong ordering and you get either deadlock, or something that is not efficient at all in parallel. MPI has a way out of this problem that is sufficient for many purposes: the combined send/recv call

```
MPI_Sendrecv( /* send data */ ....
              /* recv data */ .... );
```

This call makes it easy to exchange data between two processors: both specify the other as both target and source. However, there need not be any such relation between target and source: it is possible to receive from a predecessor in some ordering, and send to a successor in that ordering.

Above you saw some examples that had most processors doing both a send and a receive, but some only a send or only a receive. You can still use `MPI_Sendrecv` in this call if you use `MPI_PROC_NULL` for the unused source or target argument.

Exercise 1.3. Take your code from exercise 1.1 and rewrite it to use the `MPI_Sendrecv` call.

Run it and produce a trace output. Do you see the serialization behaviour of your earlier code?

1.3.1.2 Subtleties with processor synchronization

Blocking communication involves a complicated dialog between the two processors involved. Processor one says ‘I have this much data to send; do you have space for that?’, to which processor two replies ‘yes, I do; go ahead and send’, upon which processor one does the actual send. This back-and-forth (technically known as a *handshake*) takes a certain amount of communication overhead. For this reason, network hardware will sometimes forgo the handshake for small messages, and just send them regardless, knowing that the other process has a small buffer for such occasions.

One strange side-effect of this strategy is that a code that should *deadlock* according to the MPI specification does not do so. In effect, you may be shielded from your own programming mistake! Of course, if you then run a larger problem, and the small message becomes larger than the threshold, the deadlock will suddenly occur. So you find yourself in the situation that a bug only manifests itself on large problems, which are usually harder to debug. In this case, replacing every `MPI_Send` with a `MPI_Ssend` will force the handshake, even for small messages.

Conversely, you may sometimes wish to avoid the handshake on large messages. MPI as a solution for this: the `MPI_Rsend` (‘ready send’) routine sends its data immediately, but it needs the receiver to be ready for this. How can you guarantee that the receiving process is ready? You could for instance do the following (this uses non-blocking routines, which are explained below in section 1.3.2):

```
if ( receiving ) {
    MPI_Irecv()      // post non-blocking receive
    MPI_Barrier()   // synchronize
} else if ( sending ) {
    MPI_Barrier()   // synchronize
    MPI_Rsend()     // send data fast
```

When the barrier is reached, the receive has been posted, so it is safe to do a ready send. However, global barriers are not a good idea. Instead you would just synchronize the two processes involved.

Exercise 1.4. Give pseudo-code for a scheme where you synchronize the two processes through the exchange of a blocking zero-size message.

1.3.1.3 Speculative messaging

The reference for the commands introduced here can be found in section 1.10.2.1.

In some applications it makes sense that a message can come from one of a number of processes. In this case, it is possible to specify `MPI_ANY_SOURCE` as the source. To find out where the message actually came from, you would use the `MPI_SOURCE` field of the status object that is delivered by `MPI_Recv` or the `MPI_Wait...` call after an `MPI_Irecv`.

There are various scenarios where receiving from ‘any source’ makes sense. One is that of the *master-worker model*. The master task would first send data to the worker tasks, then issues a blocking wait for the data of whichever process finishes first.

1.3.2 Non-blocking communication

The reference for the commands introduced here can be found in section 1.10.4.

In the previous section you saw that blocking communication makes programming tricky if you want to avoid deadlock and performance problems. The main advantage of these routines is that you have full control about where the data is: if the send call returns the data has been successfully received, and the send buffer can be used for other purposes or de-allocated.

By contrast, the non-blocking calls `MPI_Isend` and `MPI_Irecv` do not wait for their counterpart: in effect they tell the runtime system ‘here is some data and please send it as follows’ or ‘here is some buffer space, and expect such-and-such data to come’. This is illustrated in figure 1.6.

While the use of non-blocking routines prevents deadlock, it introduces two new problems:

1. When the send call returns, the send buffer may not be safe to overwrite; when the recv call returns, you do not know for sure that the expected data is in it. Thus, you need a mechanism to make sure that data was actually sent or received.
2. With a blocking send call, you could repeatedly fill the send buffer and send it off. To send multiple messages with non-blocking calls you have to allocate multiple buffers.

For the first problem, MPI has two types of routines. The `MPI_Wait...` calls are blocking: when you issue such a call, your execution will wait until the specified requests have been completed. A typical way of using them is:

```
// start non-blocking communication
MPI_Isend( ... ); MPI_Irecv( ... );
// do work that does not depend on incoming data
...
// wait for the Isend/Irecv calls to finish
MPI_Wait( ... );
// now do the work that absolutely needs the incoming data
...
```

There are several wait calls:

- `MPI_Wait` waits for a single request. If you are indeed waiting for a single nonblocking communication to complete, this is the right routine. If you are waiting for multiple requests you could call this routine in a loop.

```
for (p=0; p<nrequests ; p++)
    MPI_Wait(request[p], &(status[p]));
```

However, this would be inefficient if the first request is fulfilled much later than the others: your waiting process would have lots of idle time. In that case, use one of the following routines.

- `MPI_Waitall` allows you to wait for a number of requests, and it does not matter in what sequence they are satisfied. Using this routine is easier to code than the loop above, and it could be more efficient.
- The ‘waitall’ routine is good if you have need all nonblocking communications to be finished before you can proceed with the rest of the program. However, sometimes it is possible to take action as each request is satisfied. In that case you could use `MPI_Waitany` and write:

```

for (p=0; p<nrequests; p++) {
    MPI_Waitany(nrequests, request_array, &index, &status);
    // operate on buffer[index]
}

```

Note that this routine takes a single status argument, passed by reference, and not an array of statuses!

- `MPI_Waitsome` is very much like `Waitany`, except that it returns multiple numbers, if multiple requests are satisfied. Now the status argument is an array of `MPI_Status` objects.

Figure 1.7 shows the trace of a non-blocking execution using `MPI_Waitall`.

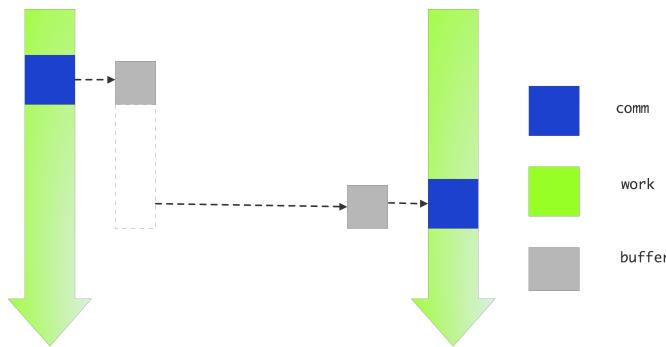


Figure 1.6: Illustration of a non-blocking communication: the sending processor immediately continues execution after issuing the send call

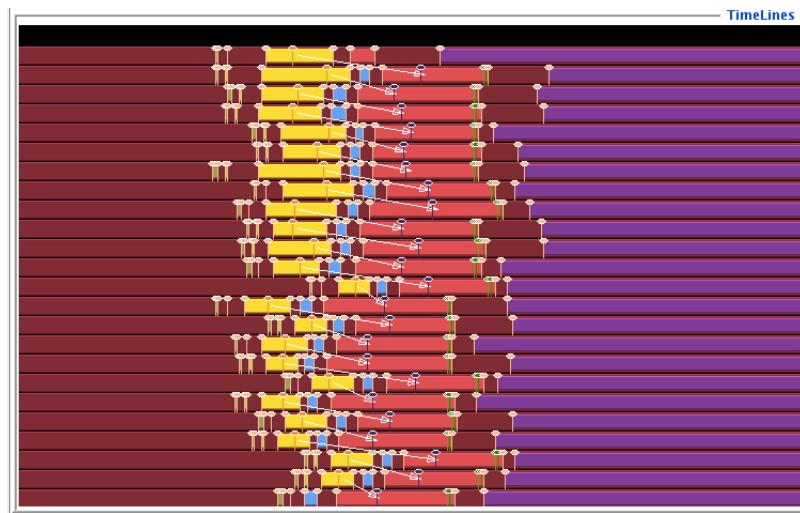


Figure 1.7: A trace of a nonblocking send between neighbouring processors

Exercise 1.5. Read section HPSC-6.4 and give pseudo-code for the distributed sparse matrix-vector product using the above idiom for using `MPI_Test...` calls. Discuss the ad-

vantages and disadvantages of this approach. The answer is not going to be black and white: discuss when you expect which approach to be preferable.

The MPI_Wait... routines are blocking. Thus, they are a good solution if the receiving process can not do anything until the data (or at least *some* data) is actually received. The MPI_Test.... calls are themselves non-blocking: they test for whether one or more requests have been fulfilled, but otherwise immediately return. This can be used in the *master-worker model*: the master process creates tasks, and sends them to whichever worker process has finished its work, but while it waits for the workers it can itself do useful work. Pseudo-code:

```
while ( not done ) {
    // create new inputs for a while
    ...
    // see if anyone has finished
    MPI_Test( .... &index, &flag );
    if ( flag ) {
        // receive processed data and send new
    }
}
```

1.3.2.1 Overlap of computation and communication

Non-blocking routines have long held the promise of letting a program *overlap its computation and communication*. The idea was that after posting the non-blocking calls the program could proceed to do non-communication work, while another part of the system would take care of the communication. Unfortunately, a lot of this communication involved activity in user space, so the solution would have been to let it be handled by a separate thread. Until recently, processors were not efficient at doing such multi-threading, so true overlap stayed a promise for the future.

1.3.2.2 More about non-blocking

Above we used MPI_Irecv, but we could have used the MPI_Recv routine. There is nothing special about a non-blocking or synchronous message once it arrives; the MPI_Recv call can match any of the send routines you have seen so far (but not MPI_Sendrecv).

1.3.3 One-sided communication

The reference for the commands introduced here can be found in section [1.10.5](#).

Above, you saw point-to-point operations of the two-sided type: they require the co-operation of a sender and receiver. This co-operation could be loose: you can post a receive with MPI_ANY_SOURCE as sender, but there had to be both a send and receive call. In this section, you will see one-sided communication routines where a process can do a ‘put’ or ‘get’ operation, writing data to or reading it from another processor, without that other processor’s involvement.

In one-sided MPI operations, also known as Remote Direct Memory Access (RDMA) or Remote Memory Access (RMA) operations, there are still two processes involved: the *origin*, which is the process that

originates the transfer, whether this is a ‘put’ or a ‘get’, and the *target* whose memory is being accessed. Unlike with two-sided operations, the target does not perform an action that is the counterpart of the action on the origin.

That does not mean that the origin can access arbitrary data on the target at arbitrary times. First of all, one-sided communication in MPI is limited to accessing only a specifically declared memory area on the target: the target declares an area of user-space memory that is accessible to other processes. This is known as a *window*. Windows limit how origin processes can access the target’s memory: you can only ‘get’ data from a window or ‘put’ it into a window; all the other memory is not reachable from other processes.

The alternative to having windows is to use *distributed shared memory* or *virtual shared memory*: memory is distributed but acts as if it shared. The so-called Partitioned Global Address Space (PGAS) languages such as Unified Parallel C (UPC) use this model. The MPI RMA model makes it possible to lock a window which makes programming slightly more cumbersome, but the implementation more efficient.

Within one-sided communication, MPI has two modes: active RMA and passive RMA. In *active RMA*, or *active target synchronization*, the target sets boundaries on the time period (the ‘epoch’) during which its window can be accessed. The main advantage of this mode is that the origin program can perform many small transfers, which are aggregated behind the scenes. Active RMA acts much like asynchronous transfer with a concluding `Waitall`.

In *passive RMA*, or *passive target synchronization*, the target process puts no limitation on when its window can be accessed. (PGAS languages such as UPC are based on this model: data is simply read or written at will.) While intuitively it is attractive to be able to write to and read from a target at arbitrary time, there are problems. For instance, it requires a remote agent on the target, which may interfere with execution of the main thread, or conversely it may not be activated at the optimal time. Passive RMA is also very hard to debug and can lead to strange deadlocks.

1.3.3.1 Windows

The reference for the commands introduced here can be found in section 1.10.5.1.

A window is a contiguous area of memory, defined with respect to a communicator: each process specifies

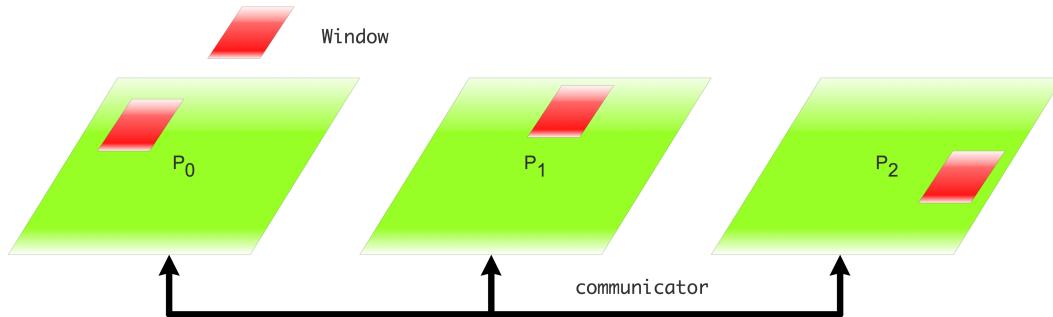


Figure 1.8: Collective definition of a window for one-sided data access

a memory area. Routine for creating and releasing windows are collective, so each process has to call them; see figure 1.8.

```
MPI_Info info;
MPI_Win window;
MPI_Win_create( /* size */, info, comm, &window );
MPI_Win_free( &window );
```

(For the `info` parameter you can often use `MPI_INFO_NULL`.) While the creation of a window is collective, each processor can specify its own window size, including zero, and even the type of the elements in it.

1.3.3.2 Active target synchronization: epochs

There are two mechanisms for *active target synchronization*, that is, one-sided communications where both sides are involved to the extent that they declare the communication epoch. In this section we look at the first mechanism, which is to use a *fence* operation:

```
MPI_Win_fence (int assert, MPI_Win win, ierr)
```

This operation is collective on the communicator of the window. It is comparable to `MPI_Wait` calls for non-blocking communication.

Unlike with wait calls, you always need two fences: one before and one after the so-called *epoch*. You can give various hints to the system about this epoch versus the ones before and after through the `assert` parameter.

```
MPI_Win_fence( (MPI_MODE_NOPUT | MPI_MODE_NOPRECEDE), win);
MPI_Get( /* operands */, win);
MPI_Win_fence( MPI_MODE_NOSUCCEED, win);
```

In between the two fences the window is exposed, and while it is you should not access it locally. If you absolutely need to access it locally, you can use an RMA operation for that. Also, there can be only one remote process that does a put; multiple accumulate accesses are allowed.

Fences are, together with other window calls, collective operations. That means they imply some amount of synchronization between processes. Consider:

```
MPI_Win_fence( ... win ... ); // start an epoch
if (mytid==0) // do lots of work
else // do almost nothing
MPI_Win_fence( ... win ... ); // end the epoch
```

and assume that all processes execute the first fence more or less at the same time. The zero process does work before it can do the second fence call, but all other processes can call it immediately. However, they can not finish that second fence call until all one-sided communication is finished, which means they wait for the zero process.

As a further restriction, you can not mix Get with Put or Accumulate calls in a single epoch. Hence, we can characterize an epoch as an *access epoch* on the origin, and as an *exposure epoch* on the target.

1. MPI

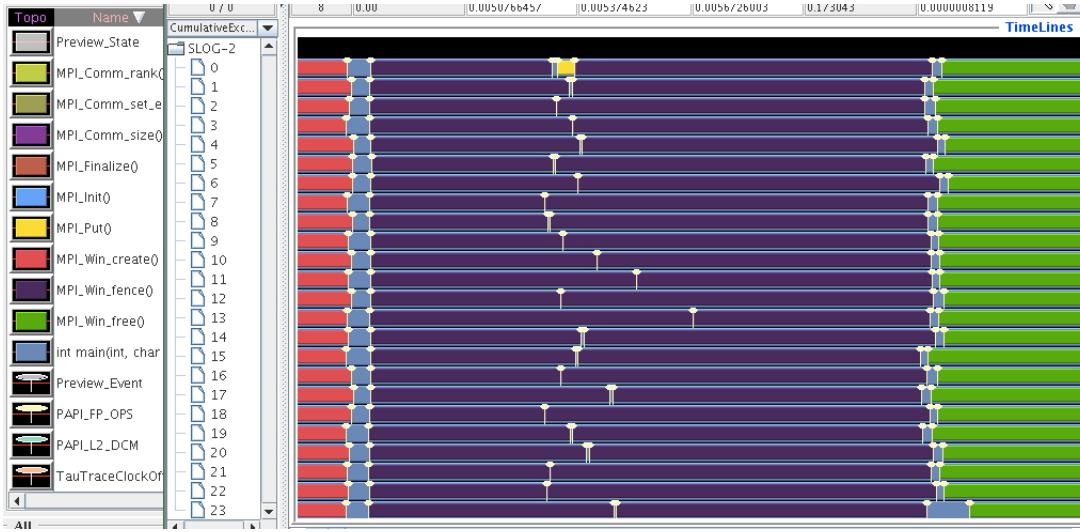


Figure 1.9: A trace of a one-sided communication epoch where process zero only originates a one-sided transfer

Assertions are an integer parameter: you can add or logical-or values. The value zero is always correct. There are two types of parameters. Local assertions are:

- MPI_MODE_NOSTORE The preceding epoch did not store anything in this window.
- MPI_MODE_NOPUT The following epoch will not store anything in this window.

Global assertions:

- MPI_MODE_NOPRECEDE This process made no RMA calls in the preceding epoch.
- MPI_MODE_NOSUCCEED This process will make no RMA calls in the next epoch.

1.3.3.3 Put, get, accumulate

Window areas are accessible to other processes in the communicator by specifying the process rank and an offset from the base of the window.

```
MPI_Put (
    void *origin_addr, int origin_count, MPI_Datatype origin_datatype,
    int target_rank,
    MPI_Aint target_disp, int target_count, MPI_Datatype target_datatype,
    MPI_Win window)
```

The MPI_Get call is very similar; a third one-sided routine is MPI_Accumulate which does a reduction operation on the results that are being put:

```
MPI_Accumulate (
    void *origin_addr, int origin_count, MPI_Datatype origin_datatype,
    int target_rank,
    MPI_Aint target_disp, int target_count, MPI_Datatype target_datatype,
```

```
MPI_Op op, MPI_Win window)
```

Exercise 1.6. Implement an ‘all-gather’ operation using one-sided communication: each processor stores a single number, and you want each processor to build up an array that contains the values from all processors. Note that you do not need a special case for a processor collecting its own value: doing ‘communication’ between a processor and itself is perfectly legal.

Accumulate is a reduction with remote result. As with `MPI_Reduce`, the order in which the operands are accumulated is undefined. The same predefined operators are available, but no user-defined ones. There is one extra operator: `MPI_REPLACE`, this has the effect that only the last result to arrive is retained.

1.3.3.4 Put vs Get

```
while(!converged(A)) {
    update(A);
    MPI_Win_fence(MPI_MODE_NOPRECEDE, win);
    for(i=0; i < toneighbors; i++)
        MPI_Put(&frombuf[i], 1, fromtype[i], toneighbor[i],
                todisp[i], 1, totype[i], win);
    MPI_Win_fence((MPI_MODE_NOSTORE | MPI_MODE_NOSUCCEED), win);
}

while(!converged(A)) {
    update_boundary(A);
    MPI_Win_fence((MPI_MODE_NOPUT | MPI_MODE_NOPRECEDE), win);
    for(i=0; i < fromneighbors; i++)
        MPI_Get(&tobuf[i], 1, totype[i], fromneighbor[i],
                fromdisp[i], 1, fromtype[i], win);
    update_core(A);
    MPI_Win_fence(MPI_MODE_NOSUCCEED, win);
}
```

1.3.3.5 More active target synchronization

The reference for the commands introduced here can be found in section [1.10.5.2](#).

There is a more fine-grained ways of doing *active target synchronization*. While fences corresponded to a global synchronization of one-sided calls, the `MPI_Win_start`, `MPI_Win_complete`, `MPI_Win_post`, `Win_wait` routines are suitable, and possibly more efficient, if only a small number of processor pairs is involved. Which routines you use depends on whether the processor is an *origin* or *target*.

If the current process is going to have the data in its window accessed, you define an *exposure epoch* by:

1. MPI

```
MPI_Win_post( /* group of origin processes */ )
MPI_Win_wait()
```

This turns the current processor into a target for access operations issued by a different process.

If the current process is going to be issuing one-sided operations, you define an *access epoch* by:

```
MPI_Win_start( /* group of target processes */ )
// access operations
MPI_Win_complete()
```

This turns the current process into the origin of a number of one-sided access operations.

Both pairs of operations declare a *group of processors*; see section 1.6.3 for how to get such a group from a communicator. On an origin processor you would specify a group that includes the targets you will interact with, on a target processor you specify a group that includes the possible origins.

1.3.3.6 Passive target synchronization

In *passive target synchronization* only the origin is actively involved: the target makes no calls whatsoever. This means that the origin process remotely locks the window on the target.

```
MPI_Win_lock (int locktype, int rank, int assert, MPI_Win win)
MPI_Win_unlock (int rank, MPI_Win win)
```

During an access epoch, a process can initiate and finish a one-sided transfer.

```
If (rank == 0) {
    MPI_Win_lock (MPI_LOCK_SHARED, 1, 0, win);
    MPI_Put (outbuf, n, MPI_INT, 1, 0, n, MPI_INT, win);
    MPI_Win_unlock (1, win);
}
```

The two lock types are:

- `MPI_LOCK_SHARED` which should be used for `Get` calls: since multiple processors are allowed to read from a window in the same epoch, the lock can be shared.
- `MPI_LOCK_EXCLUSIVE` which should be used for `Put` and `Accumulate` calls: since only one processor is allowed to write to a window during one epoch, the lock should be exclusive.

These routines make MPI behave like a shared memory system; the instructions between locking and unlocking the window effectively become *atomic operations*.

1.3.3.7 Details

Sometimes an architecture has memory that is shared between processes, or that otherwise is fast for one-sided communication. To put a window in such memory, it can be placed in memory that is especially allocated:

```
MPI_Alloc_mem() and MPI_Free_mem()
```

These calls reduce to `malloc` and `free` if there is no special memory area; SGI is an example where such memory does exist.

1.4 Collectives

The reference for the commands introduced here can be found in section 1.10.6.

Collectives are operations that involve all processes in a communicator. The simplest example is a broadcast: one processor has some data and all others need to get a copy of it. A collective is a single call, and it blocks on all processors. That does not mean that all processors exit the call at the same time: because of network latency some processors can receive their data later than others.

The collective operations discussed in this section are:

- Broadcast, reduce, and scan: in these operations a single data item is sent from or collected on a ‘root’ process.
- Gather and scatter: here the root has an array from which to send, or in which to collect, the data from the other processors.
- All-to-all, which lets all processors communicate with all others.
- Barrier, which synchronizes all processes, and which separates events before it from those after it.

There are several variants of most collectives. For instance, the gather and reduce calls have a *root of the collective* where information is collected. There is a corresponding ‘all’ variant, where the result is not left just on the root but everywhere. There are also ‘v’ variants where the amount of data coming from or going to each processor is variable.

In addition to these collective operations, there are operations that are said to be ‘collective on their communicator’, but which do not involve data movement. Collective then means that all processors must call this routine; not to do so is an error that will probably manifest itself in ‘hanging’ code. One such example is `MPI_Win_fence`.

1.4.1 Broadcast, reduce, scan

The simplest collective is the broadcast, where one process has some data that needs to be shared with all others. One scenario is that processor zero can parse the commandline arguments of the executable. The call has the following structure:

```
MPI_Bcast( data..., root , comm);
```

The root is the process that is sending its data; see figure 1.10. Typically, it will be the root of a broadcast tree. You see that there is no message tag, because collectives are blocking, so you can have only one active at a time. (In MPI 3 there are non-blocking collectives; see section 1.4.4.)

1. MPI



Figure 1.10: A simple broadcast

It is possible for the data to be an array; in that case MPI acts as if you did a separate scalar broadcast on each array index.

If a processor has only one outgoing connection, the broadcast in figure 1.10 would take a time proportional to the number of processors. One way to ameliorate that is to structure the broadcast in a tree-like fashion. This is depicted in figure 1.11. How does the communication time now depend on the number of processors?

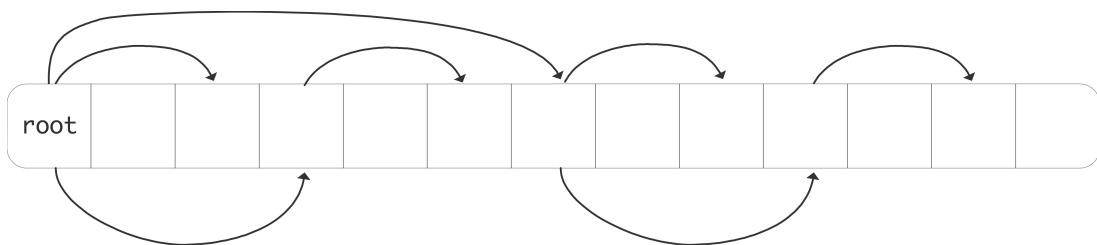


Figure 1.11: A tree-based broadcast

The theory of the complexity of collectives is described in more detail in HPSC-6.2.2.1; see also [1].

The reverse of a broadcast is a reduction:

```
MPI_Reduce( senddata, recvdata..., operator,
             root, comm );
```

Now there is a separate buffer for outgoing data, on all processors, and incoming data, only relevant on the root. Also, you have to indicate how the data is to be combined. Popular choices are `MPI_SUM`, `MPI_PROD` and `MPI_MAX`, but complicated operators such as finding the location of the maximum value exist. You can also define your own operators; section 1.4.1.2.

On processes that are not the root, the receive buffer is ignored. On the root, you have two buffers, but by specifying `MPI_IN_PLACE`, the reduction call uses the value in the receive buffer as the root's contribution to the operation. On the `Allreduce` call, `MPI_IN_PLACE` can be used for the send buffer of every process.

```
void *buffer = malloc(size*sizeof(double));
// write data into the buffer
if (mytid==0) {
    sendbuf = MPI_IN_PLACE; recvbuf = buffer;
} else {
```

```

    sendbuf = buffer;           recvbuf = NULL;
}
MPI_Reduce(sendbuf, recvbuf, size, MPI_DOUBLE, MPI_MAX, 0, comm);

```

1.4.1.1 Scan operations

The `MPI_Scan` operation also performs a reduction, but it keeps the partial results. That is, if processor i contains a number x_i , and \oplus is an operator, then the scan operation leaves $x_0 \oplus \dots \oplus x_i$ on processor i .

```
MPI_Scan( send data, recv data, operator, communicator);
```

This is an inclusive scan operation. The exclusive definition, which computes $x_0 \oplus \dots \oplus x_{i-1}$ on processor i , can easily be derived from the inclusive operation for operations such as `MPI_PLUS` or `MPI_MULT`. This is not the case for `MPI_MIN` or `MPI_MAX`, so there is an exclusive routine

```
MPI_Exscan( send data, recv data, operator, communicator);
```

with the same prototype.

The `MPI_Scan` operation is often useful with indexing data. Suppose that every processor p has a local vector where the number of elements n_p is dynamically determined. In order to translate the local numbering $0 \dots n_p - 1$ to a global numbering one does a scan with the number of local elements as input. The output is then the global number of the first local variable.

Exercise 1.7. Do you use `MPI_Scan` or `MPI_Exscan` for this operation? How would you describe the result of the other scan operation, given the same input?

It is possible to do a *segmented scan*. Let x_i be a series of numbers that we want to sum to X_i as follows. Let y_i be a series of booleans such that

$$\begin{cases} X_i = x_i & \text{if } y_i = 0 \\ X_i = X_{i-1} + x_i & \text{if } y_i = 1 \end{cases}$$

This means that X_i sums the segments between locations where $y_i = 0$ and the first subsequent place where $y_i = 1$. To implement this, you need a user-defined operator

$$\begin{pmatrix} X \\ x \\ y \end{pmatrix} = \begin{pmatrix} X_1 \\ x_1 \\ y_1 \end{pmatrix} \bigoplus \begin{pmatrix} X_2 \\ x_2 \\ y_2 \end{pmatrix} : \begin{cases} X = x_1 + x_2 & \text{if } y_2 == 1 \\ X = x_2 & \text{if } y_2 == 0 \end{cases}$$

This operator is not commutative, and it needs to be declared as such with `MPI_Op_create`.

1.4.1.2 User-defined reductions

1.4.2 Gather and scatter

In gather and scatter, the root collects information from, or conversely spreads it to, all other processes. The difference with reduce and broadcast is that it involves individual information from/to every process. Thus,

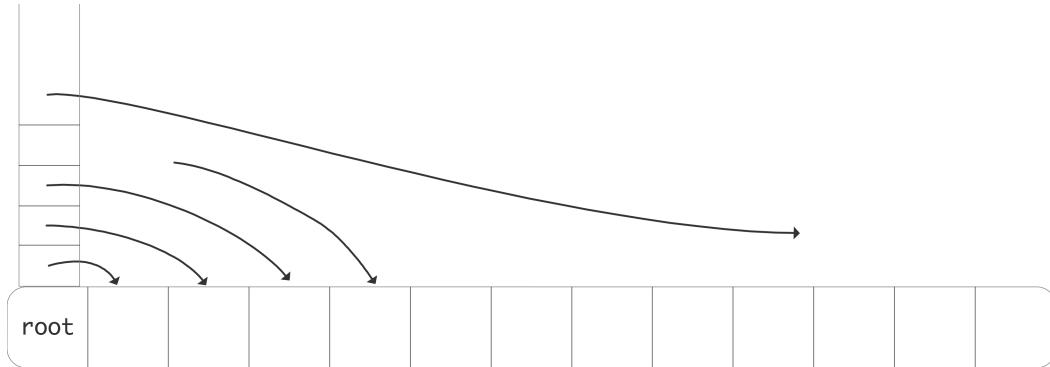


Figure 1.12: A scatter operation

the gather operation typically has an array of items, one coming from each sending process, and scatter has an array, with an individual item for each receiving process; see figure 1.12.

To make this more precise, consider, arbitrarily, the scatter operation. The root process specifies an out buffer:

```
outbuffer, outcount, outtype
```

but the `outcount` is not the length of the buffer: it is the number of elements to send to each process. On the receiving processes other than the root the `outbuffer` arguments are irrelevant.

1.4.3 Variable gather and scatter

In the gather and scatter call above each processor received or sent an identical number of items. In many cases this is appropriate, but sometimes each processor wants or contributes an individual number of items.

Looking again at the scatter, nothing changes in the receive buffer, but now the send buffer on the root is more complicated. It now consists of:

```
outbuffer, array-of-outcounts, array-of-displacements, outtype
```

1.4.4 Non-blocking collectives

MPI version 3 has non-blocking collectives.

1.4.5 Barrier and all-to-all

There are two collectives we have not mentioned yet. A barrier is a call that blocks all processes until they have all reached the barrier call. This call's simplicity is contrasted with its usefulness, which is very limited. It is almost never necessary to synchronize processes through a barrier: for most purposes it does not matter if processors are out of sync. Conversely, collectives (except the new non-blocking ones) introduce a barrier of sorts themselves.

The all-to-all call is a generalization of a scatter and gather: every process is scattering an array of data, and every process is gathering an array of data. There is also a ‘v’ variant of this routine.

1.4.6 Collectives and synchronization

Collectives, other than a barrier, have a synchronizing effect between processors. For instance, in

```
MPI_Bcast( ....data.... root);
MPI_Send(....);
```

the send operations on all processors will occur after the root executes the broadcast. Conversely, in a reduce

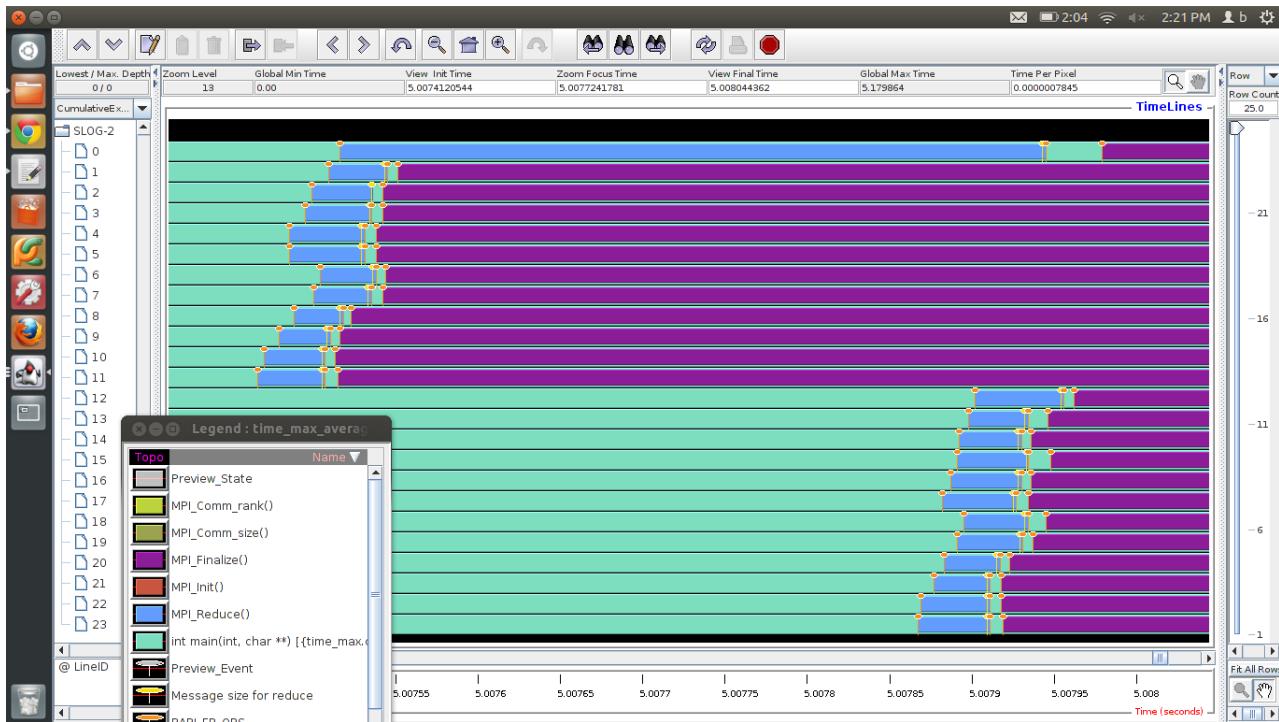


Figure 1.13: Trace of a reduction operation between two dual-socket 12-core nodes

operation the root may have to wait for other processors. This is illustrated in figure 1.13, which gives a TAU trace of a reduction operation on two nodes, with two six-core sockets (processors) each. We see that¹:

- In each socket, the reduction is a linear accumulation;
- on each node, cores zero and six then combine their result;
- after which the final accumulation is done through the network.

We also see that the two nodes are not perfectly in sync, which is normal for MPI applications. As a result, core 0 on the first node will sit idle until it receives the partial result from core 12, which is on the second node.

1. This uses mvapich version 1.6; in version 1.9 the implementation of an on-node reduction has changed to simulate shared memory.

1. MPI

While collectives synchronize in a loose sense, it is not possible to make any statements about events before and after the collectives between processors:

```
...event 1...
MPI_Bcast(...);
...event 2....
```

Consider a specific scenario:

```
switch(rank) {
    case 0:
        MPI_Bcast(buf1, count, type, 0, comm);
        MPI_Send(buf2, count, type, 1, tag, comm);
        break;
    case 1:
        MPI_Recv(buf2, count, type, MPI_ANY_SOURCE, tag, comm, status);
        MPI_Bcast(buf1, count, type, 0, comm);
        MPI_Recv(buf2, count, type, MPI_ANY_SOURCE, tag, comm, status);
        break;
    case 2:
        MPI_Send(buf2, count, type, 1, tag, comm);
        MPI_Bcast(buf1, count, type, 0, comm);
        break;
}
```

Note the MPI_ANY_SOURCE parameter in the receive calls on processor 1. One obvious execution of this would be:

1. The send from 2 is caught by processor 1;
2. Everyone executes the broadcast;
3. The send from 0 is caught by processor 1.

However, it is equally possible to have this execution:

1. Processor 0 starts its broadcast, then executes the send;
2. Processor 1's receive catches the data from 0, then it executes its part of the broadcast;
3. Processor 1 catches the data sent by 2, and finally processor 2 does its part of the broadcast.

1.5 Data types

In many cases the data you send is a single element or an array of some elementary type such as byte, int, or real. You pass the data by specifying the start address and the number of data elements.

Since MPI is a library, and not a language, there is a possibility of confusion here: you first declare your variable or array as, say, `double`, and subsequently declare it to be `MPI_DOUBLE`. First of all, you can make errors, and the compiler will most likely not catch these. Secondly, the definition of datatypes is complicated; for instance there are 32-bit and 64-bit integers. How can you keep your code both correct and portable then?

1.5.1 Elementary data types

MPI_CHAR	only for text data, do not use for small integers
MPI_UNSIGNED_CHAR	
MPI_SIGNED_CHAR	
MPI_SHORT	
MPI_UNSIGNED_SHORT	
C/C++:	MPI_INT
	MPI_UNSIGNED
	MPI_LONG
	MPI_UNSIGNED_LONG
	MPI_FLOAT
	MPI_DOUBLE
	MPI_LONG_DOUBLE
not complete, support for C99 types.	
MPI_CHARACTER	Character(Len=1)
MPI_LOGICAL	
MPI_INTEGER	
Fortran:	MPI_REAL
	MPI_DOUBLE_PRECISION
	MPI_COMPLEX
	MPI_DOUBLE_COMPLEX
	Complex(Kind=Kind(0.d0))

1.5.2 Derived datatypes

1.6 Communicators

A communicator is an object describing a group of processes. In many applications all processes work together closely coupled, and the only communicator you need is MPI_COMM_WORLD. However, there are circumstances where you want one subset of processes to operate independently of another subset. For example:

- If processors are organized in a 2×2 grid, you may want to do broadcasts inside a row or column.
- For an application that includes a producer and a consumer part, it makes sense to split the processors accordingly.

In this section we will see mechanisms for defining new communicators and sending messages between communicators.

An important reason for using communicators is the development of software libraries. If the routines in a library use their own communicator (even if it is a duplicate of the ‘outside’ communicator), there will never be a confusion between message tags inside and outside the library.

1.6.1 Basics

There are three predefined communicators:

- `MPI_COMM_WORLD` comprises all processes that were started together by `mpirun` (or some related program).
- `MPI_COMM_SELF` is the communicator that contains only the current process.
- `MPI_COMM_NULL` is the invalid communicator. Routines that construct communicators can give this as result if an error occurs.

In some applications you will find yourself regularly creating new communicators, using the mechanisms described below. In that case, you should de-allocate communicators with `MPI_Comm_free` when you're done with them.

1.6.2 Creating new communicators

There are various ways of making new communicators. We discuss three mechanisms, from simple to complicated.

1.6.2.1 Duplicating communicators

The reference for the commands introduced here can be found in section [1.10.8.1](#).

With `MPI_Comm_dup` you can make an exact duplicate of a communicator. This may seem pointless, but it is actually very useful for the design of software libraries. Image that you have a code

```
MPI_Isend(...); MPI_Irecv(...);
// library call
MPI_Waitall(...);
```

and suppose that the library has receive calls. Now it is possible that the receive in the library inadvertently catches the message that was sent in the outer environment.

To prevent this confusion, the library should duplicate the outer communicator, and send all messages with respect to its duplicate. Now messages from the user code can never reach the library software, since they are on different communicators.

1.6.2.2 Splitting a communicator

The reference for the commands introduced here can be found in section [1.10.8.2](#).

Splitting a communicator into multiple disjoint communicators can be done with `MPI_Comm_split`. This uses a ‘colour’:

```
MPI_Comm_split( old_comm, colour, new_comm, ... );
```

and all processes in the old communicator with the same colour wind up in a new communicator together. The old communicator still exists, so processes now have two different contexts in which to communicate.

Here is one example of communicator splitting. Suppose your processors are in a two-dimensional grid:

```
MPI_Comm_rank( &mytid );
proc_i = mytid % proc_column_length;
proc_j = mytid / proc_column_length;
```

You can now create a communicator per column:

```
MPI_Comm column_comm;
MPI_Comm_split( MPI_COMM_WORLD, proj_j, &column_comm );
```

and do a broadcast in that column:

```
MPI_Bcast( data, /* tag: */ 0, column_comm );
```

Because of the SPMD nature of the program, you are now doing in parallel a broadcast in every processor column. Such operations often appear in *dense linear algebra*.

1.6.2.3 Process groups

The most general mechanism is based on groups: you can extract the group from a communicator, combine different groups, and form a new communicator from the resulting group.

The group mechanism is more involved. You get the group from a communicator, or conversely make a communicator from a group with `MPI_Comm_group` and `MPI_Comm_create`:

```
MPI_Comm_group( comm, &group );
MPI_Comm_create( old_comm, group, &new_comm );
```

and groups are manipulated with `MPI_Group_incl`, `MPI_Group_excl`, `MPI_Group_difference` and a few more.

You can name your communicators with `MPI_Comm_set_name`, which could improve the quality of error messages when they arise.

1.6.3 Intra-communicators

We start by exploring the mechanisms for creating a communicator that encompasses a subset of `MPI_COMM_WORLD`.

The most general mechanism for creating communicators is through process groups: you can query the group of processes of a communicator, manipulate groups, and make a new communicator out of a group you have formed.

```
MPI_COMM_GROUP (comm, group, ierr)
MPI_COMM_CREATE (MPI_Comm comm, MPI_Group group, MPI_Comm newcomm, ierr)

MPI_GROUP_UNION(group1, group2, newgroup, ierr)
MPI_GROUP_INTERSECTION(group1, group2, newgroup, ierr)
MPI_GROUP_DIFFERENCE(group1, group2, newgroup, ierr)
```

```
MPI_GROUP_INCL(group, n, ranks, newgroup, ierr)
MPI_GROUP_EXCL(group, n, ranks, newgroup, ierr)

MPI_GROUP_SIZE(group, size, ierr)
MPI_GROUP_RANK(group, rank, ierr)
```

1.6.4 Inter-communicators

If two disjoint communicators exist, it may be necessary to communicate between them. This can of course be done by creating a new communicator that overlaps them, but this would be complicated: since the ‘inter’ communication happens in the overlap communicator, you have to translate its ordering into those of the two worker communicators. It would be easier to express messages directly in terms of those communicators, and this can be done with ‘inter-communicators’.

```
MPI_INTERCOMM_CREATE (local_comm, local_leader, bridge_comm, remote_leader,
```

After this, the intercommunicator can be used in collectives such as

```
MPI_BCAST (buff, count, dtype, root, comm, ierr)
```

- In group A, the root process passes MPI_ROOT as ‘root’ value; all others use MPI_NULL_PROC.
- In group B, all processes use a ‘root’ value that is the rank of the root process in the root group.

Gather and scatter behave similarly; the allgather is different: all send buffers of group A are concatenated in rank order, and places on all processes of group B.

Inter-communicators can be used if two groups of process work asynchronously with respect to each other; another application is fault tolerance (section 1.8.4).

1.7 Hybrid programming: MPI and threads

The reference for the commands introduced here can be found in section 1.10.11.

It is not automatic that a program or a library is *thread-safe*. A user can request a certain level of multi-threading with `MPI_Init_thread`, and the system will respond what the highest supported level is.

MPI can be thread-safe on the following levels:

- An MPI implementation can forbid any multi-threading;
- it can allow one thread to make MPI calls;
- it can allow one thread *at a time* to make MPI calls;
- it can allow arbitrary multi-threaded behaviour in MPI calls.

Some points.

- MPI can not distinguish between threads: the communicator rank identifies a process, and is therefore identical for all threads.
- A message sent to a process can be received by any thread that has issued a receive call with the right source/tag specification.
- Multi-threaded calls to an MPI routine have the semantics of an unspecified sequence of calls.
- A blocking MPI call only blocks the thread that makes it.

1.8 Leftover topics

1.8.1 Getting message information

In some circumstances the recipient may not know all details of a message.

- If you are expecting multiple incoming messages, it may be most efficient to deal with them in the order in which they arrive. For that, you have to be able to ask ‘who did this message come from, and what is in it’.
- Maybe you know the sender of a message, but the amount of data is unknown. In that case you can overallocate your receive buffer, and after the message is received ask how big it was, or you can ‘probe’ an incoming message and allocate enough data when you find out how much data is being sent.

1.8.1.1 Status object

The receive calls you saw above has a status argument. If you precisely know what is going to be sent, this argument tells you nothing new. Therefore, there is a special value `MPI_STATUS_IGNORE` that you can supply instead of a status object, which tells MPI that the status does not have to be reported. For routines such as `MPI_Waitany` where an array of statuses is needed, you can supply `MPI_STATUSES_IGNORE`.

However, if you expect data from multiple senders, or the amount of data is indeterminate, the status will give you that information.

The `MPI_Status` object is a structure with the following freely accessible members: `MPI_SOURCE`, `MPI_TAG`, and `MPI_ERROR`. There is also opaque information: the amount of data received can be retrieved by a function call to `MPI_Get_count`.

```
int MPI_Get_count(
    MPI_Status *status,
    MPI_Datatype datatype,
    int *count
);
```

This may be necessary since the `count` argument to `MPI_Recv` is the buffer size, not an indication of the actually expected number of data items.

1.8.1.2 Probing messages

MPI receive calls specify a receive buffer, and its size has to be enough for any data sent. In case you really have no idea how much data is being sent, and you don't want to overallocate the receive buffer, you can use a 'probe' call.

The calls `MPI_Probe`, `MPI_Iprobe`, accept a message, but do not copy the data. Instead, when probing tells you that there is a message, you can use `MPI_Get_count` to determine its size, allocate a large enough receive buffer, and do a regular receive to have the data copied.

1.8.2 Semantics

MPI processes are only synchronized to a certain extent, so you may wonder what guarantees there are that running a code twice will give the same result. You need to consider two cases: first of all, if the two runs are on different numbers of processors there are already numerical problems; see HPSC-[3.3.6](#).

Let us then limit ourselves to two runs on the same set of processors. In that case, MPI is deterministic as long as you do not use wildcards such as `MPI_ANY_SOURCE`. Formally, MPI messages are 'non-overtaking': two messages between the same sender-receiver pair will arrive in sequence.

1.8.3 Error handling

The reference for the commands introduced here can be found in section [1.10.9](#).

Errors in normal programs can be tricky to deal with; errors in parallel programs can be even harder. This is because in addition to everything that can go wrong with a single executable (floating point errors, memory violation) you now get errors that come from faulty interaction between multiple executables.

A few examples of what can go wrong:

- MPI errors: an MPI routine can abort for various reasons, such as receiving much more data than its buffer can accomodate. Such errors, as well as the more common type mentioned above, typically cause your whole execution to abort. That is, if one incarnation of your executable aborts, the MPI runtime will kill all others.
- Deadlocks and other hanging executions: there are various scenarios where your processes individually do not abort, but are all waiting for each other. This can happen if two processes are both waiting for a message from each other, and this can be helped by using non-blocking calls. In another scenario, through an error in program logic, one process will be waiting for more messages (including non-blocking ones) than are sent to it.

The MPI library has a general mechanism for dealing with errors that it detects. The default behaviour, where the full run is aborted, is equivalent to your code having the following call ²:

```
MPI_Comm_set_errhandler(MPI_COMM_WORLD, MPI_ERRORS_ARE_FATAL);
```

Another simple possibility is to specify

2. The routine `MPI_Errhandler_set` is deprecated.

```
MPI_Comm_set_errhandler(MPI_COMM_WORLD,MPI_ERRORS_RETURN);
```

which gives you the opportunity to write code that handles the error return value.

In most cases where an MPI error occurs a complete abort is the sensible thing, since there are few ways to recover. The second possibility can for instance be used to print out debugging information:

```
ierr = MPI_Something();
if (ierr!=0) {
    // print out information about what your programming is doing
    MPI_Abort();
}
```

For instance,

```
Fatal error in MPI_Waitall:
See the MPI_ERROR field in MPI_Status for the error code
```

You could code this as

```
MPI_Comm_set_errhandler(MPI_COMM_WORLD,MPI_ERRORS_RETURN);
ierr = MPI_Waitall(2*ntids-2,requests,status);
if (ierr!=0) {
    char errtxt[200];
    for (int i=0; i<2*ntids-2; i++) {
        int err = status[i].MPI_ERROR; int len=200;
        MPI_Error_string(err,errtxt,&len);
        printf("Waitall error: %d %s\n",err,errtxt);
    }
    MPI_Abort(MPI_COMM_WORLD,0);
}
```

One cases where errors can be handled is that of *MPI file I/O/MPI/I/O*: if an output file has the wrong permissions, code can possibly progress without writing data, or writing to a temporary file.

1.8.4 Fault tolerance

Processors are not completely reliable, so it may happen that one ‘breaks’: for software or hardware reasons it becomes unresponsive. For an MPI program this means that it becomes impossible to send data to it, and any collective operation involving it will hang. Can we deal with this case? Yes, but it involves some programming.

First of all, one of the possible MPI error return codes (section ??) is `MPI_ERR_COMM`, which can be returned if a processor in the communicator is unavailable. You may want to catch this error, and add a ‘replacement processor’ to the program. For this, the `MPI_Comm_spawn` can be used:

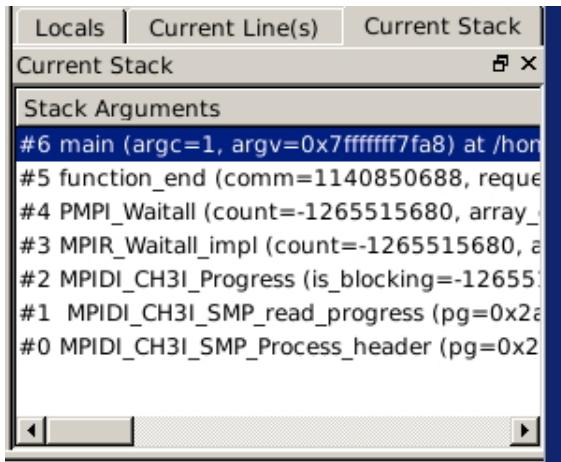


Figure 1.14: A stack trace, showing the MPI calls.

```
int MPI_Comm_spawn(char *command, char *argv[], int maxprocs, MPI_Info info,
                    int root, MPI_Comm comm, MPI_Comm *intercomm,
                    int array_of_errcodes[])
```

But this requires a change of program design: the communicator containing the new process(es) is not part of the old `MPI_COMM_WORLD`, so it is better to set up your code as a collection of inter-communicators to begin with.

1.8.5 Profiling

`MPI_Wtime` gives *wall clock time* from unspecified starting point.

MPI allows you to write your own profiling interface. To make this possible, every routine `MPI_Something` calls a routine `PMPPI_Something` that does the actual work. You can now write your `MPI_...` routine which calls `PMPPI_...`, and inserting your own profiling calls. As you can see in figure 1.14, normally only the `PMPPI` routines show up in the stack trace.

Does the standard mandate this?

1.8.6 Debugging

There are various ways of debugging an MPI program. Typically there are two cases. In the simple case your program can have a serious error in logic which shows up even with small problems and a small number of processors. In the more difficult case your program can only be run on large scale, or the problem only shows up when you run at large scale. For the second case you, unfortunately, need a dedicated debugging tool, and of course the good ones are expensive. In the first case there are some simpler solutions.

1.8.6.1 Small scale debugging

If your program hangs or crashes even with small numbers of processors, you can try debugging on your local desktop or laptop computer:

```
mpirun -np <n> xterm -e gdb yourprogram
```

This starts up a number of X terminals, each of which runs your program. The magic of `mpirun` makes sure that they all collaborate on a parallel execution of that program. If your program needs commandline arguments, you have to type those in every xterm:

```
run <argument list>
```

See appendix A.2 for more about debugging with `gdb`.

This approach is not guaranteed to work, since it depends on your ssh setup; see the discussion in <http://www.open-mpi.org/faq/?category=debugging#serial-debuggers>.

1.8.6.2 Large scale debugging

Check out `ddt` or `TotalView`.

1.8.6.3 Memory debugging of MPI programs

The commercial parallel debugging tools typically have a memory debugger. For an open source solution you can use `valgrind`, but that requires some setup during installation. See <http://valgrind.org/docs/manual/mc-manual.html#mc-manual.mpiwrap> for details.

1.8.7 Language issues

MPI is typically written in C, what if you program Fortran?

Assumed shape arrays can be a problem: they need to be copied. That's a problem with `ISEND`.

The `C++` interface is deprecated as of *MPI 2.2*. It is unclear what is happening.

1.8.8 The origin of one-sided communication in ShMem

The *Cray T3E* had a library called `shmem` which offered a type of shared memory. Rather than having a true global address space it worked by supporting variables that were guaranteed to be identical between processors, and indeed, were guaranteed to occupy the same location in memory. Variables could be declared to be shared a ‘symmetric’ pragma or directive; their values could be retrieved or set by `shmem_get` and `shmem_put` calls.

1.9 Programming projects

1.9.1 A style guide for project submissions

Here are some guidelines for how to submit assignments and projects. As a general rule, consider programming as an experimental science, and your writeup as a report on some tests you have done: explain the problem you're addressing, your strategy, your results.

Structure of your writeup Most of the exercises in this book test whether you are able to code the solution to a certain problem. That does not mean that turning in the code is sufficient, nor code plus sample output. Turn in a writeup in pdf form that was generated from a text processing program such as Word or (preferably) L^AT_EX (for a tutorial, see HPSC-B.12). Your writeup should have

- The relevant fragments of your code,
- an explanation of your algorithms or solution strategy,
- a discussion of what you observed,
- graphs of runtimes and TAU plots; see A.3.

Observe, measure, hypothesize, deduce In most applications of computing machinery we care about the efficiency with which we find the solution. Thus, make sure that you do measurements. In general, make observations that allow you to judge whether your program behaves the way you would expect it to.

Quite often your program will display unexpected behaviour. It is important to observe this, and hypothesize what the reason might be for your observed behaviour.

Including code If you include code samples in your writeup, make sure they look good. For starters, use a mono-spaced font. In L^AT_EX, you can use the `verbatim` environment or the `verbatiminput` command. In that section option the source is included automatically, rather than cut and pasted. This is to be preferred, since your writeup will stay current after you edit the source file.

Including whole source files makes for a long and boring writeup. The code samples in this book were generated as follows. In the source files, the relevant snippet was marked as

```
... boring stuff
#pragma samplex
    .. interesting! ..
#pragma end
... more boring stuff
```

The files were then processed with the following command line (actually, included in a makefile, which requires doubling the dollar signs):

```
for f in *.{c,cxx,h} ; do
    cat $x | awk 'BEGIN {f=0}
                    /#pragma end/ {f=0}
                    f==1 {print $0 > file}'
```

```
    /pragma/ {f=1; file=$2 }  
,  
done
```

which gives (in this example) a file `samplex`. Other solutions are of course possible.

Code formatting Code without proper indentation is very hard to read. Fortunately, most editors have some knowledge of the syntax of the most popular languages. The `emacs` editor will, most of the time, automatically activate the appropriate mode based on the file extension. If this does not happen, you can activate a mode by `ESC x fortran-mode` et cetera, or by putting the string `-- fortran --` in a comment on the first line of your file.

The `vi` editor also has syntax support: use the commands `:syntax on` to get syntax colouring, and `:set cindent` to get automatic indentation while you're entering text. Some of the more common questions are addressed in <http://stackoverflow.com/questions/97694/auto-indent-spaces-with-c-indentation-in-vi>

Running your code A single run doesn't prove anything. For a good report, you need to run your code for more than one input dataset (if available) and in more than one processor configuration. When you choose problem sizes, be aware that an average processor can do a billion operations per second: you need to make your problem large enough for the timings to rise above the level of random variations and startup phenomena.

When you run a code in parallel, beware that on clusters the behaviour of a parallel code will always be different between one node and multiple nodes. On a single node the MPI implementation is likely optimized to use the shared memory. This means that results obtained from a single node run will be unrepresentative. In fact, in timing and scaling tests you will often see a drop in (relative) performance going from one node to two. Therefore you need to run your code in a variety of scenarios, using more than one node.

Repository organization If you submit your work through a repository, make sure you organize your submissions in subdirectories, and that you give a clear name to all files.

1.9.2 Warmup exercises

We start with some simple exercises.

1.9.2.1 Hello world

The exercises in this section are about the routines introduced in section 1.2.3; for the reference information see section ??.

First of all we need to make sure that you have a working setup for parallel jobs. The example program `helloworld.c` does the following:

```
// helloworld.c
MPI_Init(&argc,&argv);
MPI_Comm_size(MPI_COMM_WORLD,&ntids);
MPI_Comm_rank(MPI_COMM_WORLD,&mytid);
printf("Hello, this is processor %d out of %d\n",mytid,ntids);
MPI_Finalize();
```

Compile this program and run it in parallel. Make sure that the processors do *not* all say that they are processor 0 out of 1!

1.9.2.2 Trace output

We want to make trace files of the parallel runs, for which we'll use the TAU utility of the University of Oregon. (For documentation, go to <http://www.cs.uoregon.edu/Research/tau/docs.php>.) Here are the steps:

- Load two modules:

```
module load tau
module load jdk64
```

- Recompile your program with `make yourprog`. You'll notice a lot more output: that is the TAU preprocessor.
- Now run your program, setting environment variables `TAU_TRACE` and `TAU_PROFILE` to 1, and `TRACEDIR` and `PROFILEDIR` to where you want the output to be. Big shortcut: do

```
make submit EXECUTABLE=yourprog
```

for a batch job or

```
make idevrun EXECUTABLE=yourprog
```

for an interactive parallel run. These last two set all variables for you. See if you can find where the output went...

- Now you need to postprocess the TAU output. Do `make tau EXECUTABLE=yourprog` and you'll get a file `taulog_yourprog.slog2` which you can view with the `jumpshot` program.

1.9.2.3 Collectives

It is a good idea to be able to collect statistics, so before we do anything interesting, we will look at MPI collectives; section 1.4.

Take a look at `time_max.cxx`. This program sleeps for a random number of seconds:

```
// time_max.cxx
wait = (int) ( 6.*rand() / (double)RAND_MAX );
tstart = MPI_Wtime();
sleep(wait);
```

```
tstop = MPI_Wtime();
jitter = tstop-tstart-wait;
```

and measures how long the sleep actually was:

```
if (mytid==0)
    sendbuf = MPI_IN_PLACE;
else sendbuf = (void*)&jitter;
MPI_Reduce(sendbuf, (void*)&jitter, 1, MPI_DOUBLE, MPI_MAX, 0, comm);
```

In the code, this quantity is called ‘jitter’, which is a term for random deviations in a system.

Exercise 1.8. Change this program to compute the average jitter by changing the reduction operator.

Exercise 1.9. Now compute the standard deviation

$$\sigma = \sqrt{\frac{\sum_i (x_i - m)^2}{n}}$$

where m is the average value you computed in the previous exercise.

- Solve this exercise twice: once by following the reduce by a broadcast operation and once by using an Allreduce.
- Run your code both on a single cluster node and on multiple nodes, and inspect the TAU trace. Some MPI implementations are optimized for shared memory, so the trace on a single node may not look as expected.
- Can you see from the trace how the allreduce is implemented?

Exercise 1.10. Finally, use a gather call to collect all the values on processor zero, and print them out. Is there any process that behaves very differently from the others?

For each exercise, submit code, a TAU trace, and an analysis of what you see in the traces. Submit your work by leaving a code, graphics, and a writeup in your repository.

1.9.2.4 Linear arrays of processors

In this section you are going to write a number of variations on a very simple operation: all processors pass a data item to the processor with the next higher number.

- In the file `linear-serial.c` you will find an implementation using blocking send and receive calls.
- You will change this code to use non-blocking sends and receives; they require an `MPI_Wait` call to finalize them.
- Next, you will use `MPI_Sendrecv` to arrive at a synchronous, but deadlock-free implementation.
- Finally, you will use two different one-sided scenarios.

In the reference code `linear-serial.c`, each process defines two buffers:

```
// linear-serial.c
int my_number = mytid, other_number=-1.;
```

where `other_number` is the location where the data from the left neighbour is going to be stored.

To check the correctness of the program, there is a gather operation on processor zero:

```
int *gather_buffer=NULL;
if (mytid==0) {
    gather_buffer = (int*) malloc(ntids*sizeof(int));
    if (!gather_buffer) MPI_Abort(comm,1);
}
MPI_Gather(&other_number,1,MPI_INT,
            gather_buffer,1,MPI_INT, 0,comm);
if (mytid==0) {
    int i,error=0;
    for (i=0; i<ntids; i++)
        if (gather_buffer[i]!=i-1) {
            printf("Processor %d was incorrect: %d should be %d\n",
                   i,gather_buffer[i],i-1);
            error =1;
        }
    if (!error) printf("Success!\n");
    free(gather_buffer);
}
```

1.9.2.4.1 Coding with blocking calls Passing data to a neighbouring processor should be a very parallel operation. However, if we code this naively, with `MPI_Send` and `MPI_Recv`, we get an unexpected serial behaviour, as was explained in section 1.3.1.

```
if (mytid<ntids-1)
    MPI_Ssend( /* data: */ &my_number,1,MPI_INT,
               /* to: */ mytid+1, /* tag: */ 0, comm);
if (mytid>0)
    MPI_Recv( /* data: */ &other_number,1,MPI_INT,
              /* from: */ mytid-1, 0, comm, &status);
```

(Note that this uses an `Ssend`; see section 1.3.1.2 for the explanation why.)

Exercise 1.11. Compile and run this code, and generate a TAU trace file. Confirm that the execution is serial. Does replacing the `Ssend` by `Send` change this?

Let's clean up the code a little.

Exercise 1.12. First write this code more elegantly by using `MPI_PROC_NULL`.

1.9.2.4.2 A better blocking solution The easiest way to prevent the serialization problem of the previous exercises is to use the `MPI_Sendrecv` call. This routine acknowledges that often a processor will have a receive call whenever there is a send. For border cases where a send or receive is unmatched you can use `MPI_PROC_NULL`.

Exercise 1.13. Rewrite the code using `MPI_Sendrecv`. Confirm with a TAU trace that execution is no longer serial.

Note that the `Sendrecv` call itself is still blocking, but at least the ordering of its constituent send and recv are no longer ordered in time.

1.9.2.4.3 Non-blocking calls The other way around the blocking behaviour is to use `Irecv` and `Isend` calls, which do not block. Of course, now you need a guarantee that these send and receive actions are concluded; in this case, use `MPI_Waitall`.

Exercise 1.14. Implement a fully parallel version by using `MPI_Isend` and `MPI_Irecv`.

1.9.2.4.4 One-sided communication Another way to have non-blocking behaviour is to use one-sided communication. During a Put or Get operation, execution will only block while the data is being transferred out of or into the origin process, but it is not blocked by the target. Again, you need a guarantee that the transfer is concluded; here use `MPI_Win_fence`.

Exercise 1.15. Write two versions of the code: one using `MPI_Put` and one with `MPI_Get`. Make TAU traces.

Investigate blocking behaviour through TAU visualizations.

Exercise 1.16. If you transfer a large amount of data, and the target processor is occupied, can you see any effect on the origin? Are the fences synchronized?

1.9.3 Mandelbrot set

If you've never heard the name *Mandelbrot set*, you probably recognize the picture. Its formal definition is as follows:

A point c in the complex plane is part of the Mandelbrot set if the series x_n defined by

$$\begin{cases} x_0 = 0 \\ x_{n+1} = x_n^2 + c \end{cases}$$

satisfies

$$\forall n : |x_n| \leq 2.$$

It is easy to see that only points c in the bounding circle $|c| < 2$ qualify, but apart from that it's hard to say much without a lot more thinking. Or computing; and that's what we're going to do.

In this set of exercises you are going to take an example program `mandel_main.cxx` and extend it to use a variety of MPI programming constructs. This program has been set up as a master-worker model: there is one master processor (for a change this is the last processor, rather than zero) which gives out work to, and accepts results from, the worker processors. It then takes the results and construct an image file from them.

1.9.3.1 Invocation

The `mandel_main` program is called as

```
mpirun -np 123 mandel_main steps 456 iters 789
```

where the `steps` parameter indicates how many steps in x, y direction there are in the image, and `iters` gives the maximum number of iterations in the `belong` test.

If you forget the parameter, you can call the program with

```
mandel_serial -h
```

and it will print out the usage information.

1.9.3.2 Tools

The driver part of the Mandelbrot program is simple. There is a `circle` object that can generate coordinates

```
// mandel.h
class circle {
public :
    circle(int pxls,int bound,int bs);
    void next_coordinate(struct coordinate& xy);
    int is_valid_coordinate(struct coordinate xy);
    void invalid_coordinate(struct coordinate& xy);
```

and a global routine that tests whether a coordinate is in the set, at least up to an iteration bound. It returns zero if the series from the given starting point has not diverged, or the iteration number in which it diverged if it did so.

```
int belongs(struct coordinate xy,int itbound) {
    double x=xy.x, y=xy.y; int it;
    for (it=0; it<itbound; it++) {
        double xx,yy;
        xx = x*x - y*y + xy.x;
        yy = 2*x*y + xy.y;
        x = xx; y = yy;
        if (xx+yy>4.) {
            return it;
        }
    }
    return 0;
}
```

In the former case, the point could be in the Mandelbrot set, and we colour it black, in the latter case we give it a colour depending on the iteration number.

```

if (iteration==0)
    memset(colour,0,3*sizeof(float));
else {
    float rfloat = ((float) iteration) / workcircle->infty;
    colour[0] = rfloat;
    colour[1] = max((float)0,(float)(1-2*rfloat));
    colour[2] = max((float)0,(float)(2*(rfloat-.5)));
}

```

We use a fairly simple code for the worker processes: they execute a loop in which they wait for input, process it, return the result.

```

void queue::wait_for_work(MPI_Comm comm,circle *workcircle) {
    MPI_Status status; int ntids;
    MPI_Comm_size(comm,&ntids);
    int stop = 0;

    while (!stop) {
        struct coordinate xy;
        int res;

        MPI_Recv(&xy,1,coordinate_type,ntids-1,0, comm,&status);
        stop = !workcircle->is_valid_coordinate(xy);
        if (stop) res = 0;
        else {
            res = belongs(xy,workcircle->infty);
        }
        MPI_Send(&res,1,MPI_INT,ntids-1,0, comm);
    }
    return;
}

```

A very simple solution using blocking sends on the master is given:

```

// mandel_serial.cxx
class serialqueue : public queue {
private :
    int free_processor;
public :
    serialqueue(MPI_Comm queue_comm,circle *workcircle)
        : queue(queue_comm,workcircle) {
        free_processor=0;
    };
    /**
        The 'addtask' routine adds a task to the queue. In this

```

simple case it immediately sends the task to a worker and waits for the result, which is added to the image.

This routine is only called with valid coordinates; the calling environment will stop the process once an invalid coordinate is encountered.

```
/*
int addtask(struct coordinate xy) {
    MPI_Status status; int contribution, err;

    err = MPI_Send(&xy, 1, coordinate_type,
                   free_processor, 0, comm); CHK(err);
    err = MPI_Recv(&contribution, 1, MPI_INT,
                   free_processor, 0, comm, &status); CHK(err);

    coordinate_to_image(xy, contribution);
    total_tasks++;
    free_processor++;
    if (free_processor==ntids-1)
        // wrap around to the first again
        free_processor = 0;
    return 0;
};
```

Exercise 1.17. Explain why this solution is very inefficient. Make a trace of its execution that bears this out.

1.9.3.3 Bulk task scheduling

The previous section showed a very inefficient solution, but that was mostly intended to set up the code base. If all tasks take about the same amount of time, you can give each process a task, and then wait on them all to finish. A first way to do this is with non-blocking sends.

Exercise 1.18. Code a solution where you give a task to all worker processes using non-blocking sends and receives, and then wait for these tasks with `MPI_Waitall` to finish before you give a new round of data to all workers. Make a trace of the execution of this and report on the total time.

You can do this by writing a new class that inherits from `queue`, and that provides its own `addtask` method:

```
// mandel_bulk.cxx
class bulkqueue : public queue {
public :
    bulkqueue(MPI_Comm queue_comm, circle *workcircle)
        : queue(queue_comm, workcircle) {
```

You will also have to override the `complete` method: when the circle object indicates that all coordinates have been generated, not all workers will be busy, so you need to supply the proper `MPI_Waitall` call.

1.9.3.4 Collective task scheduling

Another implementation of the bulk scheduling of the previous section would be through using collectives.

Exercise 1.19. Code a solution which uses scatter to distribute data to the worker tasks, and gather to collect the results. Is this solution more or less efficient than the previous?

1.9.3.5 Asynchronous task scheduling

At the start of section 1.9.3.3 we said that bulk scheduling mostly makes sense if all tasks take similar time to complete. In the Mandelbrot case this is clearly not the case.

Exercise 1.20. Code a fully dynamic solution that uses `MPI_Probe` or `MPI_Waitany`. Make an execution trace and report on the total running time.

1.9.3.6 One-sided solution

Let us reason about whether it is possible (or advisable) to code a one-sided solution to computing the Mandelbrot set. With active target synchronization you could have an exposure window on the host to which the worker tasks would write. To prevent conflicts you would allocate an array and have each worker write to a separate location in it. The problem here is that the workers may not be sufficiently synchronized because of the differing time for computation.

Consider then passive target synchronization. Now the worker tasks could write to the window on the master whenever they have something to report; by locking the window they prevent other tasks from interfering. After a worker writes a result, it can get new data from an array of all coordinates on the master.

It is hard to get results into the image as they become available. For this, the master would continuously have to scan the results array. Therefore, constructing the image is easiest done when all tasks are concluded.

1.9.4 Data parallel grids

1.9.4.1 A realistic programming example

In this section we will gradually build a semi-realistic example program. To get you started some pieces have already been written: as a starting point look at `code/mpi/c/grid.cxx`.

1.9.4.1.1 Description of the problem With this example you will investigate several strategies for implementing a simple iterative method. Let's say you have a two-dimensional grid of datapoints $G = \{g_{ij} : 0 \leq i < n_i, 0 \leq j < n_j\}$ and you want to compute G' where

$$g'_{ij} = 1/4 \cdot (g_{i+1,j} + g_{i-1,j} + g_{i,j+1} + g_{i,j-1}). \quad (1.1)$$

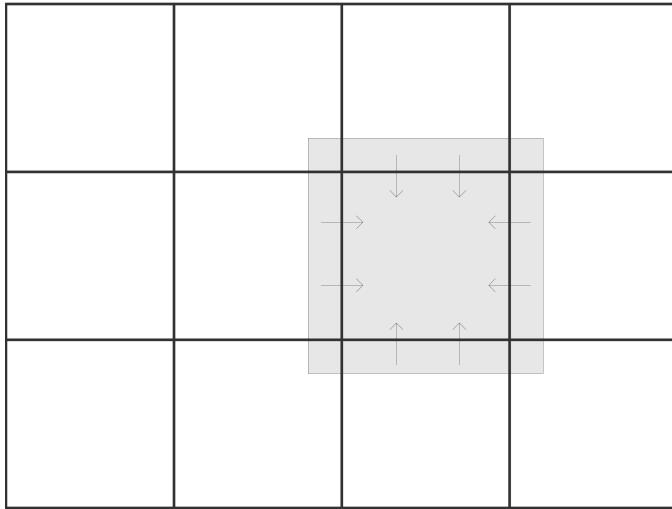


Figure 1.15: A grid divided over processors, with the ‘ghost’ region indicated

This is easy enough to implement sequentially, but in parallel this requires some care.

Let’s divide the grid G and divide it over a two-dimension grid of $p_i \times p_j$ processors. (Other strategies exist, but this one scales best; see section HPSC-6.4.) Formally, we define two sequences of points

$$0 = i_0 < \dots < i_{p_i} < i_{p_i+1} = n_i, \quad 0 < j_0 < \dots < j_{p_j} < i_{p_j+1} = n_j$$

and we say that processor (p, q) computes g_{ij} for

$$i_p \leq i < i_{p+1}, \quad j_q \leq j < j_{q+1}.$$

From formula (1.1) you see that the processor then needs one row of points on each side surrounding its part of the grid. A picture makes this clear; see figure 1.15. These elements surrounding the processor’s own part are called the *halo* or *ghost region* of that processor.

The problem is now that the elements in the halo are stored on a different processor, so communication is needed to gather them. In the upcoming exercises you will have to use different strategies for doing so.

1.9.4.1.2 Code basics The program needs to read the values of the grid size and the processor grid size from the commandline, as well as the number of iterations. This routine does some error checking: if the number of processors does not add up to the size of `MPI_COMM_WORLD`, a nonzero error code is returned.

```
ierr = parameters_from_commandline
(argc, argv, comm, &ni, &nj, &pi, &pj, &nit);
if (ierr) return MPI_Abort(comm, 1);
```

From the processor parameters we make a processor grid object:

```
processor_grid *pgrid = new processor_grid(comm,pi,pj);
```

and from the numerical parameters we make a number grid:

```
number_grid *grid = new number_grid(pgrid,ni,nj);
```

Number grids have a number of methods defined. To set the value of all the elements belonging to a processor to that processor's number:

```
grid->set_test_values();
```

To set random values:

```
grid->set_random_values();
```

If you want to visualize the whole grid, the following call gathers all values on processor zero and prints them:

```
grid->gather_and_print();
```

Next we need to look at some data structure details.

The definition of the `number_grid` object starts as follows:

```
class number_grid {  
public:  
    processor_grid *pgrid;  
    double *values,*shadow;
```

where `values` contains the elements owned by the processor, and `shadow` is intended to contain the values plus the ghost region. So how does `shadow` receive those values? Well, the call looks like

```
grid->build_shadow();
```

and you will need to supply the implementation of that. Once you've done so, there is a routine that prints out the shadow array of each processor

```
grid->print_shadow();
```

This routine does the sequenced printing that you implemented in exercise ??.

In the file `code/mpi/c/grid_impl.cxx` you can see several uses of the macro `INDEX`. This translates from a two-dimensional coordinate system to one-dimensional. Its main use is letting you use (i,j) coordinates for indexing the processor grid and the number grid: for processors you need the translation to the linear rank, and for the grid you need the translation to the linear array that holds the values.

A good example of the use of `INDEX` is in the `number_grid::relax` routine: this takes points from the shadow array and averages them into a point of the `values` array. (To understand the reason for this

particular averaging, see HPSC-4.2.2.2 and HPSC-5.5.3.) Note how the INDEX macro is used to index in a $i\text{length} \times j\text{length}$ target array values, while reading from a $(i\text{length} + 2) \times (j\text{length} + 2)$ source array shadow.

```
for (i=0; i<iLength; i++) {
    for (j=0; j<jLength; j++) {
        int c=0;
        double new_value=0.;
        for (c=0; c<5; c++) {
            int ioff=i+1+iOffsetsets[c], joff=j+1+jOffsetsets[c];
            new_value += coefficients[c] *
                shadow[ INDEX(ioff, joff, iLength+2, jLength+2) ];
        }
        values[ INDEX(i, j, iLength, jLength) ] = new_value/8.;
    }
}
```

1.10 Reference to the routines

This section gives reference information and illustrative examples of the use of MPI. While the code snippets given here should be enough, full programs can be found in the repository for this book <https://bitbucket.org/VictorEijkhout/parallel-computing-book>.

1.10.1 Basics

1.10.1.1 MPI setup

This reference section gives the syntax for routines introduced in section 1.2.1.

Every MPI program has to start with

```
MPI_Init(&argc, &argv);
```

where `argc` and `argv` are the arguments of a C language main program:

```
int main(int argc, char **argv) {
    ....
    return 0;
}
```

The regular way to conclude an MPI program is through

```
MPI_Finalize();
```

but an abnormal end to a run can be forced by

```
    MPI_Abort(comm,value);
```

This aborts execution on all processes associated with the communicator, but many implementations simply abort all processes. The `value` parameter is returned to the environment.

The commandline arguments can only be guaranteed to be passed correctly to process zero. Here is a fragment of code that shows use of commandline arguments. The program `examples/mpi/c/init.c` takes a single integer commandline argument. If the user forgets to specify an argument or specifies `-h`, a usage message is printed and the program aborts, otherwise the parameter is broadcast to all processes.

```
// init.c
MPI_Comm_rank(comm,&mytid);
if (mytid==0) {
    if ( argc==1 || // the program is called without parameter
        ( argc>1 && !strcmp(argv[1],"-h") ) // user asked for help
        ) {
        printf("\nUsage: init [0-9]+\n");
        MPI_Abort(comm,1);
    }
    input_argument = atoi(argv[1]);
}
MPI_Bcast(&input_argument,1,MPI_INT,0,comm);
```

1.10.1.2 MPI ranks and communicator sizes

This reference section gives the syntax for routines introduced in section 1.2.3.

Every MPI process has its own local storage. So if you pretend that all these small arrays are really together one big array, you need to know where each piece fits in the whole. For this you need to know at the very least how many processes there are and what the rank of a process is.

The following example creates a distributed array that will contains the values of a function at equidistant points in the interval $[0, 1]$. The code implements this as follows:

1. each process has an array of 10 points,
2. so the distributed array has a length of 10 times the number of processes;
3. of this array, each process has 10 points, starting at 10 times the process id.
4. To fill in values in the local array, the process iterates only over its local points,
5. but it needs to calculate the global coordinate of those points.

```
// local.cxx
int nlocalpoints = 10,
    ntotal_points = ntids*nlocalpoints,
    my_global_start = mytid*nlocalpoints;
double stepsize = 1./(ntotal_points-1);
array = new double[nlocalpoints];
for (int i=0; i<nlocalpoints; i++)
```

```
array[i] = f( (i+my_global_start)*stepsize );
```

1.10.1.3 Send and receive buffers

The data is specified as a number of elements in a buffer. The same MPI routine can be used with data of different types, so the standard indicates such buffers as *choice*. The specification of this differs per language:

- In C it is an address, so the clean way is to pass it as `(void*)&myvar`.
- Fortran compilers may complain about type mismatches. This can not be helped.

1.10.1.4 Types

`double` vs `MPI_DOUBLE`, Fortran especially.

Addresses have type `MPI_Aint` or `INTEGER (KIND=MPI_ADDRESS_KIND)` in Fortran. The start of the address range is given in `MPI_BOTTOM`.

1.10.2 Blocking communication

This reference section gives the syntax for routines introduced in section 1.3.1.

The basic send command is

```
int MPI_Send(void *buf,  
            int count, MPI_Datatype datatype, int dest, int tag,  
            MPI_Comm comm)
```

http://www.mcs.anl.gov/research/projects/mpi/www/www3/MPI_Send.html This routine may not blocking for small messages; to force blocking behaviour use `MPI_Ssend` with the same argument list. http://www.mcs.anl.gov/research/projects/mpi/www/www3/MPI_Ssend.html

The basic blocking receive command is

```
int MPI_Recv(void *buf,  
            int count, MPI_Datatype datatype, int source, int tag,  
            MPI_Comm comm, MPI_Status *status)
```

http://www.mcs.anl.gov/research/projects/mpi/www/www3/MPI_Recv.html The `count` argument indicates the maximum length of a message; the actual length of the message can be determined with `MPI_Get_count`; see section 1.3.1.3.

The following code is guaranteed to block, since a `MPI_Recv` always blocks:

```
// recvblock.c  
other = 1-mytid;  
MPI_Recv(&recvbuf,1,MPI_INT,other,0,comm,&status);
```

```
MPI_Send(&sendbuf,1,MPI_INT,other,0,comm);
printf("This statement will not be reached on %d\n",mytid);
```

On the other hand, if we put the send call before the receive, code may not block for small messages that fall under the *eager limit*. In this example we send gradually larger messages. From the screen output you can see what the largest message was that fell under the eager limit; after that the code hangs because of a deadlock.

```
// sendblock.c
int other = 1-mytid;
/* loop over increasingly large messages */
for (int size=1; size<2000000000; size*=10) {
    sendbuf = (int*) malloc(size*sizeof(int));
    recvbuf = (int*) malloc(size*sizeof(int));
    if (!sendbuf || !recvbuf) {
        printf("Out of memory\n"); MPI_Abort(comm,1);
    }
    MPI_Send(sendbuf,size,MPI_INT,other,0,comm);
    MPI_Recv(recvbuf,size,MPI_INT,other,0,comm,&status);
    /* If control reaches this point, the send call
       did not block. If the send call blocks,
       we do not reach this point, and the program will hang.
    */
    if (mytid==0)
        printf("Send did not block for size %d\n",size);
    free(sendbuf); free(recvbuf);
}
```

If you want a code to behave the same for all message sizes, you force the send call to be blocking by using `MPI_Ssend`:

```
// ssendblock.c
other = 1-mytid;
sendbuf = (int*) malloc(sizeof(int));
recvbuf = (int*) malloc(sizeof(int));
size = 1;
MPI_Ssend(sendbuf,size,MPI_INT,other,0,comm);
MPI_Recv(recvbuf,size,MPI_INT,other,0,comm,&status);
printf("This statement is not reached\n");
```

1.10.2.1 Receive status

This reference section gives the syntax for routines introduced in section [1.3.1.3](#).

In section 1.3.1.3 we mentioned the master-worker model as one opportunity for inspecting the MPI_SOURCE field of the MPI_Status object. Here is a small example: the tasks perform a variable amount of work (modeled here by a random wait) before sending a message to the master. The master waits for any source, and inspects the status field to report where the message comes from.

```
// anysource.c
if (mytid==ntids-1) {
    int *recv_buffer;
    MPI_Status status;

    recv_buffer = (int*) malloc((ntids-1)*sizeof(int));

    for (int p=0; p<ntids-1; p++) {
        ierr = MPI_Recv(recv_buffer+p, 1, MPI_INT, MPI_ANY_SOURCE, 0, comm,
                        &status); CHK(ierr);
        int sender = status.MPI_SOURCE;
        printf("Message from sender=%d: %d\n",
               sender, recv_buffer[p]);
    }
} else {
    float randomfraction = (rand() / (double)RAND_MAX);
    int randomwait = (int) ( ntids * randomfraction );
    printf("process %d waits for %e/%d=%d\n",
           mytid, randomfraction, ntids, randomwait);
    sleep(randomwait);
    ierr = MPI_Send(&randomwait, 1, MPI_INT, ntids-1, 0, comm); CHK(ierr);
}
```

1.10.3 Deadlock-free blocking messages

This reference section gives the syntax for routines introduced in section 1.3.1.1.

If messages are send roughly in pairs, the MPI_Sendrecv call is an easy way to prevent deadlock. Here you specify both the target of a send and the source of a receive, which can be same in case of a pairwise exchange of data, but they need not be the same.

```
int MPI_Sendrecv(
    void *sendbuf, int sendcount, MPI_Datatype sendtype,
                int dest, int sendtag,
    void *recvbuf, int recvcount, MPI_Datatype recvtype,
                int source, int recvtag,
    MPI_Comm comm, MPI_Status *status)
```

As an example we set up a ring of three processors: each process sends to its right neighbour, and receives from its left neighbour.

```
// sendrecv.c
int
    right = (mytid+1)%3,
    left = (mytid+2)%3;
int my_data, other_data;
MPI_Sendrecv( &my_data, 1, MPI_INTEGER, right, 0,
&other_data, 1, MPI_INTEGER, left, 0,
comm, MPI_STATUS_IGNORE);
```

1.10.4 Non-blocking communication

This reference section gives the syntax for routines introduced in section 1.3.2.

The non-blocking routines have much the same parameter list as the blocking ones, with the addition of an MPI_Request parameter. The MPI_Isend routine does not have a ‘status’ parameter, which has moved to the ‘wait’ routine.

```
int MPI_Isend(void *buf,
    int count, MPI_Datatype datatype, int dest, int tag,
    MPI_Comm comm, MPI_Request *request)
```

http://www.mcs.anl.gov/research/projects/mpi/www/www3/MPI_Isend.html

```
int MPI_Irecv(void *buf,
    int count, MPI_Datatype datatype, int source, int tag,
    MPI_Comm comm, MPI_Request *request)
```

http://www.mcs.anl.gov/research/projects/mpi/www/www3/MPI_Irecv.html

There are various ‘wait’ routines. Since you will often do at least one send and one receive, this routine is useful:

```
int MPI_Waitall(int count, MPI_Request array_of_requests[],
    MPI_Status array_of_statuses[])
```

http://www.mcs.anl.gov/research/projects/mpi/www/www3/MPI_Waitall.html

Here is a simple code that does a non-blocking exchange between two processors:

```
// irecvnonblock.c
MPI_Request request[2]
MPI_Status status[2];
other = 1-mytid;
MPI_Irecv(&recvbuf, 1, MPI_INT, other, 0, comm, &(request[0]));
MPI_Isend(&sendbuf, 1, MPI_INT, other, 0, comm, &(request[1]));
MPI_Waitall(2, request, status);
```

1. MPI

It is possible to omit the status array by specifying `MPI_STATUSES_IGNORE`. Other routines are `MPI_Wait` for a single request, and `MPI_Waitsome`, `MPI_Waitany`.

The above fragment is unrealistically simple. In a more general scenario we have to manage send and receive buffers: we need as many buffers as there are simultaneous non-blocking sends and receives.

```
// irecvloop.c
MPI_Request requests =
    (MPI_Request*) malloc( 2*ntids*sizeof(MPI_Request) );
recv_buffers = (int*) malloc( ntids*sizeof(int) );
send_buffers = (int*) malloc( ntids*sizeof(int) );
for (int p=0; p<ntids; p++) {
    int left_p = (p-1) % ntids,
        right_p = (p+1) % ntids;
    send_buffer[p] = ntids-p;
    MPI_Isend(sendbuffer+p,1,MPI_INT, right_p,0, requests+2*p);
    MPI_Irecv(recvbuffer+p,1,MPI_INT, left_p,0, requests+2*p+1);
}
MPI_Waitall(2*ntids,requests,MPI_STATUSES_IGNORE);
```

Instead of waiting for all messages, we can wait for any message to come with `MPI_Waitany`, and process the receive data as it comes in.

```
// irecv_source.c
if (mytid==ntids-1) {
    int *recv_buffer;
    MPI_Request *request;
    recv_buffer = (int*) malloc((ntids-1)*sizeof(int));
    request = (MPI_Request*) malloc((ntids-1)*sizeof(MPI_Request));

    for (int p=0; p<ntids-1; p++) {
        ierr = MPI_Irecv(recv_buffer+p,1,MPI_INT, p,0,comm,
                         request+p); CHK(ierr);
    }
    for (int p=0; p<ntids-1; p++) {
        int index, sender;
        MPI_Waitany(ntids-1,request,&index,MPI_STATUS_IGNORE);
        printf("Message from %d: %d\n", index, recv_buffer[index]);
    }
}
```

Note the `MPI_STATUS_IGNORE` parameter: we know everything about the incoming message, so we do not need to query a status object. Contrast this with the example in section [1.10.2.1](#).

1.10.5 One-sided communication

This reference section gives the syntax for routines introduced in section [1.3.3](#).

1.10.5.1 Windows and epochs

This reference section gives the syntax for routines introduced in section 1.3.3.1.

```
MPI_WIN_CREATE (void *base, MPI_Aint size,
                 int disp_unit, MPI_Info info,
                 MPI_Comm comm, MPI_Win *win, ierr)
```

The data array must not be PARAMETER of static const.

Here is a single put operation. Note that the window create and window fence calls are collective, so they have to be performed on all processors of the communicator that was used in the create call.

```
// putfence.c
MPI_Win the_window;
MPI_Win_create(&other_number,1,sizeof(int),
               MPI_INFO_NULL,comm,&the_window);
MPI_Win_fence(0,the_window);
if (mytid==0) {
    MPI_Put( /* data on origin: */ &my_number, 1,MPI_INT,
            /* data on target: */ other,0,      1,MPI_INT,
            the_window);
}
MPI_Win_fence(0,the_window);
MPI_Win_free(&the_window);
```

The MPI_Info parameter can be used to pass implementation-dependent information:

```
MPI_Info info;
MPI_Info_create(&info);
MPI_Info_set(info,"no_locks","true");
MPI_Win_create( ... info ... &win);
MPI_Info_free(&info);
```

It is always valid to use MPI_INFO_NULL.

1.10.5.2 More active target synchronization

This reference section gives the syntax for routines introduced in section 1.3.3.5.

You start and complete an exposure epoch with :

```
int MPI_Win_post(MPI_Group group, int assert, MPI_Win win)
int MPI_Win_wait(MPI_Win win)
```

In other words, this turns your window into the *target* for a remote access.

You start and complete an access epoch with :

```
int MPI_Win_start(MPI_Group group, int assert, MPI_Win win)
int MPI_Win_complete(MPI_Win win)
```

In other words, these calls border the access to a remote window, with the current processor being the *origin* of the remote access.

In the following snippet a single processor puts data on one other. Note that they both have their own definition of the group, and that the receiving process only does the post and wait calls.

```
// postwaitwin_two.c
if (mytid==origin) {
    MPI_Group_incl(all_group,1,&target,&two_group);
    // access
    MPI_Win_start(two_group,0,the_window);
    MPI_Put( /* data on origin: */ &my_number, 1,MPI_INT,
             /* data on target: */ target,0, 1,MPI_INT,
             the_window);
    MPI_Win_complete(the_window);
}

if (mytid==target) {
    MPI_Group_incl(all_group,1,&origin,&two_group);
    // exposure
    MPI_Win_post(two_group,0,the_window);
    MPI_Win_wait(the_window);
}
```

1.10.6 Collectives

This reference section gives the syntax for routines introduced in section 1.4.

```
int MPI_Bcast( void *buffer, int count, MPI_Datatype datatype, int root,
               MPI_Comm comm )

int MPI_Reduce(void *sendbuf, void *recvbuf, int count, MPI_Datatype datatype,
               MPI_Op op, int root, MPI_Comm comm)
```

On processes that are not the root, the receive buffer is ignored. On the root, you have two buffers, but by specifying `MPI_IN_PLACE`, the reduction call uses the value in the receive buffer as the root's contribution to the operation. On the Allreduce call, `MPI_IN_PLACE` can be used for the send buffer of every process.

The scan operations are

```
int MPI_Scan(void* sendbuf, void* recvbuf,
             int count, MPI_Datatype datatype, MPI_Op op, MPI_Comm comm )
```

and

```
int MPI_Exscan(void* sendbuf, void* recvbuf,
               int count, MPI_Datatype datatype, MPI_Op op, MPI_Comm comm )
```

The MPI_Op operations do not return an error code.

The result of the exclusive scan is undefined on processor 0, and on processor 1 it is a copy of the send value of processor 1. In particular, the MPI_Op need not be called on these two processors.

1.10.7 Cancelling messages

In section 1.3.1.3 we showed a master-worker example where the master accepts in arbitrary order the messages from the workers. Here we will show a slightly more complicated example, where only the result of the first task to complete is needed. Thus, we issue an MPI_Recv with MPI_ANY_SOURCE as source. When a result comes, we broadcast its source to all processes. All the other workers then use this information to cancel their message with an MPI_Cancel operation.

```
// cancel.c
if (mytid==ntids-1) {
    MPI_Status status;
    ierr = MPI_Recv(dummy, 0, MPI_INT, MPI_ANY_SOURCE, 0, comm,
                    &status); CHK(ierr);
    first_tid = status.MPI_SOURCE;
    ierr = MPI_Bcast(&first_tid, 1, MPI_INT, ntids-1, comm); CHK(ierr);
    printf("first msg came from %d\n", first_tid);
} else {
    float randomfraction = (rand() / (double)RAND_MAX);
    int randomwait = (int) (ntids * randomfraction );
    MPI_Request request;
    printf("process %d waits for %e/%d=%d\n",
           mytid, randomfraction, ntids, randomwait);
    sleep(randomwait);
    ierr = MPI_Isend(dummy, 0, MPI_INT, ntids-1, 0, comm,
                     &request); CHK(ierr);
    ierr = MPI_Bcast(&first_tid, 1, MPI_INT, ntids-1, comm
                     ); CHK(ierr);
    if (mytid!=first_tid) {
        ierr = MPI_Cancel(&request); CHK(ierr);
    }
}
```

1.10.8 Communicators

1.10.8.1 Communicator duplication

This reference section gives the syntax for routines introduced in section 1.6.2.1.

In section 1.8.2 it was explained that MPI messages are non-overtaking. This may lead to confusing situations, witness the following snippet:

```
// commdup_wrong.cxx
class library {
private:
    MPI_Comm comm;
    int mytid, ntid, other;
    MPI_Request *request;
public:
    library(MPI_Comm incomm) {
        comm = incomm;
        MPI_Comm_rank(comm, &mytid);
        other = 1-mytid;
        request = new MPI_Request[2];
    };
    int communication_start();
    int communication_end();
};
int ierr = MPI_Isend(&sdata, 1, MPI_INT, other, 1, comm, &(request[0])); CHK(ierr);
my_library.communication_start();
int ierr = MPI_Irecv(&rdata, 1, MPI_INT, other, MPI_ANY_TAG, comm, &(request[1]));
ierr = MPI_Waitall(2, request, status); CHK(ierr);
my_library.communication_end();
```

This models a main program that does a simple message exchange, and it makes two calls to library routines. Unbeknown to the user, the library also issues send and receive calls, and they turn out to interfere:

```
int library::communication_start() {
    int sdata=6, rdata, ierr;
    ierr = MPI_Isend(&sdata, 1, MPI_INT, other, 2, comm, &(request[0])); CHK(ierr);
    ierr = MPI_Irecv(&rdata, 1, MPI_INT, other, MPI_ANY_TAG, comm, &(request[1])); CHK(ierr);
    return 0;
}

int library::communication_end() {
    MPI_Status status[2];
    int ierr;
    ierr = MPI_Waitall(2, request, status); CHK(ierr);
    return 0;
}
```

Here

- The main program does a send,
- the library call `function_start` does a send and a receive; because the receive can match either send, it is paired with the first one;
- the main program does a receive, which will be paired with the send of the library call;
- both the main program and the library do a wait call, and in both cases all requests are successfully fulfilled, just not the way you intended.

The solution is to give the library a separate communicator with `MPI_Comm_dup`. Newly created communicators should be released again with `MPI_Comm_free`.

```
// commdup_right.cxx
class library {
private:
    MPI_Comm comm;
    int mytid, ntids, other;
    MPI_Request *request;
public:
    library(MPI_Comm incomm) {
        MPI_Comm_dup(incomm, &comm);
        MPI_Comm_rank(comm, &mytid);
        other = 1-mytid;
        request = new MPI_Request[2];
    };
    ~library() {
        MPI_Comm_free(&comm);
    }
    int communication_start();
    int communication_end();
};
```

1.10.8.2 Splitting communicators

This reference section gives the syntax for routines introduced in section 1.6.2.2.

The command `MPI_Comm_split` takes a communicator, and divides it into a number of disjoint communicators. It does this by assigning processes to the same subcommunicator if they have the same user-specified ‘colour’ value.

```
int MPI_Comm_split(MPI_Comm comm, int color, int key,
                    MPI_Comm *newcomm)
```

The ranking of processes in the new communicator is determined by a ‘key’ value. Most of the time, there is no reason to use a relative ranking that is different from the global ranking, so the `MPI_Comm_rank` value of the global communicator is a good choice.

```
// mvp2d.cxx
row_number = ntids % ntids_i;
col_number = ntids / ntids_j;
MPI_Comm_split(global_comm, row_number, mytid, &row_comm);
MPI_Comm_split(global_comm, col_number, mytid, &col_comm);
```

1.10.9 Error handling

This reference section gives the syntax for routines introduced in section 1.8.3.

MPI operators (`MPI_Op`) do not return an error code. In case of an error they call `MPI_Abort`; if `MPI_ERRORS_RETURN` is the error handler, errors may be silently ignored.

1.10.10 More utility stuff

This reference section gives the syntax for routines introduced in section 1.8.5.

MPI has a *wall clock* timer: `MPI_Wtime`

```
double MPI_Wtime(void);
```

which gives the number of seconds from a certain point in the past. Thus, you would write:

```
double tstart,tstop,elapsed;
tstart = MPI_Wtime();
tstop = MPI_Wtime();
elapsed = tstop-tstart;
```

The timer has a resolution of `MPI_Wtick`:

```
double MPI_Wtick(void);
```

Timing in parallel is a tricky issue. For instance, most clusters do not have a central clock, so you can not relate start and stop times on one process to those on another. You can test for a global clock as follows :

```
int *v,flag;
MPI_Attr_get( comm, MPI_WTIME_IS_GLOBAL, &v, &flag );
if (mytid==0) printf(``Time synchronized? %d->%d\n'',flag,*v);
```

1.10.11 Multi-threading

This reference section gives the syntax for routines introduced in section 1.7.

```
int MPI_Init_thread( int *argc, char ***argv, int required, int *provided )
```

- `MPI_THREAD_SINGLE`: each MPI process can only have a single thread.

- MPI_THREAD_FUNNELED: an MPI process can be multithreaded, but all MPI calls need to be done from a single thread.
- MPI_THREAD_SERIALIZED: processes can sustain multiple threads that make MPI calls, but these threads can not be simultaneous: they need to be for instance in an OpenMP *critical section*.
- MPI_THREAD_MULTIPLE: processes can be fully generally multi-threaded.

1.11 Review questions

If you answer that a statement is false, give a one-line explanation.

1. True or false: mpicc is a compiler.
2. True or false: mpirun can only be used for interactive parallel runs.
3. What is the function of a hostfile?
4. True or false: in each communicator, processes are numbered consecutively from zero.
5. If a processor issues a send to another processor, and that other processor issues a receive from the first at *exactly the same time*, the send and receive call return immediately
6. Describe a deadlock scenario involving three processors.
7. True or false: a message sent with MPI_Isend from one processor can be received with an MPI_Recv call on another processor.
8. True or false: a message sent with MPI_Send from one processor can be received with an MPI_Irecv on another processor.
9. Why does the MPI_Irecv call not have an MPI_Status argument?
10. Give a simple model for the time a send operation takes.
11. Give a simple model for the time a broadcast of a single scalar takes.
12. What is the relation between the concepts of ‘origin’, ‘target’ and ‘window’ in one-sided communication.
13. What are the three routines for one-sided data transfer?
14. Give an example of a collective call with and without a root processor.
15. Give two examples of derived datatypes.
16. Give an example where the sender uses a different type to send than the receiver uses in the corresponding receive call. Name the types involved.

1.12 Literature

Online resources:

- MPI 1 Complete reference:
<http://www.netlib.org/utk/papers/mpi-book/mpi-book.html>
- Official MPI documents:
<http://www.mpi-forum.org/docs/>
- List of all MPI routines:
<http://www.mcs.anl.gov/research/projects/mpi/www3/>

1. MPI

Tutorial books on MPI:

- Using MPI [2] by some of the original authors.

Chapter 2

Ruminations on parallelism

The basic problem in designing a parallel algorithm is to find the proper assignment of work to processes. We discuss some possibilities.

2.1 Fully parallel work

Some algorithms involve operations that are truly independent. For instance, if you add two vectors

$$\bar{z} \leftarrow \bar{x} + \bar{y},$$

all the element additions

$$z_i \leftarrow x_i + y_i$$

are independent. Thus, if we had as many processors as vector elements, each could do one addition. We could write this as:

```
parallel_for p in Processors:  
    i = index_of_processor()  
    z[i] = x[i]+y[i]
```

2.2 Introduction: Quicksort

Here is a simple recursive implementation of quicksort.

```
def QuickSort( array )  
    if len(array)==1:  
        return array  
    else:  
        median = find_median( array )  
        left_array = [ x for x in array if x<median ]  
        right_array = [ x for x in array if x>=median ]
```

2. Ruminations on parallelism

```
        return QuickSort( left_array ) + QuickSort( right_array )
def find_median( array ):
    return array[0] # simple solution, probably bad.
```

The important steps are:

- Finding the median. We can do this in a simple way, which is potentially bad, or in more sophisticated ways; however, no general statement about the complexity of this can be made, so we'll just ignore this;
- Splitting the array in two halves;
- Sorting the two subarrays.

There is clear parallelism in this algorithm: the two recursive calls are independent and can therefore be done in parallel.

Let's attempt a complexity analysis of this parallelization scheme. We start with an array of length N , and we assume that we always find the perfect median at zero cost. In that case, the first sort call takes time N to split the array in two halves; the simultaneous recursive calls take time $N/2$ each to split their subarrays, et cetera. We see that the running time is approximately $2N$, which is better than the originally expected $N \log N$, but not a whole lot.

The problem is of course that splitting the array is not done in parallel, although it probably could be. Here is how we could do it with two processors:

```
N = len(array)
new_array = [ None for i in range(N) ]
low_pointer = 0
high_pointer = N
median = find_median( array )

def make_left_array():
    global array,new_array
    for i in range(N):
        if array[i]<median:
            new_array[low_pointer] = array[i]
            low_pointer += 1
    return new_array[:low_pointer]

def make_right_array():
    global array,new_array
    for i in range(N):
        if array[i]>=median:
            high_pointer -= 1
            new_array[high_pointer] = array[i]
    return new_array[high_pointer:]
```

In this code, one process finds the elements below the median, and the other one find elements at least as large as the median. Now we are indeed using two processes, but they each still go through the whole array, so the running time is still N .

(There is a further problem with this algorithm, namely that it uses twice as many memory locations for each level of recursive function calls. This can be eliminated by the ‘Dutch National Flag’ algorithm http://en.wikipedia.org/wiki/Dutch_national_flag_problem, but this needlessly complicates the discussion.)

As a way of halving the amount of work per process, consider a solution where both processes find elements of both type, but they each consider only half the array.

```
N = len(array)
new_array = [ None for i in range(N) ]
low_pointer = 0
high_pointer = N-1
median = find_median( array )

def make_left_array():
    global array,new_array
    for i in range(0:N:2):
        place_element[i]
    return new_array[:low_pointer]

def make_right_array():
    global array,new_array
    for i in range(1:N:2):
        place_element[i]
    return new_array[high_pointer:]

def place_element( i ):
    if array[i]<median:
        new_array[low_pointer] = array[i]
        low_pointer += 1
    elif array[i]>=median:
        high_pointer -= 1
        new_array[high_pointer] = array[i]
```

Now both processes access shared data which has many problems associated with it.

2.3 Data partitioning

In the previous section we looked at the parallelization of quicksort purely by considering a division of the work. We will now look at strategies of parallelizing by looking at the data involved.

2. Ruminations on parallelism

Certain problems can be described as constructing an output data object consisting of many elements, with often independence of the element calculations. Linear algebra is a source of such examples:

- In the matrix-vector product every element of the output vector can be computed individually.
- The same holds for the matrix-matrix product.
- On the other hand, computing and applying an *LU* factorization (see HPSC-[5.3](#)) has a great deal of sequentiality. But even there, sub-operations such as updating a matrix block with a low rank matrix show a great deal of parallelism.

From this description it is clear that one obvious partitioning scheme would be to partition the output over the processes. This is known as the *owner computes* scheme: each output element is uniquely owned by one process, which performs all or most of the computations for constructing that element.

Chapter 3

Hybrid computing

MPI-2 provides precise interaction with multi-threaded programs MPI_THREAD_SINGLE MPI_THREAD_FUNNELLED (OpenMP loops) MPI_THREAD_SERIAL (Open MP single) MPI_THREAD_MULTIPLE

Chapter 4

Support libraries

ParaMesh

Global Arrays

PETSc

Hdf5 and Silo

Appendix A

Practical tutorials

here are some tutorials

A.1 Managing projects with Make

The *Make* utility helps you manage the building of projects: its main task is to facilitate rebuilding only those parts of a multi-file project that need to be recompiled or rebuilt. This can save lots of time, since it can replace a minutes-long full installation by a single file compilation. *Make* can also help maintaining multiple installations of a program on a single machine, for instance compiling a library with more than one compiler, or compiling a program in debug and optimized mode.

Make is a Unix utility with a long history, and traditionally there are variants with slightly different behaviour, for instance on the various flavours of Unix such as HP-UX, AIX, IRIX. These days, it is advisable, no matter the platform, to use the GNU version of *Make* which has some very powerful extensions; it is available on all Unix platforms (on Linux it is the only available variant), and it is a *de facto* standard. The manual is available at <http://www.gnu.org/software/make/manual/make.html>, or you can read the book [3].

There are other build systems, most notably Scons and Bjam. We will not discuss those here. The examples in this tutorial will be for the C and Fortran languages, but *Make* can work with any language, and in fact with things like TeX that are not really a language at all; see section A.1.6.

A.1.1 A simple example

Purpose. In this section you will see a simple example, just to give the flavour of *Make*.

A.1.1.1 C

Make the following files:

foo.c

bar.c

bar.h and a makefile:

Makefile

The makefile has a number of rules like

```
foo.o : foo.c
<TAB>cc -c foo.c
```

which have the general form

```
target : prerequisite(s)
<TAB>rule(s)
```

where the rule lines are indented by a TAB character.

A rule, such as above, states that a ‘target’ file `foo.o` is made from a ‘prerequisite’ `foo.c`, namely by executing the command `cc -c foo.c`. The precise definition of the rule is:

- if the target `foo.o` does not exist or is older than the prerequisite `foo.c`,
- then the command part of the rule is executed: `cc -c foo.c`
- If the prerequisite is itself the target of another rule, than that rule is executed first.

Probably the best way to interpret a rule is:

- if any prerequisite has changed,
- then the target needs to be remade,
- and that is done by executing the commands of the rule.

If you call `make` without any arguments, the first rule in the makefile is evaluated. You can execute other rules by explicitly invoking them, for instance `make foo.o` to compile a single file.

Exercise. Call `make`.

Expected outcome. The above rules are applied: `make` without arguments tries to build the first target, `fooprog`. In order to build this, it needs the prerequisites `foo.o` and `bar.o`, which do not exist. However, there are rules for making them, which `make` recursively invokes. Hence you see two compilations, for `foo.o` and `bar.o`, and a link command for `fooprog`.

Caveats. Typos in the makefile or in file names can cause various errors. In particular, make sure you use tabs and not spaces for the rule lines. Unfortunately, debugging a makefile is not simple. *Make*’s error message will usually give you the line number in the make file where the error was detected.

Exercise. Do `make clean`, followed by `mv foo.c boo.c` and `make` again. Explain the error message. Restore the original file name.

Expected outcome. `Make` will complain that there is no rule to make `foo.c`. This error was caused when `foo.c` was a prerequisite, and was found not to exist. `Make` then went looking for a rule to make it.

Now add a second argument to the function `bar`. This requires you to edit `bar.c` and `bar.h`: go ahead and make these edits. However, it also requires you to edit `foo.c`, but let us for now ‘forget’ to do that. We will see how `Make` can help you find the resulting error.

Exercise. Call `make` to recompile your program. Did it recompile `foo.c`?

Expected outcome. Even though conceptually `foo.c` would need to be recompiled since it uses the `bar` function, `Make` did not do so because the makefile had no rule that forced it.

In the makefile, change the line

```
foo.o : foo.c
```

to

```
foo.o : foo.c bar.h
```

which adds `bar.h` as a prerequisite for `foo.o`. This means that, in this case where `foo.o` already exists, `Make` will check that `foo.o` is not older than any of its prerequisites. Since `bar.h` has been edited, it is younger than `foo.o`, so `foo.o` needs to be reconstructed.

Exercise. Confirm that the new makefile indeed causes `foo.o` to be recompiled if `bar.h` is changed. This compilation will now give an error, since you ‘forgot’ to edit the use of the `bar` function.

A.1.1.2 Fortran

Make the following files:

`foomain.F`

`foomod.F` and a makefile:

`Makefile` If you call `make`, the first rule in the makefile is executed. Do this, and explain what happens.

Exercise. Call `make`.

Expected outcome. The above rules are applied: `make` without arguments tries to build the first target, `foomain`. In order to build this, it needs the prerequisites `foomain.o` and `foomod.o`, which do not exist. However, there are rules for making them, which `make` recursively invokes. Hence you see two compilations, for `foomain.o` and `foomod.o`, and a link command for `fooprog`.

Caveats. Typos in the makefile or in file names can cause various errors. Unfortunately, debugging a makefile is not simple. You will just have to understand the errors, and make the corrections.

Exercise. Do `make clean`, followed by `mv foo.c boo.c` and `make` again. Explain the error message. Restore the original file name.

Expected outcome. `Make` will complain that there is no rule to make `foo.c`. This error was caused when `foo.c` was a prerequisite, and was found not to exist. `Make` then went looking for a rule to make it.

Now add an extra parameter to `func` in `foomod.F` and recompile.

Exercise. Call `make` to recompile your program. Did it recompile `foomain.F`?

Expected outcome. Even though conceptually `foomain.F` would need to be recompiled, `Make` did not do so because the makefile had no rule that forced it.

Change the line

```
foomain.o : foomain.F
```

to

```
foomain.o : foomain.F foomod.F
```

which adds `foomod.F` as a prerequisite for `foomain.o`. This means that, in this case where `foomain.o` already exists, *Make* will check that `foomain.o` is not older than any of its prerequisites. Since `foomod.F` has been edited, it is younger than `foomain.o`, so `foomain.o` needs to be reconstructed.

Exercise. Confirm that the corrected makefile indeed causes `foomain.F` to be recompiled.

A.1.2 Variables and template rules

Purpose. In this section you will learn various work-saving mechanism in *Make*, such as the use of variables, and of template rules.

A.1.2.1 Makefile variables

It is convenient to introduce variables in your makefile. For instance, instead of spelling out the compiler explicitly every time, introduce a variable in the makefile:

```
CC = gcc  
FC = gfortran
```

and use `$(CC)` or `$(FC)` on the compile lines:

```
foo.o : foo.c  
        $(CC) -c foo.c  
foomain.o : foomain.F  
        $(FC) -c foomain.F
```

Exercise. Edit your makefile as indicated. First do `make clean`, then `make foo` (C) or `make fooprog` (Fortran).

Expected outcome. You should see the exact same compile and link lines as before.

Caveats. Unlike in the shell, where braces are optional, variable names in a makefile have to be in braces or parentheses. Experiment with what happens if you forget the braces around a variable name.

One advantage of using variables is that you can now change the compiler from the commandline:

```
make CC="icc -O2"  
make FC="gfortran -g"
```

Exercise. Invoke *Make* as suggested (after `make clean`). Do you see the difference in your screen output?

Expected outcome. The compile lines now show the added compiler option `-O2` or `-g`.

Make also has built-in variables:

- `$@` The target. Use this in the link line for the main program.
- `$^` The list of prerequisites. Use this also in the link line for the program.
- `$<` The first prerequisite. Use this in the compile commands for the individual object files.

Using these variables, the rule for `fooprog` becomes

```
fooprog : foo.o bar.o  
        ${CC} -o $@ $^
```

and a typical compile line becomes

```
foo.o : foo.c bar.h  
        ${CC} -c $<
```

You can also declare a variable

```
THEPROGRAM = fooprog
```

and use this variable instead of the program name in your makefile. This makes it easier to change your mind about the name of the executable later.

Exercise. Construct a commandline so that your makefile will build the executable `fooprog_v2`.

Expected outcome. You need to specify the `THEPROGRAM` variable on the commandline using the syntax `make VAR=value`.

Caveats. Make sure that there are no spaces around the equals sign in your commandline.

A.1.2.2 Template rules

In your makefile, the rules for the object files are practically identical:

- the rule header (`foo.o : foo.c`) states that a source file is a prerequisite for the object file with the same base name;
- and the instructions for compiling (`${CC} -c $<`) are even character-for-character the same, now that you are using *Make*'s built-in variables;
- the only rule with a difference is

```
foo.o : foo.c bar.h  
        ${CC} -c $<
```

where the object file depends on the source file and another file.

We can take the commonalities and summarize them in one rule¹:

```
% .o : %.c
      ${CC} -c $<
% .o : %.F
      ${FC} -c $<
```

This states that any object file depends on the C or Fortran file with the same base name. To regenerate the object file, invoke the C or Fortran compiler with the `-c` flag. These template rules can function as a replacement for the multiple specific targets in the makefiles above, except for the rule for `foo.o`.

The dependence of `foo.o` on `bar.h` can be handled by adding a rule

```
foo.o : bar.h
```

with no further instructions. This rule states, ‘if the prerequisite file `bar.h` changed, file `foo.o` needs updating’. Make will then search the makefile for a different rule that states how this updating is done.

Exercise. Change your makefile to incorporate these ideas, and test.

A.1.3 Wildcards

Your makefile now uses one general rule for compiling all your source files. Often, these source files will be all the `.c` or `.F` files in your directory, so is there a way to state ‘compile everything in this directory’? Indeed there is. Add the following lines to your makefile, and use the variable `COBJECTS` or `FOBJECTS` wherever appropriate.

```
# wildcard: find all files that match a pattern
CSOURCES := ${wildcard *.c}
# pattern substitution: replace one pattern string by another
COBJECTS := ${patsubst %.c,%.o,${SRC} }

FSOURCES := ${wildcard *.F}
FOBJECTS := ${patsubst %.F,%.o,${SRC} }
```

A.1.4 Miscellania

A.1.4.1 What does this makefile do?

Above you learned that issuing the `make` command will automatically execute the first rule in the makefile. This is convenient in one sense², and inconvenient in another: the only way to find out what possible actions a makefile allows is to read the makefile itself, or the – usually insufficient – documentation.

-
1. This mechanism is the first instance you’ll see that only exists in GNU make, though in this particular case there is a similar mechanism in standard make. That will not be the case for the wildcard mechanism in the next section.
 2. There is a convention among software developers that a package can be installed by the sequence `./configure ; make ; make install`, meaning: Configure the build process for this computer, Do the actual build, Copy files to some system directory such as `/usr/bin`.

A better idea is to start the makefile with a target

```
info :  
    @echo "The following are possible:"  
    @echo "  make"  
    @echo "  make clean"
```

Now make without explicit targets informs you of the capabilities of the makefile. The at-sign at the start of the commandline means ‘do not echo this command to the terminal’, which makes for cleaner terminal output; remove the at signs and observe the difference in behaviour.

A.1.4.2 Phony targets

The example makefile contained a target `clean`. This uses the *Make* mechanisms to accomplish some actions that are not related to file creation: calling `make clean` causes *Make* to reason ‘there is no file called `clean`, so the following instructions need to be performed’. However, this does not actually cause a file `clean` to spring into being, so calling `make clean` again will make the same instructions being executed.

To indicate that this rule does not actually make the target, declare

```
.PHONY : clean
```

One benefit of declaring a target to be phony, is that the *Make* rule will still work, even if you have a file named `clean`.

A.1.4.3 Predefined variables and rules

Calling `make -p yourtarget` causes `make` to print out all its actions, as well as the values of all variables and rules, both in your makefile and ones that are predefined. If you do this in a directory where there is no makefile, you’ll see that `make` actually already knows how to compile `.c` or `.F` files. Find this rule and find the definition of the variables in it.

You see that you can customize `make` by setting such variables as `CFLAGS` or `FFLAGS`. Confirm this with some experimentation. If you want to make a second makefile for the same sources, you can call `make -f othermakefile` to use this instead of the default *Makefile*.

A.1.4.4 Using the target as prerequisite

Suppose you have two different targets that are treated largely the same. You would want to write:

```
PROGS = myfoo other  
${PROGS} : ${@:o}  
    ${CC} -o $@ ${@:o} ${list of libraries goes here}
```

and saying `make myfoo` would cause

```
cc -c myfoo.c
cc -o myfoo myfoo.o ${list of libraries}
```

and likewise for `make other`. What goes wrong here is the use of `$@.o` as prerequisite. In Gnu Make, you can repair this as follows:

```
.SECONDEXPANSION:
${PROGS} : $$@.o
```

A.1.5 Shell scripting in a Makefile

Purpose. In this section you will see an example of a longer shell script appearing in a makefile rule.

In the makefiles you have seen so far, the command part was a single line. You can actually have as many lines there as you want. For example, let us make a rule for making backups of the program you are building.

Add a backup rule to your makefile. The first thing it needs to do is make a backup directory:

```
.PHONY : backup
backup :
    if [ ! -d backup ] ; then
        mkdir backup
    fi
```

Did you type this? Unfortunately it does not work: every line in the command part of a makefile rule gets executed as a single program. Therefore, you need to write the whole command on one line:

```
backup :
    if [ ! -d backup ] ; then mkdir backup ; fi
```

or if the line gets too long:

```
backup :
    if [ ! -d backup ] ; then \
        mkdir backup ; \
    fi
```

Next we do the actual copy:

```
backup :
    if [ ! -d backup ] ; then mkdir backup ; fi
    cp myprog backup/myprog
```

But this backup scheme only saves one version. Let us make a version that has the date in the name of the saved program.

The Unix `date` command can customize its output by accepting a format string. Type the following:
`date` This can be used in the makefile.

Exercise. Edit the `cp` command line so that the name of the backup file includes the current date.

Expected outcome. Hint: you need the backquote. Consult the Unix tutorial if you do not remember what backquotes do.

If you are defining shell variables in the command section of a makefile rule, you need to be aware of the following. Extend your `backup` rule with a loop to copy the object files:

```
backup :  
    if [ ! -d backup ] ; then mkdir backup ; fi  
    cp myprog backup/myprog  
    for f in ${OBJS} ; do \  
        cp $f backup ; \  
    done
```

(This is not the best way to copy, but we use it for the purpose of demonstration.) This leads to an error message, caused by the fact that *Make* interprets `$f` as an environment variable of the outer process. What works is:

```
backup :  
    if [ ! -d backup ] ; then mkdir backup ; fi  
    cp myprog backup/myprog  
    for f in ${OBJS} ; do \  
        cp $$f backup ; \  
    done
```

(In this case *Make* replaces the double dollar by a single one when it scans the commandline. During the execution of the commandline, `$f` then expands to the proper filename.)

A.1.6 A Makefile for **LATEX**

A.2 Debugging

When a program misbehaves, *debugging* is the process of finding out *why*. There are various strategies of finding errors in a program. The crudest one is debugging by print statements. If you have a notion of where in your code the error arises, you can edit your code to insert print statements, recompile, rerun, and see if the output gives you any suggestions. There are several problems with this:

- The edit/compile/run cycle is time consuming, especially since
- often the error will be caused by an earlier section of code, requiring you to edit, compile, and rerun repeatedly. Furthermore,
- the amount of data produced by your program can be too large to display and inspect effectively, and
- if your program is parallel, you probably need to print out data from all processors, making the inspection process very tedious.

For these reasons, the best way to debug is by the use of an interactive *debugger*, a program that allows you to monitor and control the behaviour of a running program. In this section you will familiarize yourself with *gdb*, which is the open source debugger of the *GNU* project. Other debuggers are proprietary, and typically come with a compiler suite. Another distinction is that *gdb* is a commandline debugger; there are graphical debuggers such as *ddd* (a frontend to *gdb*) or *DDT* and *TotalView* (debuggers for parallel codes). We limit ourselves to *gdb*, since it incorporates the basic concepts common to all debuggers.

In this tutorial you will debug a number of simple programs with *gdb* and *valgrind*. The files can be downloaded from <http://tinyurl.com/ISTC-debug-tutorial>.

A.2.1 Invoking *gdb*

There are three ways of using *gdb*: using it to start a program, attaching it to an already running program, or using it to inspect a *core dump*. We will only consider the first possibility.

Here is an example of how to start *gdb* with program that has no arguments (Fortran users, use *hello.F*):

```
tutorials/gdb/c/hello.c
%% cc -g -o hello hello.c
# regular invocation:
%% ./hello
hello world
# invocation from gdb:
%% gdb hello
GNU gdb 6.3.50-20050815 # .... version info
Copyright 2004 Free Software Foundation, Inc. .... copyright info ....
(gdb) run
Starting program: /home/eijkhout/tutorials/gdb/hello
Reading symbols for shared libraries +. done
hello world

Program exited normally.
```

A. Practical tutorials

```
(gdb) quit  
%%
```

Important note: the program was compiled with the *debug flag* `-g`. This causes the *symbol table* (that is, the translation from machine address to program variables) and other debug information to be included in the binary. This will make your binary larger than strictly necessary, but it will also make it slower, for instance because the compiler will not perform certain optimizations³.

To illustrate the presence of the symbol table do

```
%% cc -g -o hello hello.c  
%% gdb hello  
GNU gdb 6.3.50-20050815 # ..... version info  
(gdb) list
```

and compare it with leaving out the `-g` flag:

```
%% cc -o hello hello.c  
%% gdb hello  
GNU gdb 6.3.50-20050815 # ..... version info  
(gdb) list
```

For a program with commandline input we give the arguments to the `run` command (Fortran users use `say.F`):

tutorials/gdb/c/say.c

```
%% cc -o say -g say.c  
%% ./say 2  
hello world  
hello world  
%% gdb say  
.... the usual messages ...  
(gdb) run 2  
Starting program: /home/eijkhout/tutorials/gdb/c/say 2  
Reading symbols for shared libraries +. done  
hello world  
hello world  
  
Program exited normally.
```

3. Compiler optimizations are not supposed to change the semantics of a program, but sometimes do. This can lead to the nightmare scenario where a program crashes or gives incorrect results, but magically works correctly with compiled with debug and run in a debugger.

A.2.2 Finding errors

Let us now consider some programs with errors.

A.2.2.1 C programs

```
tutorials/gdb/c/square.c
```

```
%% cc -g -o square square.c
%% ./square
5000
Segmentation fault
```

The *segmentation fault* (other messages are possible too) indicates that we are accessing memory that we are not allowed to, making the program abort. A debugger will quickly tell us where this happens:

```
%% gdb square
(gdb) run
50000

Program received signal EXC_BAD_ACCESS, Could not access memory.
Reason: KERN_INVALID_ADDRESS at address: 0x000000000000eb4a
0x00007fff824295ca in __svfscanf_l ()
```

Apparently the error occurred in a function `__svfscanf_l`, which is not one of ours, but a system function. Using the `backtrace` (or `bt`, also `where` or `w`) command we quickly find out how this came to be called:

```
(gdb) backtrace
#0 0x00007fff824295ca in __svfscanf_l ()
#1 0x00007fff8244011b in fscanf ()
#2 0x0000000100000e89 in main (argc=1, argv=0x7fff5fbfc7c0) at square.c:7
```

We take a close look at line 7, and see that we need to change `nmax` to `&nmax`.

There is still an error in our program:

```
(gdb) run
50000

Program received signal EXC_BAD_ACCESS, Could not access memory.
Reason: KERN_PROTECTION_FAILURE at address: 0x000000010000f000
0x0000000100000ebe in main (argc=2, argv=0x7fff5fbfc7a8) at square1.c:9
9         squares[i] = 1./(i*i); sum += squares[i];
```

We investigate further:

A. Practical tutorials

```
(gdb) print i
$1 = 11237
(gdb) print squares[i]
Cannot access memory at address 0x10000f000
```

and we quickly see that we forgot to allocate `squares`.

By the way, we were lucky here: this sort of memory errors is not always detected. Starting our programm with a smaller input does not lead to an error:

```
(gdb) run
50
Sum: 1.625133e+00

Program exited normally.
```

A.2.2.2 Fortran programs

Compile and run the following program:

`tutorials/gdb/f/square.F` It should abort with a message such as ‘Illegal instruction’. Running the program in `gdb` quickly tells you where the problem lies:

```
(gdb) run
Starting program: tutorials/gdb//fsquare
Reading symbols for shared libraries +++. done

Program received signal EXC_BAD_INSTRUCTION, Illegal instruction/operand.
0x0000000100000da3 in square () at square.F:7
7           sum = sum + squares(i)
```

We take a close look at the code and see that we did not allocate `squares` properly.

A.2.3 Memory debugging with Valgrind

Insert the following allocation of `squares` in your program:

```
squares = (float *) malloc( nmax*sizeof(float) );
```

Compile and run your program. The output will likely be correct, although the program is not. Can you see the problem?

To find such subtle memory errors you need a different tool: a memory debugging tool. A popular (because open source) one is *valgrind*; a common commercial tool is *purify*.

`tutorials/gdb/c/square1.c` Compile this program with `cc -o square1 square1.c` and run it with `valgrind square1` (you need to type the input value). You will lots of output, starting with:

```
%% valgrind square1
==53695== Memcheck, a memory error detector
==53695== Copyright (C) 2002-2010, and GNU GPL'd, by Julian Seward et al.
==53695== Using Valgrind-3.6.1 and LibVEX; rerun with -h for copyright info
==53695== Command: a.out
==53695==
10
==53695== Invalid write of size 4
==53695==   at 0x100000EB0: main (square1.c:10)
==53695==     Address 0x10027e148 is 0 bytes after a block of size 40 alloc'd
==53695==       at 0x1000101EF: malloc (vg_replace_malloc.c:236)
==53695==     by 0x100000E77: main (square1.c:8)
==53695==
==53695== Invalid read of size 4
==53695==   at 0x100000EC1: main (square1.c:11)
==53695==     Address 0x10027e148 is 0 bytes after a block of size 40 alloc'd
==53695==       at 0x1000101EF: malloc (vg_replace_malloc.c:236)
==53695==     by 0x100000E77: main (square1.c:8)
```

Valgrind is informative but cryptic, since it works on the bare memory, not on variables. Thus, these error messages take some exegesis. They state that a line 10 writes a 4-byte object immediately after a block of 40 bytes that was allocated. In other words: the code is writing outside the bounds of an allocated array. Do you see what the problem in the code is?

Note that valgrind also reports at the end of the program run how much memory is still in use, meaning not properly freed.

If you fix the array bounds and recompile and rerun the program, valgrind still complains:

```
==53785== Conditional jump or move depends on uninitialised value(s)
==53785==   at 0x10006FC68: __ dtoa (in /usr/lib/libSystem.B.dylib)
==53785==     by 0x10003199F: __ vfprintf (in /usr/lib/libSystem.B.dylib)
==53785==       by 0x1000738AA: vfprintf_l (in /usr/lib/libSystem.B.dylib)
==53785==         by 0x1000A1006: printf (in /usr/lib/libSystem.B.dylib)
==53785==           by 0x100000EF3: main (in ./square2)
```

Although no line number is given, the mention of `printf` gives an indication where the problem lies. The reference to an ‘uninitialized value’ is again cryptic: the only value being output is `sum`, and that is not uninitialized: it has been added to several times. Do you see why valgrind calls it uninitialized all the same?

A.2.4 Stepping through a program

Often the error in a program is sufficiently obscure that you need to investigate the program run in detail. Compile the following program

tutorials/gdb/c/roots.c and run it:

```
%% ./roots
```

A. Practical tutorials

```
sum: nan
```

Start it in gdb as follows:

```
%% gdb roots
GNU gdb 6.3.50-20050815 (Apple version gdb-1469) (Wed May 5 04:36:56 UTC 2005)
Copyright 2004 Free Software Foundation, Inc.

...
(gdb) break main
Breakpoint 1 at 0x100000ea6: file root.c, line 14.
(gdb) run
Starting program: tutorials/gdb/c/roots
Reading symbols for shared libraries +. done

Breakpoint 1, main () at roots.c:14
14      float x=0;
```

Here you have done the following:

- Before calling `run` you set a *breakpoint* at the main program, meaning that the execution will stop when it reaches the main program.
- You then call `run` and the program execution starts;
- The execution stops at the first instruction in main.

If execution is stopped at a breakpoint, you can do various things, such as issuing the `step` command:

```
Breakpoint 1, main () at roots.c:14
14      float x=0;
(gdb) step
15      for (i=100; i>-100; i--)
(gdb)
16      x += root(i);
(gdb)
```

(if you just hit return, the previously issued command is repeated). Do a number of `steps` in a row by hitting return. What do you notice about the function and the loop?

Switch from doing `step` to doing `next`. Now what do you notice about the loop and the function?

Set another breakpoint: `break 17` and do `cont`. What happens?

Rerun the program after you set a breakpoint on the line with the `sqrt` call. When the execution stops there do `where` and `list`.

- If you set many breakpoints, you can find out what they are with `info breakpoints`.
- You can remove breakpoints with `delete n` where `n` is the number of the breakpoint.
- If you restart your program with `run` without leaving `gdb`, the breakpoints stay in effect.
- If you leave `gdb`, the breakpoints are cleared but you can save them: `save breakpoints <file>`. Use `source <file>` to read them in on the next `gdb` run.

A.2.5 Inspecting values

Run the previous program again in gdb: set a breakpoint at the line that does the `sqrt` call before you actually call `run`. When the program gets to line 8 you can do `print n`. Do `cont`. Where does the program stop?

If you want to repair a variable, you can do `set var=value`. Change the variable `n` and confirm that the square root of the new value is computed. Which commands do you do?

If a problem occurs in a loop, it can be tedious keep typing `cont` and inspecting the variable with `print`. Instead you can add a condition to an existing breakpoint: the following:

```
condition 1 if (n<0)
```

or set the condition when you define the breakpoint:

```
break 8 if (n<0)
```

Another possibility is to use `ignore 1 50`, which will not stop at breakpoint 1 the next 50 times.

Remove the existing breakpoint, redefine it with the condition `n<0` and rerun your program. When the program breaks, find for what value of the loop variable it happened. What is the sequence of commands you use?

A.2.6 Further reading

A good tutorial: <http://www.dirac.org/linux/gdb/>.

Reference manual: http://www.ofb.net-gnu/gdb/gdb_toc.html.

A.3 Tracing

A.3.1 TAU profiling and tracing

TAU <http://www.cs.uoregon.edu/Research/tau/home.php> is a utility for profiling and tracing your parallel programs. Profiling is the gathering and displaying of bulk statistics, for instance showing you which routines take the most time, or whether communication takes a large portion of your runtime. When you get concerned about performance, a good profiling tool is indispensable.

Tracing is the construction and displaying of time-dependent information on your program run, for instance showing you if one process lags behind others. For understanding a program's behaviour, and the reasons behind profiling statistics, a tracing tool can be very insightful.

TAU works by adding *instrumentation* to your code: in effect it is a source-to-source translator that takes your code and turns it into one that generates run-time statistics. Doing this instrumentation is fortunately simple: start by having this code fragment in your makefile:

```
ifdef TACC_TAU_DIR
    CC = tau_cc.sh
else
    CC = mpicc
endif

% : %.c
${CC} -o $@ $^
```

To use TAU, do `module load tau`. You have to set the environment variable `TAU_TRACE` to 1; it's advisable to set `TRACEDIR` to some directory for all the TAU output. Likewise set `TAU_PROFILE` to 1 and set `PROFILEDIR`.

Bibliography

- [1] Ernie Chan, Marcel Heimlich, Avi Purkayastha, and Robert van de Geijn. Collective communication: theory, practice, and experience. *Concurrency and Computation: Practice and Experience*, 19:1749–1783, 2007.
- [2] W. Gropp, E. Lusk, and A. Skjellum. *Using MPI*. The MIT Press, 1994.
- [3] Robert Mecklenburg. *Managing Projects with GNU Make*. O'Reilly Media, 3rd edition edition, 2004. Print ISBN:978-0-596-00610-5 ISBN 10:0-596-00610-1 Ebook ISBN:978-0-596-10445-0 ISBN 10:0-596-10445-6.

Appendix B

Index and list of acronyms

AVX Advanced Vector Extensions

BSP Bulk Synchronous Parallel

CAF Co-array Fortran

DAG Directed Acyclic Graph

DSP Digital Signal Processing

FPU Floating Point Unit

FFT Fast Fourier Transform

FSA Finite State Automaton

HPC High-Performance Computing

HPF High Performance Fortran

MIC Many Integrated Cores

MIMD Multiple Instruction Multiple Data

MPI Message Passing Interface

MTA Multi-Threaded Architecture

NUMA Non-Uniform Memory Access

PGAS Partitioned Global Address Space

PDE Partial Differential Equation

PRAM Parallel Random Access Machine

RDMA Remote Direct Memory Access

RMA Remote Memory Access

SAN Storage Area Network

SaaS Software as-a Service

SFC Space-Filling Curve

SIMD Single Instruction Multiple Data

SIMT Single Instruction Multiple Thread

SM Streaming Multiprocessor

SMP Symmetric Multi Processing

SOR Successive Over-Relaxation

SP Streaming Processor

SPMD Single Program Multiple Data

SPD symmetric positive definite

SSE SIMD Streaming Extensions

TLB Translation Look-aside Buffer

UMA Uniform Memory Access

UPC Unified Parallel C

WAN Wide Area Network

Index

active target synchronization, 18, 19, 21
argc, 50
argv, 50
atomic operations, 22

batch
 job, 7
 scheduler, 7
Beowulf cluster, 6
breakpoint, 86
buffers, 11

C++, 37
C99, 29
choice, 52
collective
 root of the, 23
collectives, 23–28
communication
 blocking, 10–14
 non-blocking, 15–17
 one-sided, 17–23
 overlap with computation, 17
 two-sided, 10–17
communicator, 9, 29–32
core dump, 81
Cray
 T3E, 37
critical section, 63

ddd, 81
DDT, 81
ddt, 37
deadlock, 11, 14
debug flag, 82
debugger, 81
debugging, 81–87

dense linear algebra, 31
distributed shared memory, 18

eager limit, 53
emacs, 39
epoch, 19
 access, 19, 22, 57
 exposure, 19, 21, 57
ethernet, 8

fence, 19
Fortran, 9

gdb, 81–87
ghost region, 48
GNU, 81
 gdb, see gdb
group of
 processors, 22

halo, 48
handshake, 14

ibrun, 7
instrumentation, 88

Make, 72–80
Mandelbrot set, 43
master-worker, 43
master-worker model, 14, 17
MPI
 2.2, 37
 I/O, 35
 semantics, 34
 MPI_Abort, 9, 50
 MPI_Aint, 52
 MPI_ANY_SOURCE, 14, 34, 59
 MPI_BOTTOM, 52

INDEX

MPI_Cancel, 59
MPI_Comm_create, 31
MPI_Comm_dup, 30, 61
MPI_Comm_free, 30, 61
MPI_Comm_group, 31
MPI_COMM_NULL, 30
MPI_Comm_rank, 9
MPI_COMM_SELF, 30
MPI_Comm_set_errhandler, 34
MPI_Comm_set_name, 31
MPI_Comm_size, 9
MPI_Comm_split, 30, 61
MPI_COMM_WORLD, 9, 30
MPI_ERROR, 35
MPI_Error_string, 35
MPI_ERRORS_ARE_FATAL, 34
MPI_ERRORS_RETURN, 34, 62
MPI_Finalize, 50
MPI_Get_count, 33, 34
MPI_Group_difference, 31
MPI_Group_excl, 31
MPI_Group_incl, 31
MPI_IN_PLACE, 24, 58
MPI_Info, 57
MPI_INFO_NULL, 19, 57
MPI_Init, 50
MPI_Init_thread, 32
MPI_Iprobe, 34
MPI_Irecv, 15
MPI_Isend, 15
MPI_LOCK_EXCLUSIVE, 22
MPI_LOCK_SHARED, 22
MPI_MAX, 24
MPI_MODE_NOPRECEDE, 20
MPI_MODE_NOPUT, 20
MPI_MODE_NOSTORE, 20
MPI_MODE_NOSUCCEED, 20
MPI_Op, 62
MPI_Op_create, 25
MPI_Probe, 34
MPI_PROC_NULL, 13, 43
MPI_PROD, 24
MPI_REPLACE, 21
MPI_Rsend, 14
MPI_Scan, 25
MPI_Sendrecv, 13, 43, 54
MPI_SOURCE, 14, 54
MPI_Ssend, 14
MPI_Status, 33, 54
MPI_STATUS_IGNORE, 33, 56
MPI_STATUSES_IGNORE, 33
MPI_SUM, 24
MPI_THREAD_FUNNELED, 63
MPI_THREAD_MULTIPLE, 63
MPI_THREAD_SERIALIZED, 63
MPI_THREAD_SINGLE, 62
MPI_Wait, 15
MPI_Wait..., 15
MPI_Waitall, 15
MPI_Waitany, 15, 33, 56
MPI_Waitsome, 16
MPI_Win_complete, 57
MPI_Win_fence, 19, 43
MPI_Win_post, 57
MPI_Win_start, 57
MPI_Win_wait, 57
MPI_Wtick, 62
MPI_Wtime, 36, 62
MPI_WTIME_IS_GLOBAL, 62
mpirun, 7, 8, 30
origin, 17, 21, 58
owner computes, 68
passive target synchronization, 18, 22
PMPI_..., 36
purify, 84
RMA
 active, 18
 passive, 18
segmentation fault, 83
segmented scan, 25
serialization
 unexpected, 12
shmem, 37
ssh, 7

symbol table, 82

target, 18, 21, 57

active synchronization, see active target synchronization

passive synchronization, see passive target synchronization

TAU, 88

thread-safe, 32

TotalView, 37, 81

valgrind, 37, 84–85

vi, 39

virtual shared memory, 18

wall clock, 62

wall clock time, 36

window, 18–20