

Collective Communications inside MPI

Jerome Vienne

Outline

- ❑ What is a Collective ?
- ❑ Behind Gather
- ❑ Overview of algorithms
 - Binomial Tree
 - Barrier Algorithms
 - Allgather
 - ...
- ❑ Current state of research

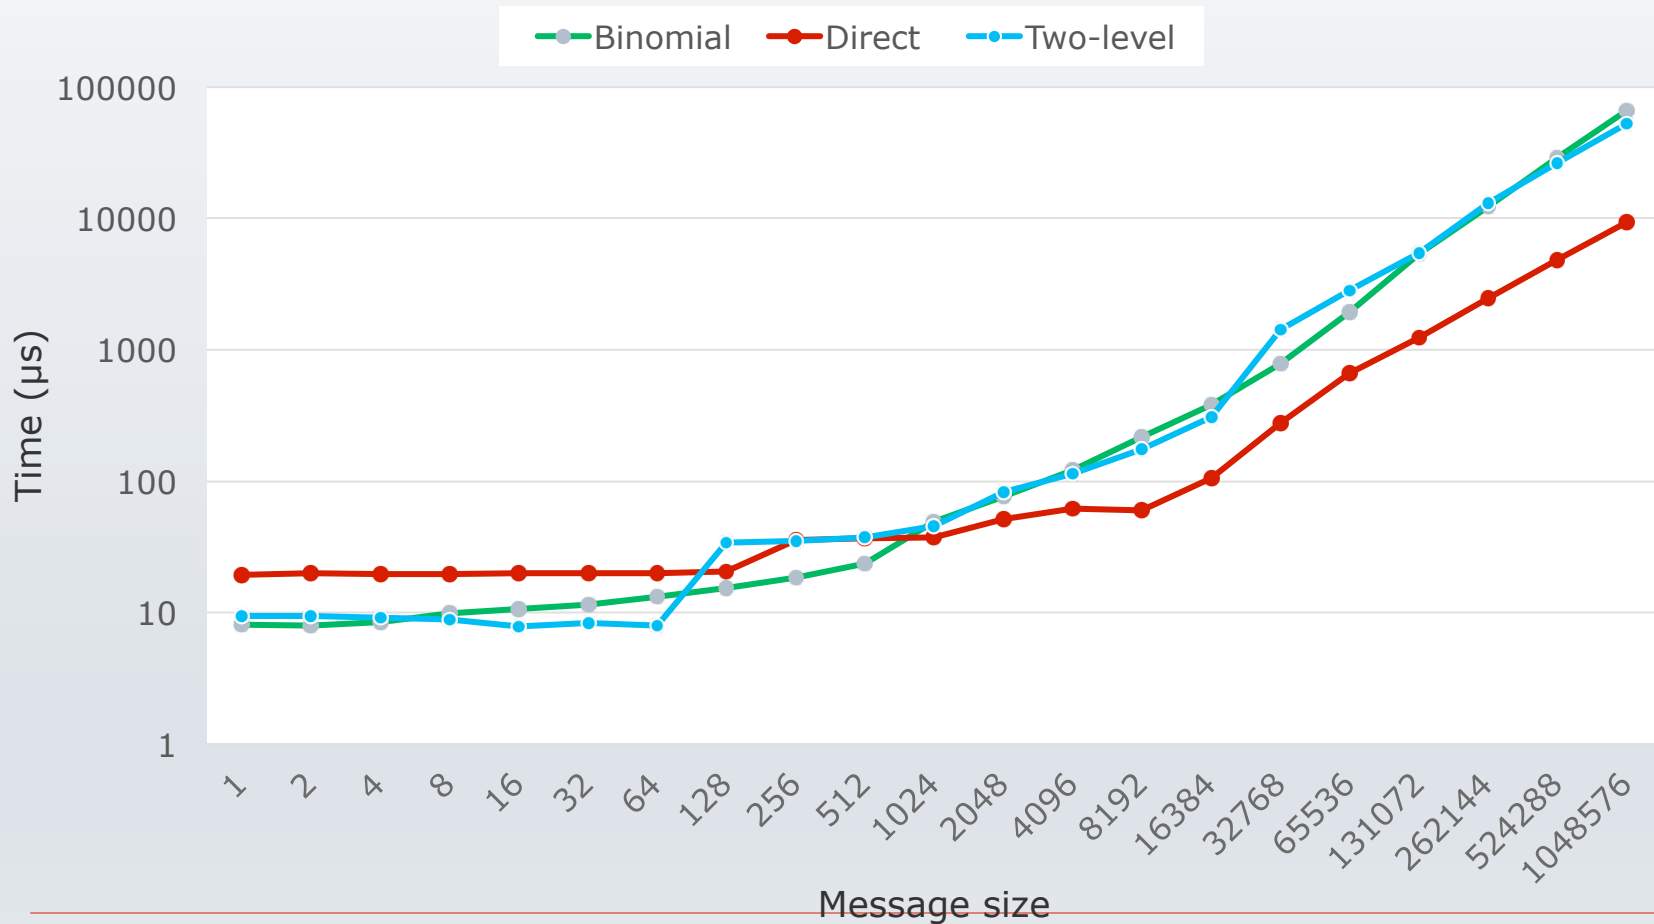
What is a Collective ?

- ❑ Collective communication is a method of communication which involves participation of **all** processes in a communicator.
- ❑ MPI-1: Blocking Collectives
- ❑ MPI-3: Non-blocking Collectives

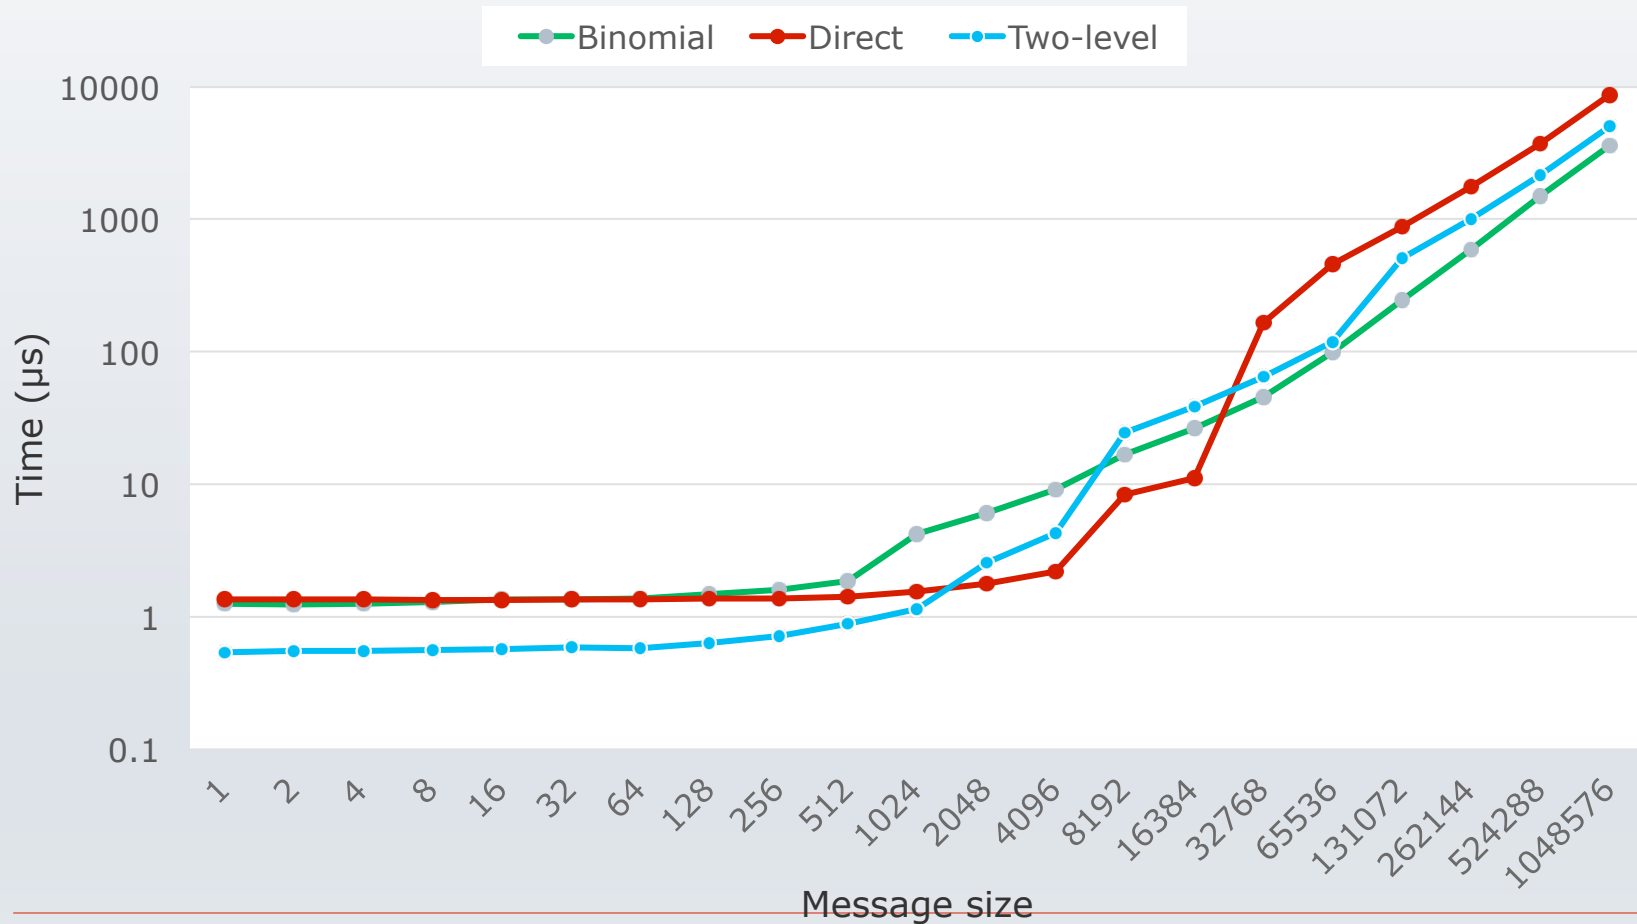
Goal of this presentation

- ❑ Why do we need different algorithms ?
- ❑ Which one is the best ? How to do the tuning ?
- ❑ Example of algorithms.
- ❑ What is the current state of research in this area ?

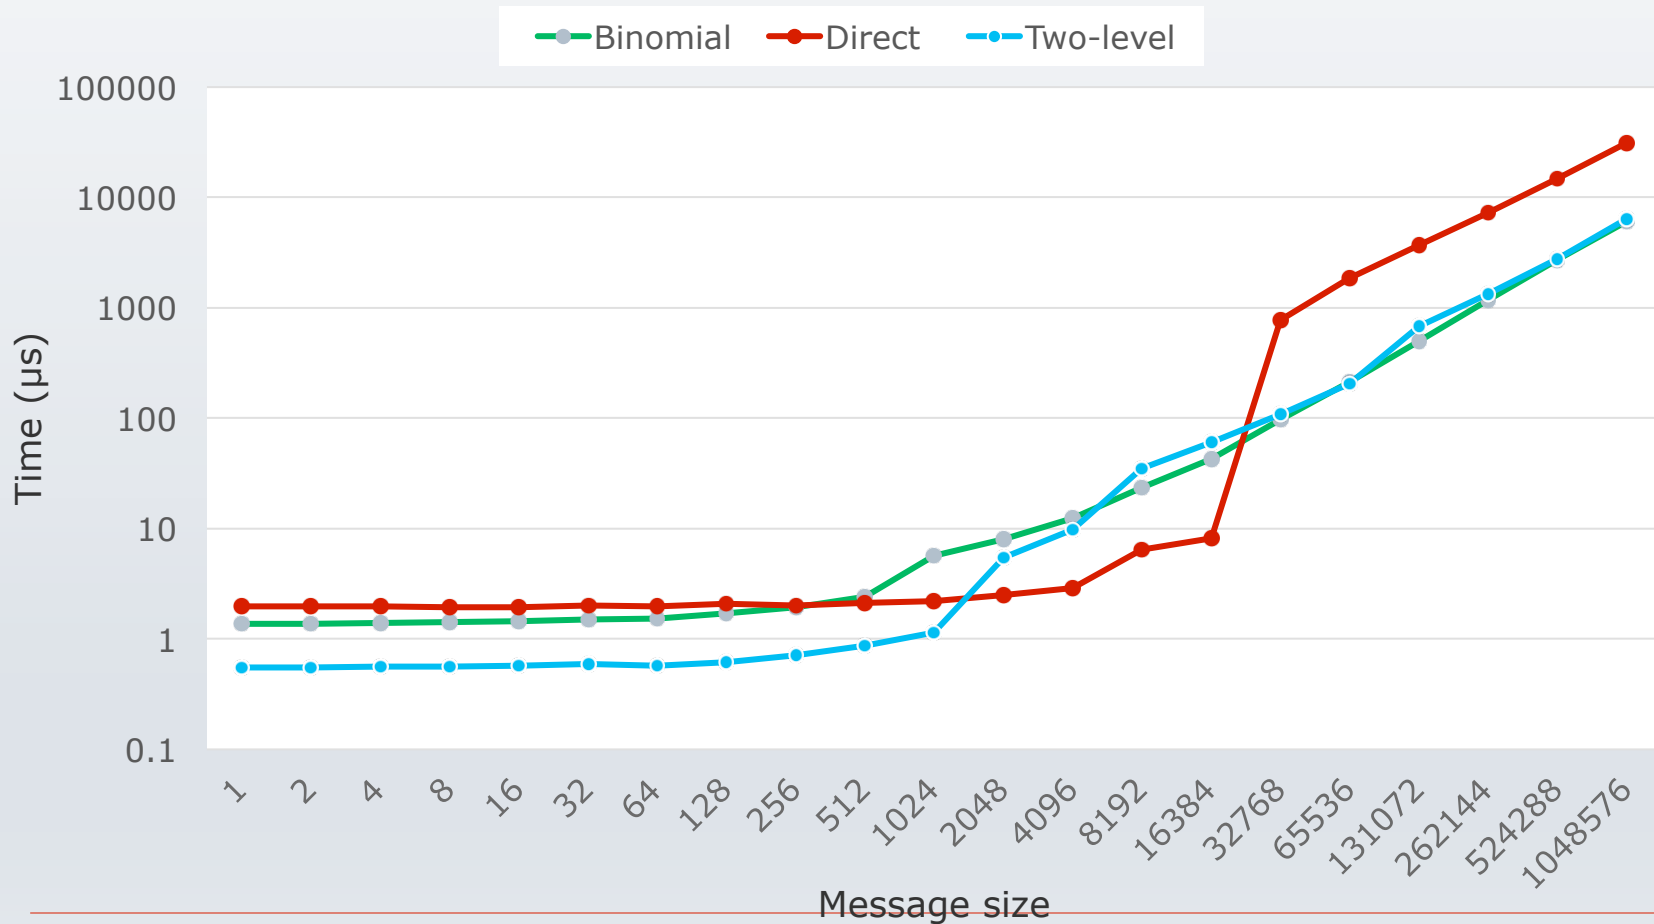
Behind Gather (64 cores/IMB)



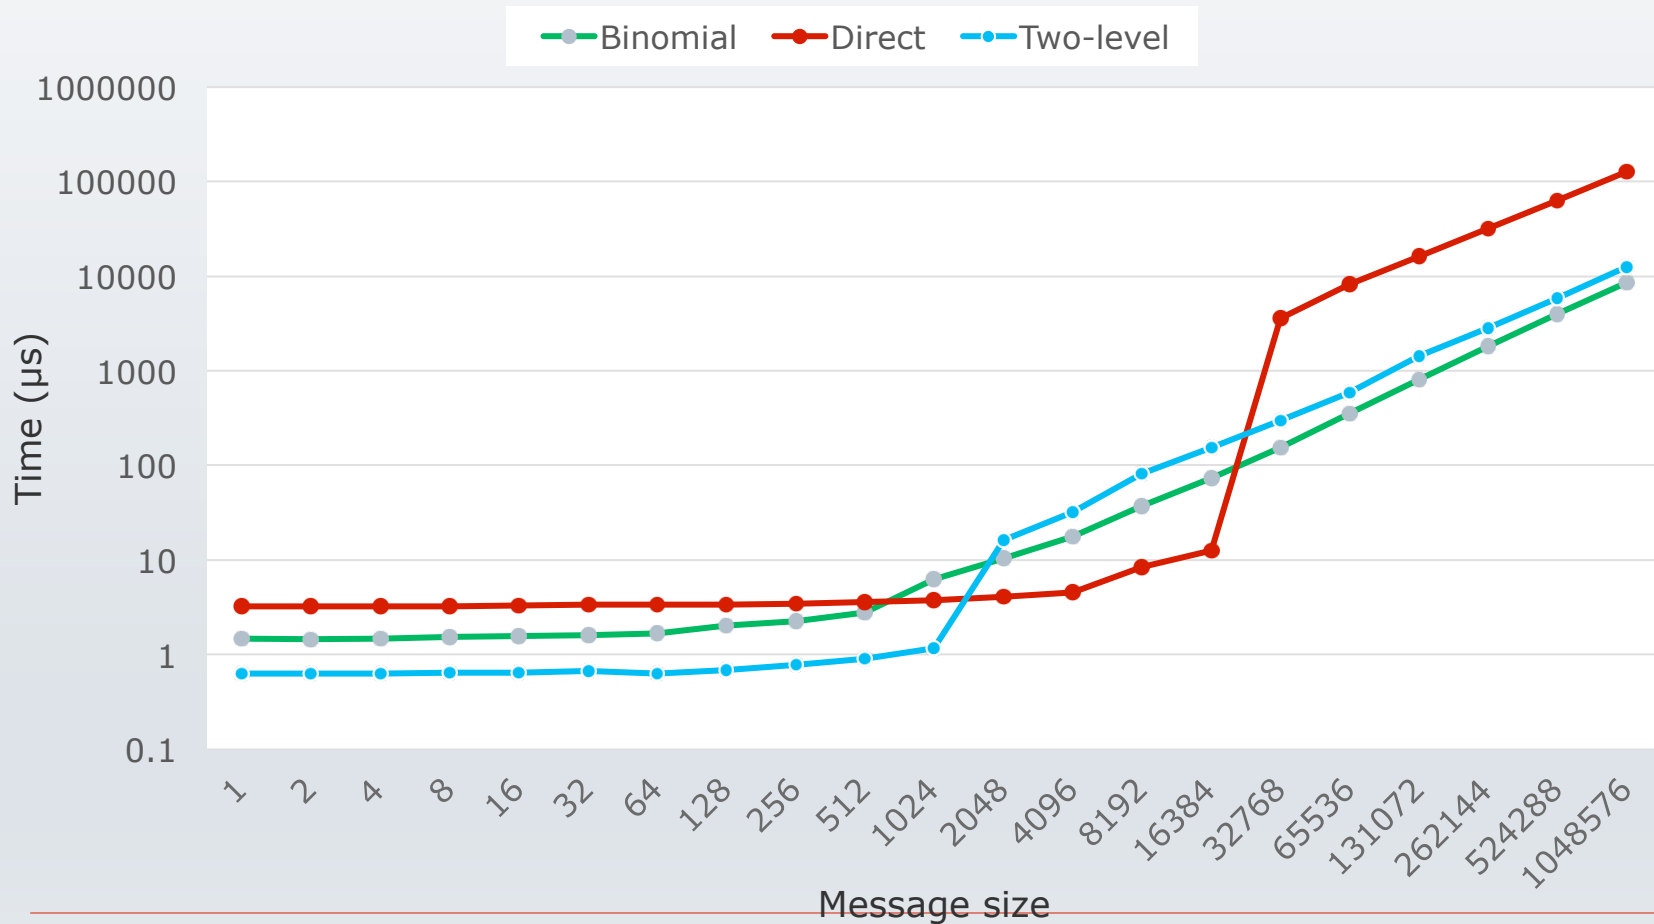
Behind Gather (64 cores/OMB)



Behind Gather (256 cores/OMB)



Behind Gather (1024 cores/OMB)



General Notes on Optimizing Collectives

- ❑ 2 components for collective communications – latency and bandwidth
- ❑ Latency(α) – time when the collective completes with the first byte
- ❑ Bandwidth(β) – rate at which collective proceeds after the first byte transmission
- ❑ Cost for communication – $\alpha + \beta m$ (Hockney)
- ❑ Latency is critical for small message sizes and bandwidth for large message sizes

MVAPICH Algorithm choices

- ☐ Depend on Platform
- ☐ Depend on system size
- ☐ Based on OSU Micro Benchmarks (OMB) results
- ☐ For each Platform:
 - Step 1: Point-to-Point Tuning
 - Step 2: Collective Tuning

Overview of algorithms

- ☐ Binomial Tree
- ☐ Barrier (Butterfly, Dissemination, Tournament...)
- ☐ Allgather
- ☐ Alltoall
- ☐ Reduce-Scatter for commutative operations: Recursive halving algorithm

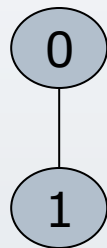
Binomial Tree

- **Definition (Binomial Tree)** The *binomial tree of order $k \geq 0$* with root R is the tree B_k defined as follows
 1. If $k=0$, $B_k = \{R\}$. i.e., the binomial tree of order zero consists of a single node, R .
 2. If $k>0$, $B_k = \{R, B_0, B_1, \dots, B_{k-1}\}$. i.e., the binomial tree of order $k>0$ comprises the root R , and k binomial subtrees, $B_0 - B_{k-1}$.
- B_k contains 2^k nodes
- The height of B_k is k
- The number of nodes at level l in B_k , where $0 \leq l \leq k$, is given by the *binomial coefficient* kC_l

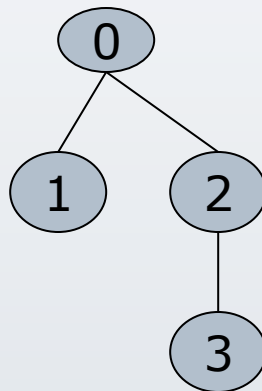
Binomial Trees



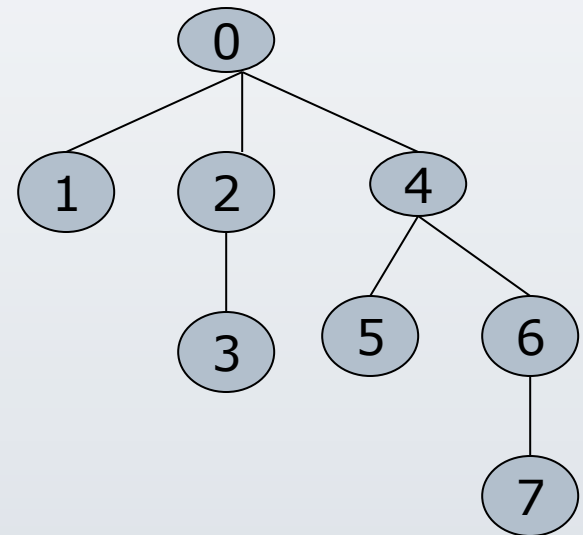
B₀



B₁

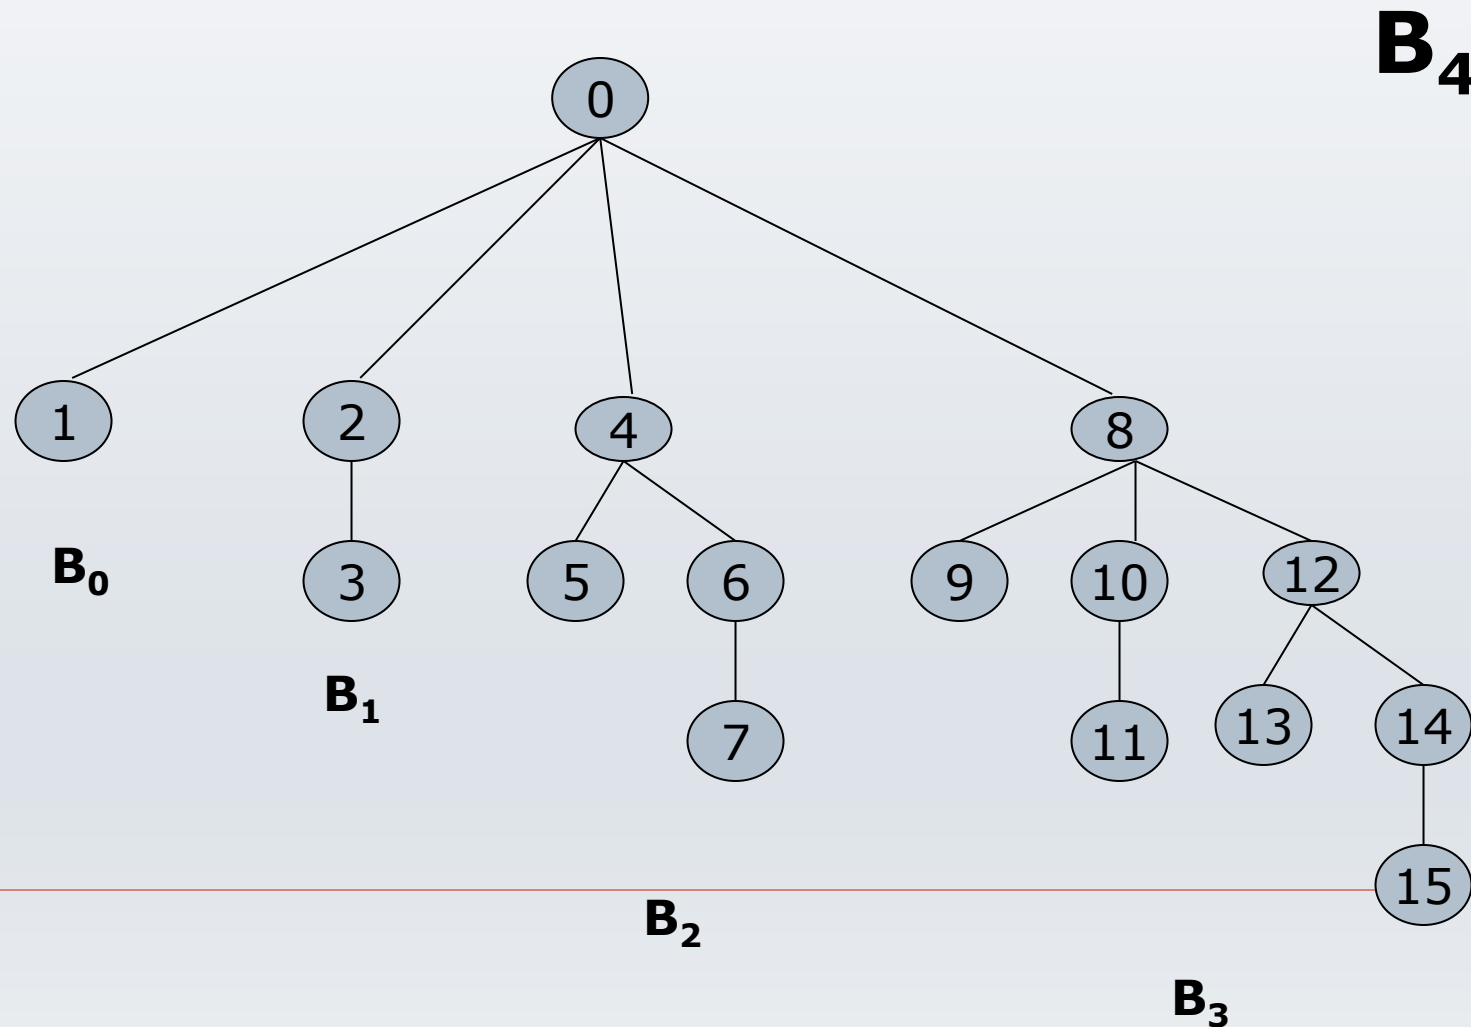


B₂



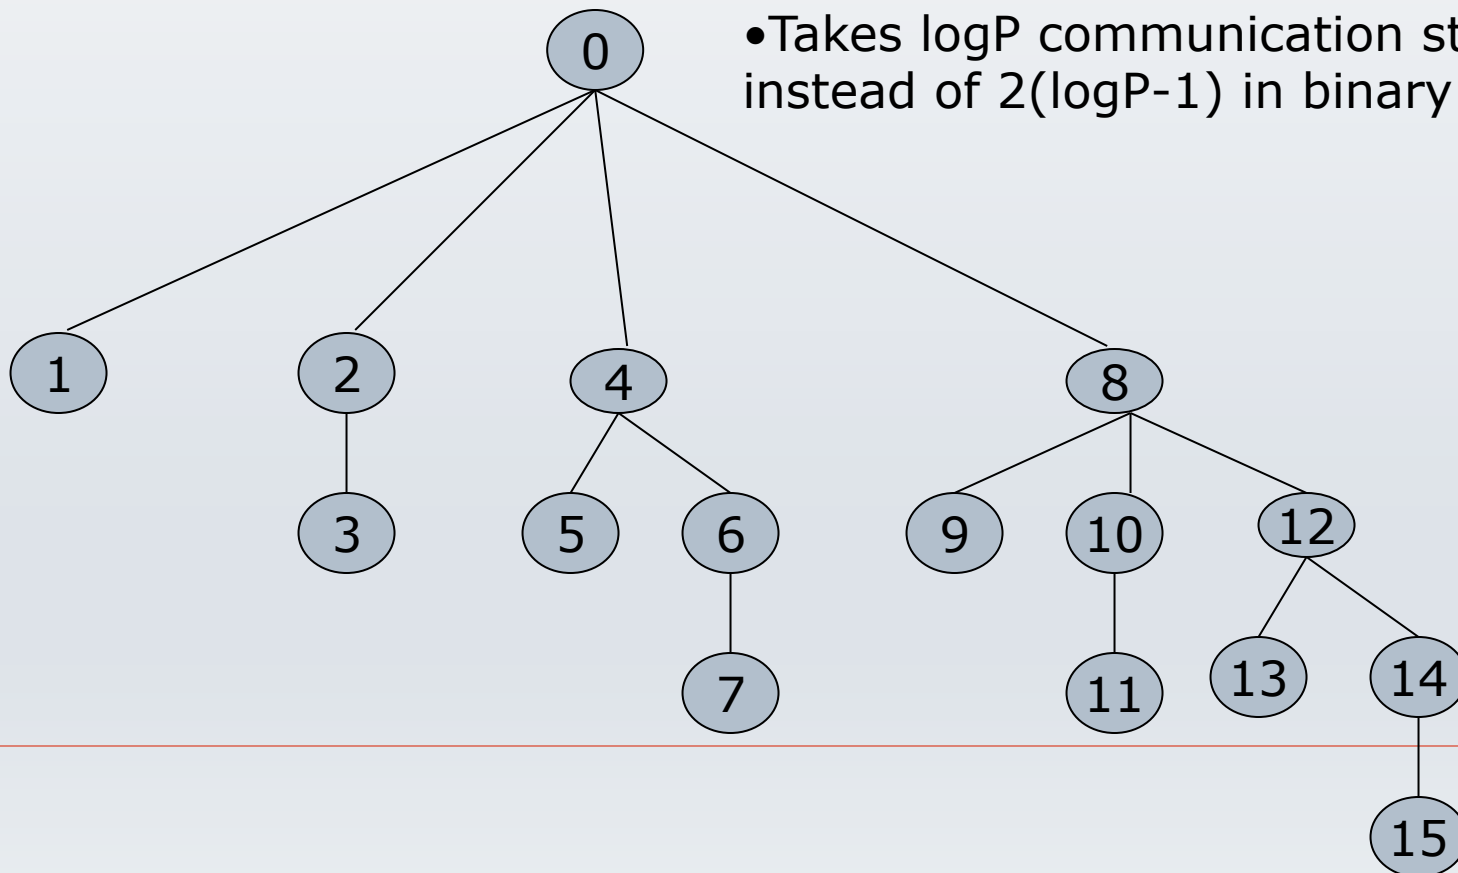
B₃

Binomial Trees



Binomial Trees

- Broadcast, Scatter and Gather usually implemented by binomial
- Takes $\log P$ communication steps instead of $2(\log P - 1)$ in binary



Barrier Algorithms

- ❑ **Butterfly barrier** by Eugene Brooks II
- ❑ In round k , i synchronizes with $i \oplus 2^k$ pairwise.
- ❑ If p not power of 2, existing procs. stand for missing ones.
- ❑ Worstcase – $2\log P$ pairwise synchronizations by a processor



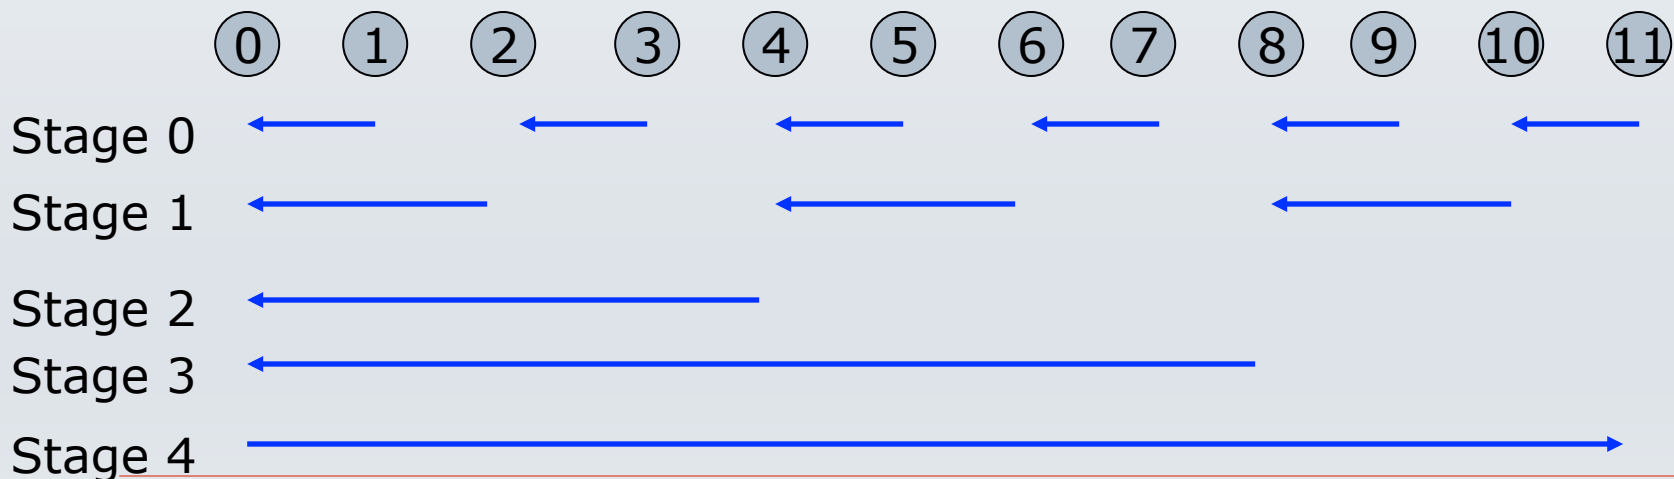
Barrier Algorithms

- ❑ **Dissemination barrier** by Hensgen, Finkel and Manser
- ❑ In round k , i signals $(i+2^k) \bmod P$
- ❑ No pairwise synchronization
- ❑ Same as butterfly but with different partners
- ❑ Almost $\log(\text{next power of } 2 > P)$ on critical path irrespective of P



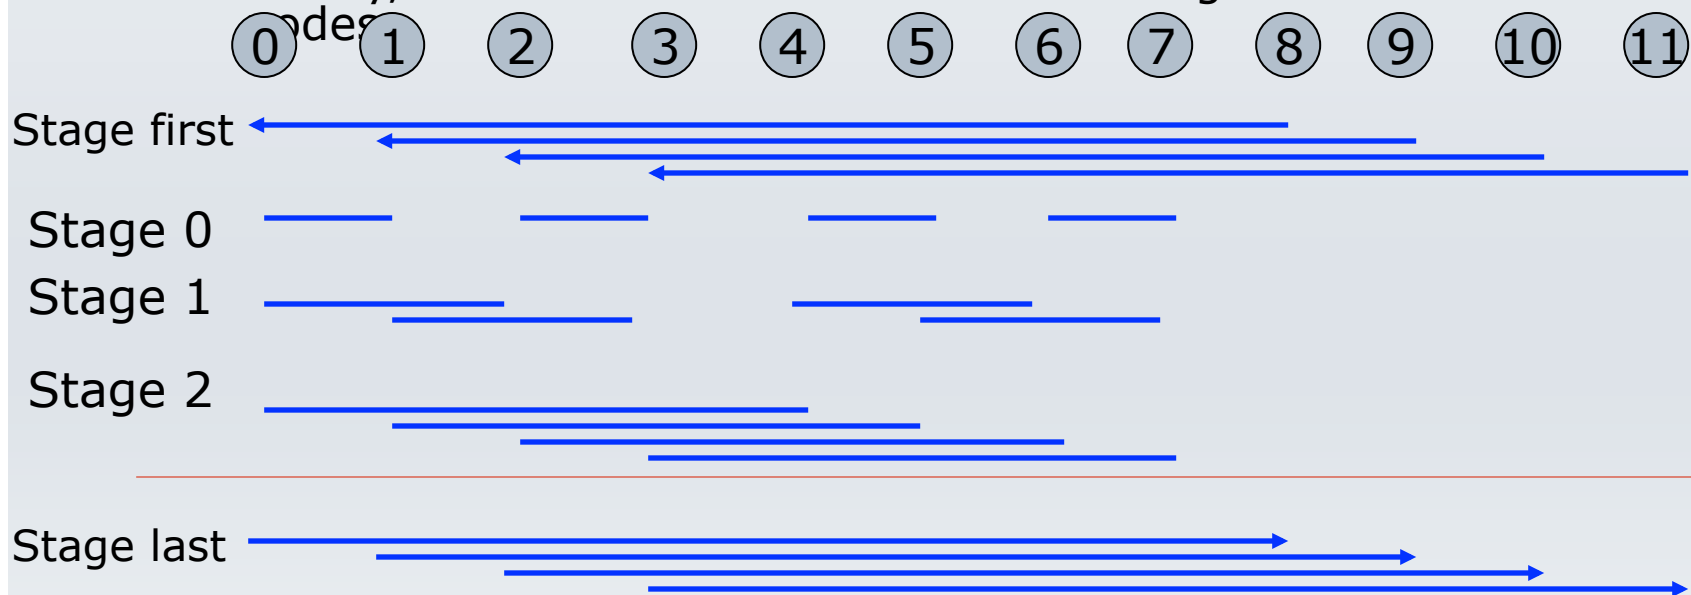
Barrier Algorithms

- ❑ **Tournament barrier** by Hensgen, Finkel and Manser
- ❑ In the 1st round, each pair of nodes (players) synchronize (play a game)
- ❑ The receiver will be considered as the winner of the game
- ❑ In the 2nd round, the winners of the 1st round will synchronize (play games)
- ❑ The receiver in the 2nd round will advance to the 3rd round
- ❑ This process continues till there is 1 winner left in the tournament
- ❑ The single winner then broadcasts a message to all the other nodes
- ❑ At each round k , proc. j receives a message from proc. i , where $i = j - 2^k$



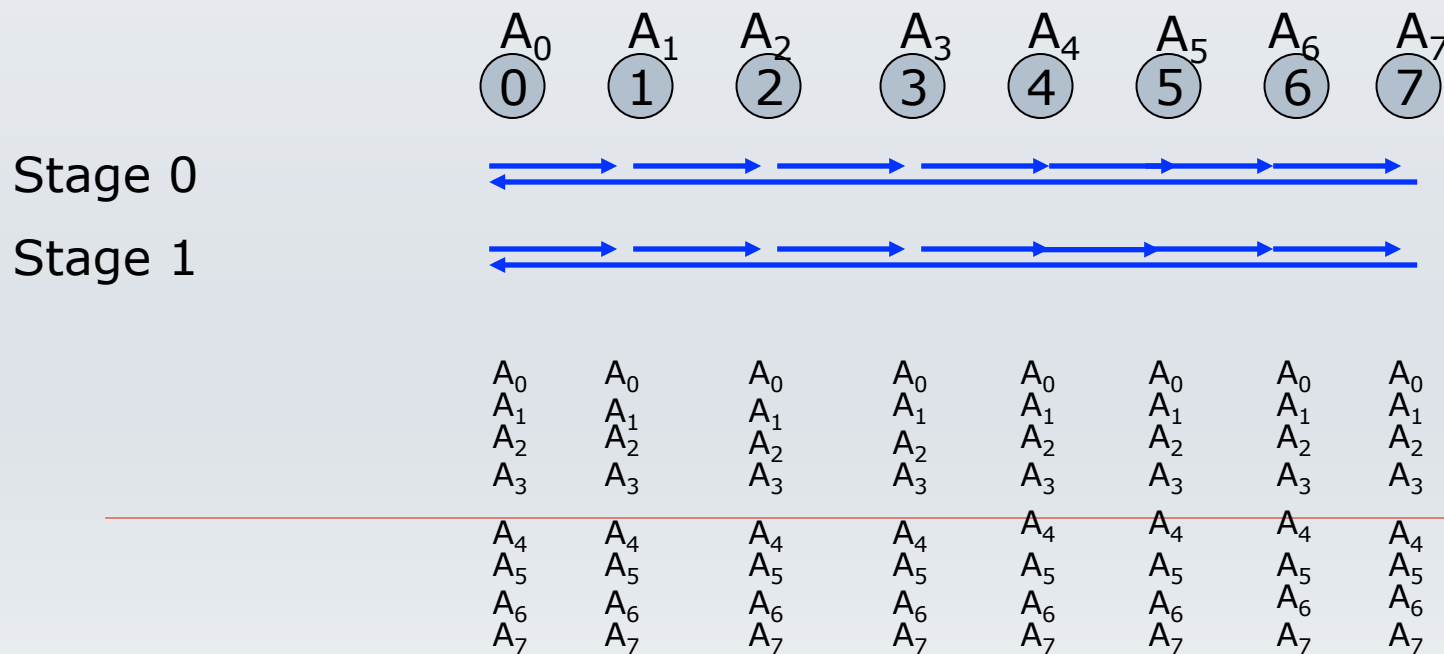
Barrier Algorithms

- ❑ **MVAPICH Barrier (pairwise exchange with recursive doubling)**
- ❑ Same as butterfly barrier.
- ❑ If nodes not equal to power, find the nearest power of 2, i.e. $m = 2^n$
- ❑ The last surfeit nodes, i.e. $\text{surfeit} = \text{size} - m$, initially send messages to the first surfeit number of nodes
- ❑ The first m nodes then perform butterfly barrier
- ❑ Finally, the first surfeit nodes send messages to the last surfeit



Allgather implementation

- ❑ In general, optimized allxxx operations depend on hardware topology, network contentions etc.
- ❑ Circular/ring allgather
- ❑ Each process receives from left and sends to right
- ❑ P steps



Bruck's Allgather

- Similar to dissemination barrier
- $\log P$ steps



AlltoAll

□ The naive implementation

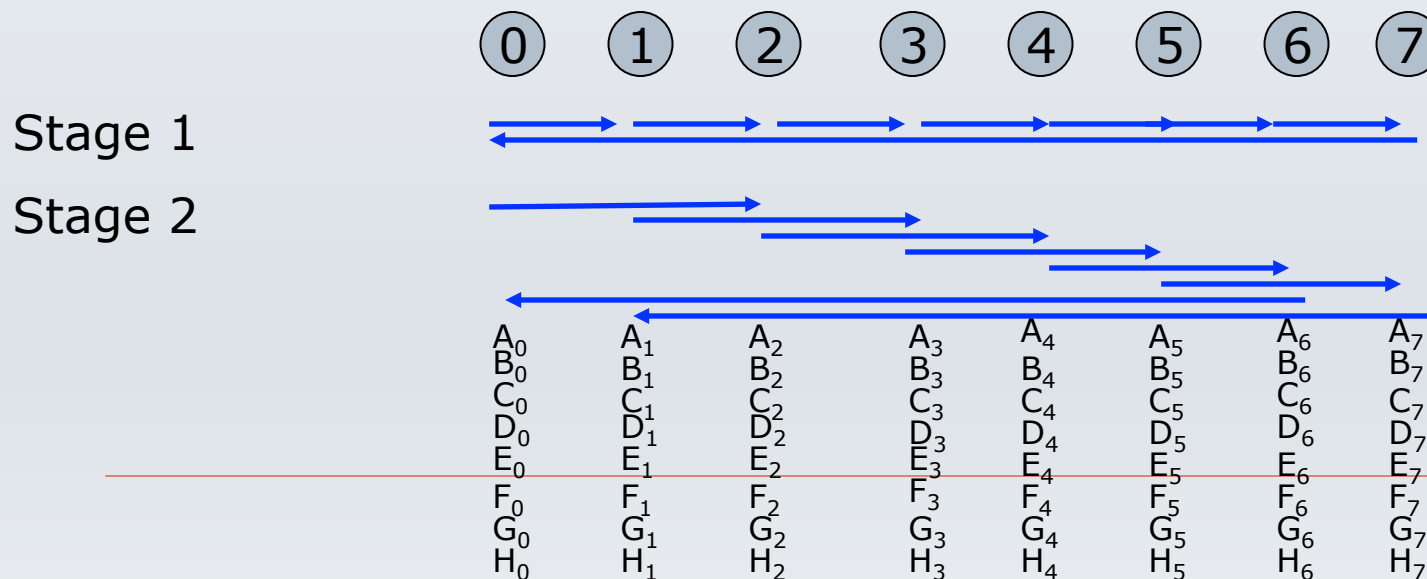
```
for all procs. i in order{  
    if i # my proc., then send to i and recv from i  
}
```

□ MPICH implementation – similar to naïve, but doesn't do it in order

```
for all procs. i in order{  
    dest = (my_proc+i)modP  
    src = (myproc-i+P)modP  
    send to dest and recv from src  
}
```

AlltoAll implementation

- Circular alltoall
- For step k in $\{1..P\}$, proc. i sends to $(i+k) \bmod P$ and receives from $(i-k + P) \bmod P$



Reduce and AllReduce

- ❑ Reduce and allreduce can be implemented with tree algorithms, e.g. binary tree
- ❑ But in tree based algorithms, some processors are not involved in computation
- ❑ Rolf Rabenseifner of Stuttgart – algorithms for reduce and allreduce

Reduce-Scatter for commutative operations: Recursive halving algorithm

- ❑ Recursive doubling – in the first step, communication is with the neighboring process. In each step, the communication distance doubles
- ❑ Recursive halving – reverse of recursive doubling
- ❑ At the first step
 - a process communicates with another process $P/2$ away
 - sends data needed by the other half
 - Receives data needed by its half
 - Performs operation
- ❑ Next step – distance $P/4$ away and so on...
- ❑ $\lg P$ steps

Reduce-Scatter for non-commutative operations: Recursive doubling algorithm

- ❑ In the first step, data (all data except the one needed for its result) is exchanged with the neighboring process
- ❑ In the next step, $(n-2n/p)$ data (all except the one needed by it and the one needed by process it communicated with the previous step) is communicated with process that is distance 2 apart
- ❑ In the third step $(n-4n/p)$ data with process that is distance 4 apart and so on...
- ❑ $\lg P$ steps

Example – Broadcast (MPICH)

- Binomial Broadcast
 - $\log p$ steps
 - Amount of data communicated at each step - n
 - $\text{cost} = \log p (\alpha + n\beta)$
- scatter and allgather
 - Divide message into p segments
 - Scatter the p segments to p processes using binomial scatter – $\log p \alpha + (n/p)(p-1) \beta$
 - Scattered data collected at all processes using ring allgather – $(p-1) \alpha + (n/p)(p-1) \beta$
 - $\text{cost} = (\log p + p-1) \alpha + 2(n/p)(p-1) \beta$
- Hence binomial broadcast for small messages and (scatter+allgather) for long messages

Two-level Algorithms of MVAPICH

- ❑ Take into account the node topology
- ❑ Notion of “node leader”
- ❑ Inter-node communications are between node leaders
- ❑ Intra-node communications are between the node leader and other processes
- ❑ Used for Bcast, Reduce, Gather...

MPICH Algorithms

- Allgather
 - Bruck Algorithm (variation of dissemination) (< 80 KB) and non-power-of-two
 - Recursive doubling (< 512 KB) for power-of-2 processes
 - ring (> 512 KB) and (80-512 KB) for any processes
- Broadcast
 - Binomial (< 12 KB), binomial scatter + ring allgather (> 512 KB)
- Alltoall
 - Bruck's algorithm (for < 256 bytes)
 - Post all irecv's and isends (for medium size messages. 256 bytes – 32 KB)
 - Pairwise exchange (for long messages and power-of-2 processors) – $p-1$ steps. In each step k , each process i exchanges data with $(i \text{ xor } k)$
 - For non-power of 2, an algorithm in which in each step, k , process i sends data to $(i+k)$ and receives from $(i-k)$

MPICH Algorithms

- Reduce-scatter
 - For commutative operations:
 - Recursive halving (< 512 KB), pairwise exchange (> 512 KB; $p-1$ steps; rank+ i at step i)
 - For non-commutative:
 - Recursive doubling (< 512 bytes), pairwise exchange (> 512 bytes)
- Reduce
 - Binomial algorithm (< 2 KB), Rabenseifner (> 2 KB)
- AllReduce
 - Recursive doubling (short) , Rabenseifner (long messages)

Current state of research

- ☐ Socket level leaders
- ☐ Auto - Selection of best algorithms (Tuning)
- ☐ Usage of CMA (Cross Memory Attach)
- ☐ Support for GPUs and MICs
- ☐ Topology aware collectives
- ☐ Energy Aware collectives
- ☐ Etc....

References

- ❑ Thakur et. al. – Optimization of Collective Communication Operations in MPICH. IJHPCA 2005.
- ❑ Thakur et. al. - Improving the Performance of Collective Operations in MPICH. EuroPVM/MPI 2003.
- ❑ <http://mvapich.cse.ohio-state.edu/publications/>