

Parallel Computing for Science & Engineering Spring 2013: MPI introduction

Instructors:

Victor Eijkhout, Research Scientist, TACC

Kent Milfeld, Research Associate, TACC



Outline

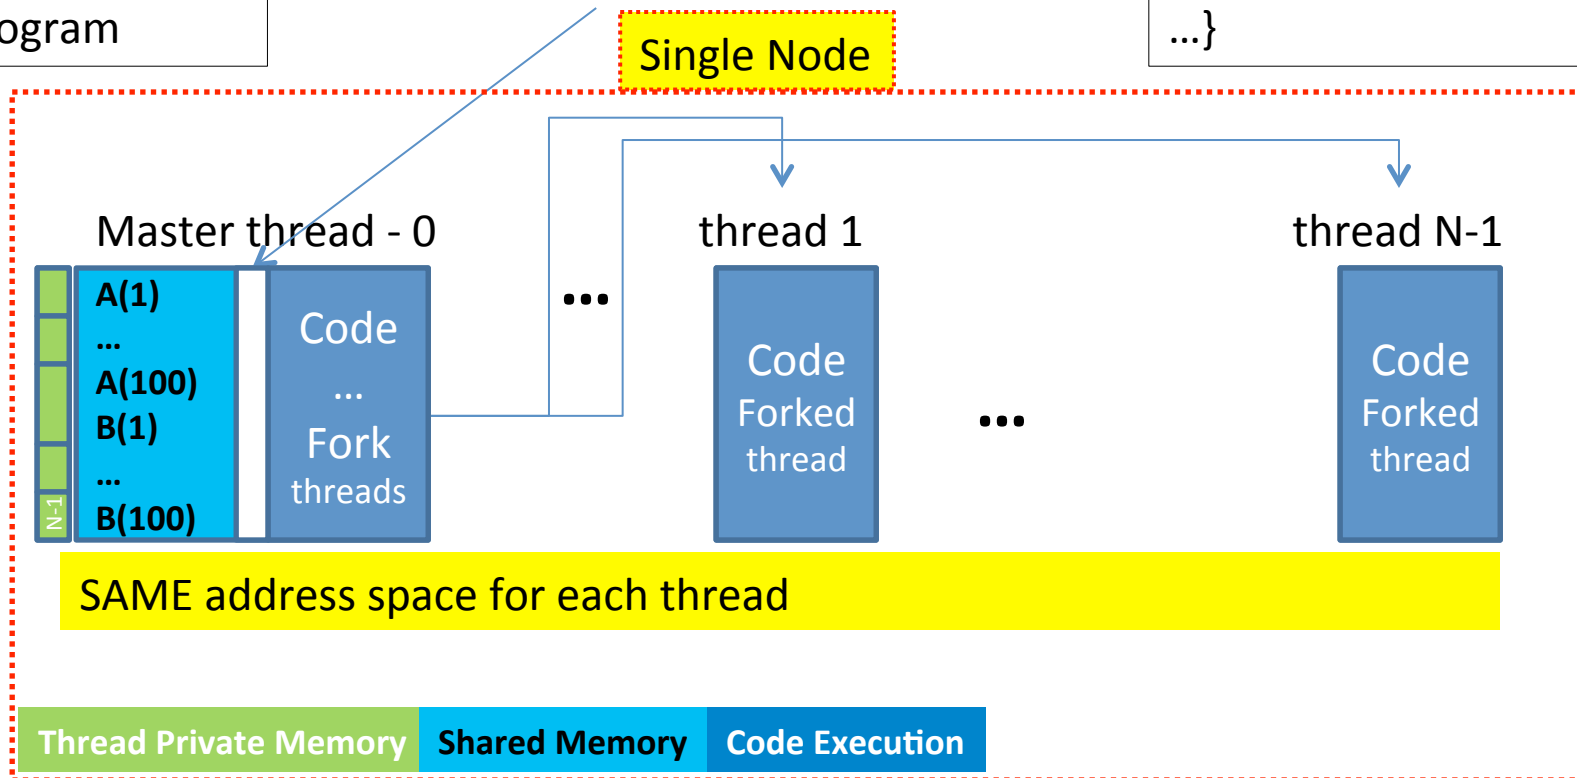
- Executing OpenMP and MPI
- Paradigm/Key Concepts/Advantages
- MPI History version 1 and 2, implementations
- Compiling, Running
- MPI Initialize, Finalize and task-id/task-count
- MPI Communicators

OpenMP (shared memory)

```
program myomp
real*8 a(100), b(100)
...
!$omp parallel
... end program
```

```
Compile → a.out
Set OMP_NUM_THREADS to N
./a.out
```

```
int main () {
double a(100), b(100);
...
#pragma omp parallel
...}
```

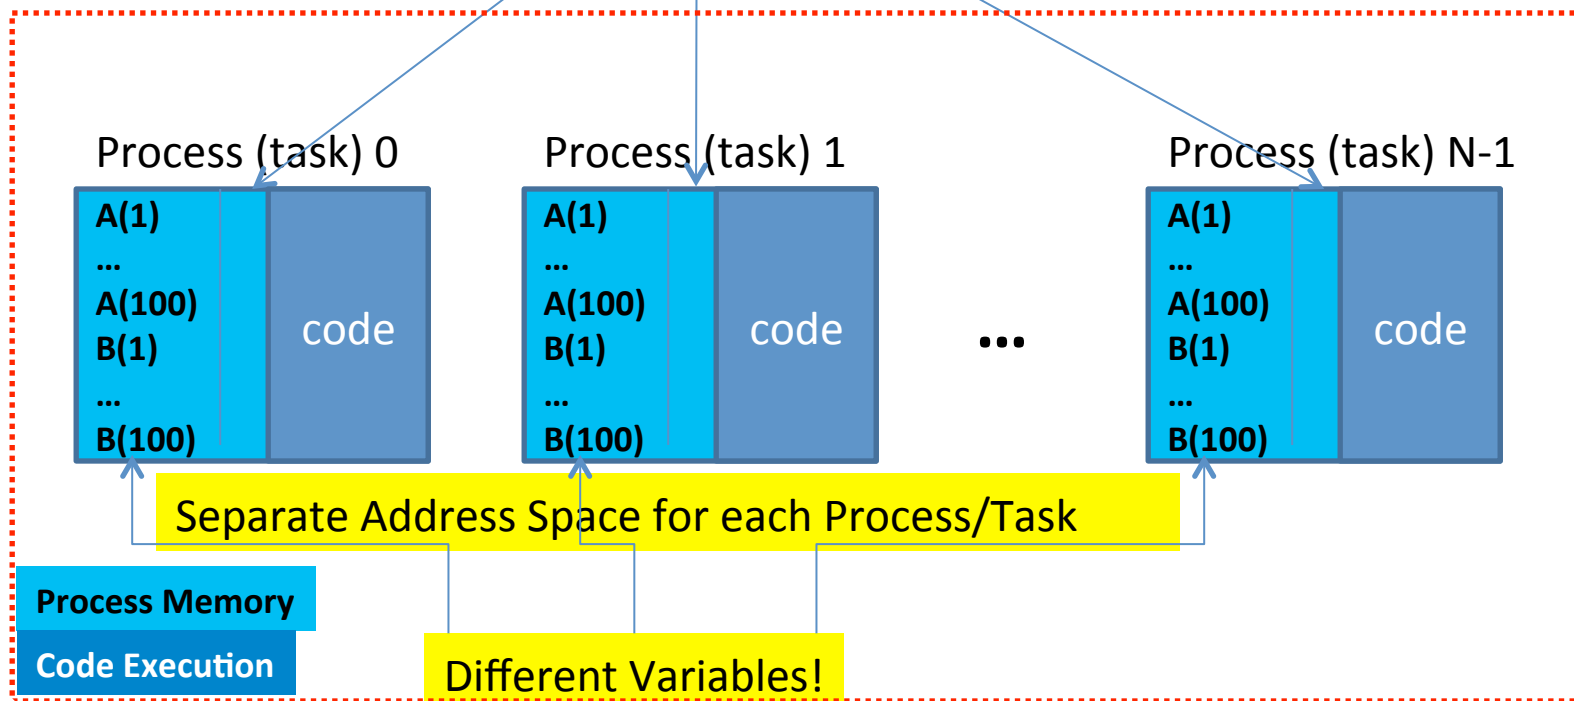


MPI (distributed memory)

```
program mympi  
real*8 a(100), b(100)  
...  
end program
```

Compile → a.out
Launch on N cores
ibrun ./a.out

```
int main () {  
double a(100), b(100);  
...  
}
```



Message Passing Paradigm

- A Parallel **MPI Program is launched as separate processes (tasks)**, each with their own address space.
 - Requires partitioning data across tasks.
- **Data is explicitly moved** from task to task
 - A task accesses the data of another task through a transaction called “message passing” in which a copy of the data (message) is transferred (passed) from one task to another.
- There are two classes of message passing (transfers)
 - **Point-to-Point messages** involve only two tasks
 - **Collective messages** involve a set of tasks
- Access to subsets of complex data structures is simplified
 - A **data subset is described as a single Data Type** entity
- Transfers use **synchronous or asynchronous protocols**
- Messaging can be arranged into efficient topologies

Key Concepts-- Summary

- Used to create parallel **SPMD** programs on distributed-memory machines with explicit message passing
- Routines available for
 - Point-to-Point Communication
 - Collective Communication
 - 1-to-many
 - many-to-1
 - many-to-many
 - Data Types
 - Synchronization (barriers, non-blocking MP)
 - Parallel IO
 - Topologies

Advantages of Message Passing

- Universality
 - Message passing model works on separate processors connected by any network (and even on shared memory systems)
 - matches the hardware of most of today's parallel supercomputers as well as ad hoc networks of computers
- Performance/Scalability
 - Scalability is the most compelling reason why message passing will remain a permanent component of HPC (High Performance Computing)
 - As modern systems increase core counts, management of the memory hierarchy (including distributed memory) is the key to extracting the highest performance
 - Each message passing process only directly uses its local data, avoiding complexities of process-shared data, and allowing compilers and cache management hardware to function without contention.

MPI-1

- MPI-1 - Message Passing Interface (v. 1.2)
 - Library
 - Specification: defined by committee of vendors, implementers, and parallel programmers
 - Designed with SPMD (single program, multiple data) technique in mind.
- Available on almost all parallel machines in C/C++ and Fortran
- About 125 routines
 - 6 basic routines
 - the rest are extensions that can simplify algorithm implementation and optimize performance

MPI-1

Web

www-unix.mcs.anl.gov/mpi/

www.mcs.anl.gov/research/projects/mpich2/

hwww.mpi-forum.org/

Books

Using MPI, by Gropp, Lusk, and Skjellum

MPI Annotated Reference Manual, by Marc Snir, *et al*

Parallel Programming with MPI, by Peter Pacheco

Using MPI-2, by Gropp, Lusk and Thakur

Getting Started

www.mcs.anl.gov/research/projects/mpi/tutorial/gropp/talk.html

<http://ci-tutor.ncsa.illinois.edu/>

www.nersc.gov/nusers/help/tutorials/mpi/intro/ (simple, direct)

<https://computing.llnl.gov/?set=training&page=index>

Advanced: www.mcs.anl.gov/research/projects/mpi/tutorial/

Standard

www.mpi-forum.org/docs/

MPI-1 Implementations

- Many parallel machine, HPC interconnect, and commercial software vendors have optimized versions
 - Hardware vendors: IBM, Sun, HP, Intel
 - Interconnect vendors: Myricom, Quadrics, InfiniBand*
 - * InfiniBand: open source drivers/university MPI collaboration
 - <http://www.openib.org/>
 - <http://mvapich.cse.ohio-state.edu/>
 - Software vendors: MPI/Pro, Platform MPI (was Scali MPI), etc.
- Others
 - MPICH,
 - MPICH-G2, Globus-based www-unix.mcs.anl.gov/mpi/mpich/www3.niu.edu/mpi/
 - MPICH-VMI, Virtual Machine Interface <http://vmi.ncsa.uiuc.edu/>
 - MVAPICH <http://mvapich.cse.ohio-state.edu/>
 - OpenMPI (the MPI formerly LAM/MPI) <http://www.open-mpi.org/>

MPI-2

- Includes features left out of MPI-1
 - One-sided communications
 - Dynamic process control
 - More complicated collectives
 - MPI-IO
- Implementations
 - not quickly undertaken after the standard document was released (in 1997)
 - now OpenMPI, MPICH2 (and its descendants), and the vendor implementations are pretty complete or fully complete

Compiling MPI Programs

- Generally use a special compiler or compiler wrapper script
 - not defined by the standard
 - consult your implementation
 - handles correct include path, library path, and libraries
- MPICH-style (the most common)
 - C

```
mpicc -o mycexe mycode.c
```
 - Fortran

```
mpif90 -o myfexe mycode.f
```

Running MPI Programs

- MPI programs require some help to get started
 - what computers should I run on?
 - how do I access them?
- MPICH-style

```
mpirun -np 10 -machinefile mach ./a.out
```
- When batch systems are involved, all bets are off
@TACC Lonestar/Longhorn/Ranger (via a job script)


```
ibrun tacc_affinity ./a.out
```

 - SGE (Sun Grid Engine) batch utility handles the rest

The Parallel Code

- Parallel executables are nothing more than independent processes launches by ssh commands: `ssh <nodename> <environment> executable`.
 - Executables need organization info (initialize).
 - Executable needs to synchronize.
 - Program needs to know its id and # of execs.
 - Executable needs to clean up at end.

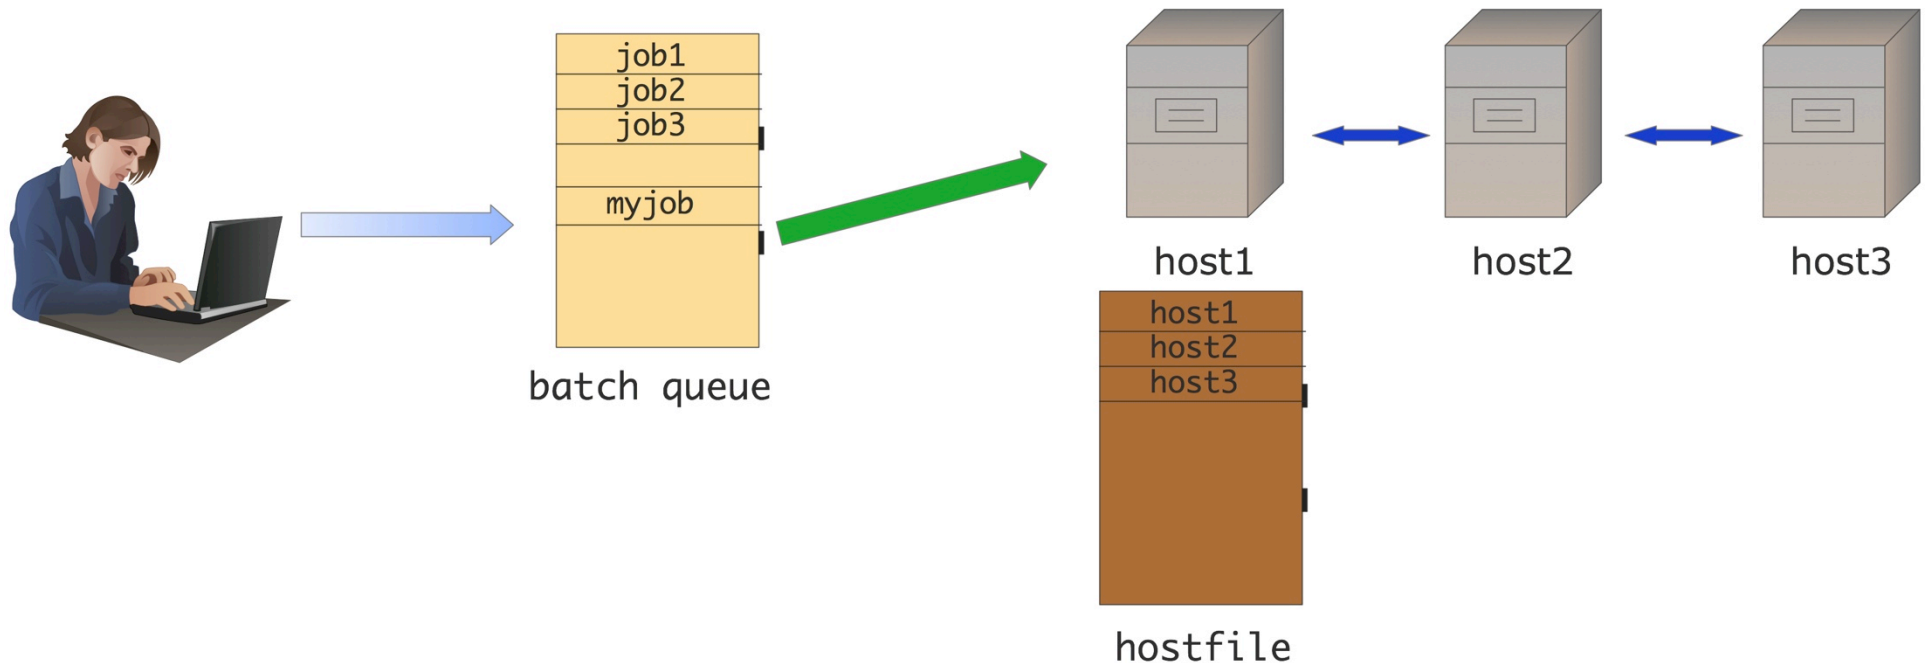
Interactive scenario



User type:

`mpirun -np 5 hostfile ./myprogram arguments`

Batch scenario



- User submits batch job to queue, executed later by scheduler

Minimal MPI program

- Every MPI program needs these...
 - C version

```
#include <mpi.h>
...
ierr=MPI_Init(&argc, &argv);
ierr=MPI_Comm_size(MPI_COMM_WORLD, &npes);
ierr=MPI_Comm_rank(MPI_COMM_WORLD, &iam);
...
ierr=MPI_Finalize();
```

In C MPI routines are functions which return the error value

Minimal MPI program

- Every MPI program needs these...

– Fortran version

```
include 'mpif.h'           or      use mpi
...
call MPI_Init(ierr)
call MPI_Comm_size(MPI_COMM_WORLD, npes, ierr)
call MPI_Comm_rank(MPI_COMM_WORLD, iam, ierr)
...
call MPI_Finalize(ierr)
```

In Fortran, MPI routines are subroutines with the last parameter as the error value

MPI Initialization & Termination

- All processes must initialize and finalize MPI (each is a **collective call***).
 - **MPI_Init** : starts up the MPI runtime environment
 - **MPI_Finalize** : shuts down the MPI runtime environment
- Must include header files – provides basic MPI definitions and types.
 - Header File

Fortran 77	Fortran 90	C/C++
include 'mpif.h'	use mpi	#include <mpi.h>

- Format of MPI calls

Fortran 77/90 binding (upper or lower case)	C/C++ binding
CALL MPI_XYYY(parameters..., ierr)	ierr = MPI_Xyyy(parameters...)

* Means the entire group of tasks must execute this call.

Run Parameters

- **MPI_Comm_size** : gets the number of processes in a run
Integer
(typically called just after **MPI_Init**).
- **MPI_Comm_rank** : gets the process ID (rank) of the current process,
integer between 0 and $NP-1$ inclusive
(typically called just after **MPI_Init**).

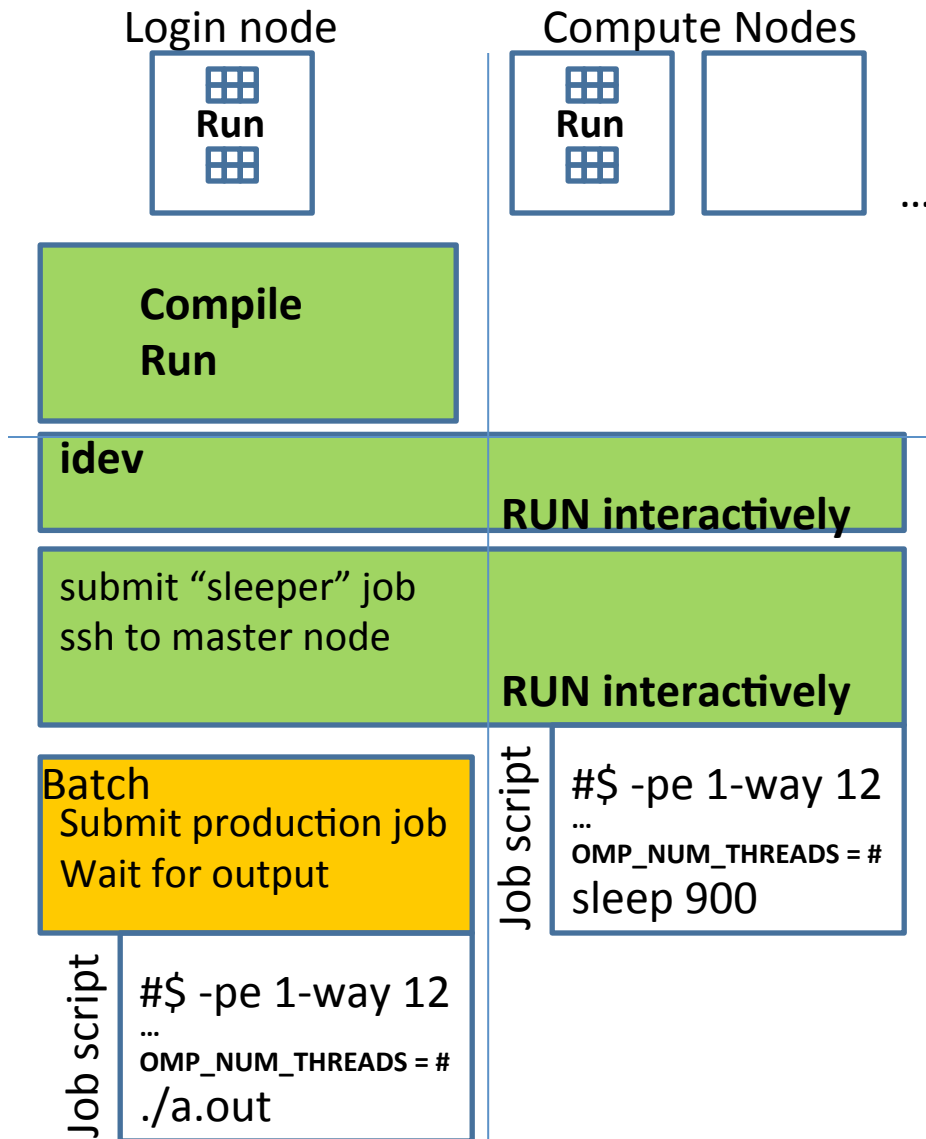
Communicators

- Communicators
 - MPI uses a communicator objects (and groups) to identify a **set of processes which communicate only within their set.**
 - MPI_COMM_WORLD is defined in the MPI include file as **all processes** (ranks) of your job
 - **Required** parameter **for most MPI calls**
 - You **can create subsets** of MPI_COMM_WORLD
- Rank
 - Unique **process ID** within a communicator
 - Assigned by the system when the process initializes (for MPI_COMM_WORLD)
 - Processors within a communicator are assigned numbers **0 to n-1** (C/F90)
 - Used to specify sources and destinations of messages, process specific indexing and operations.

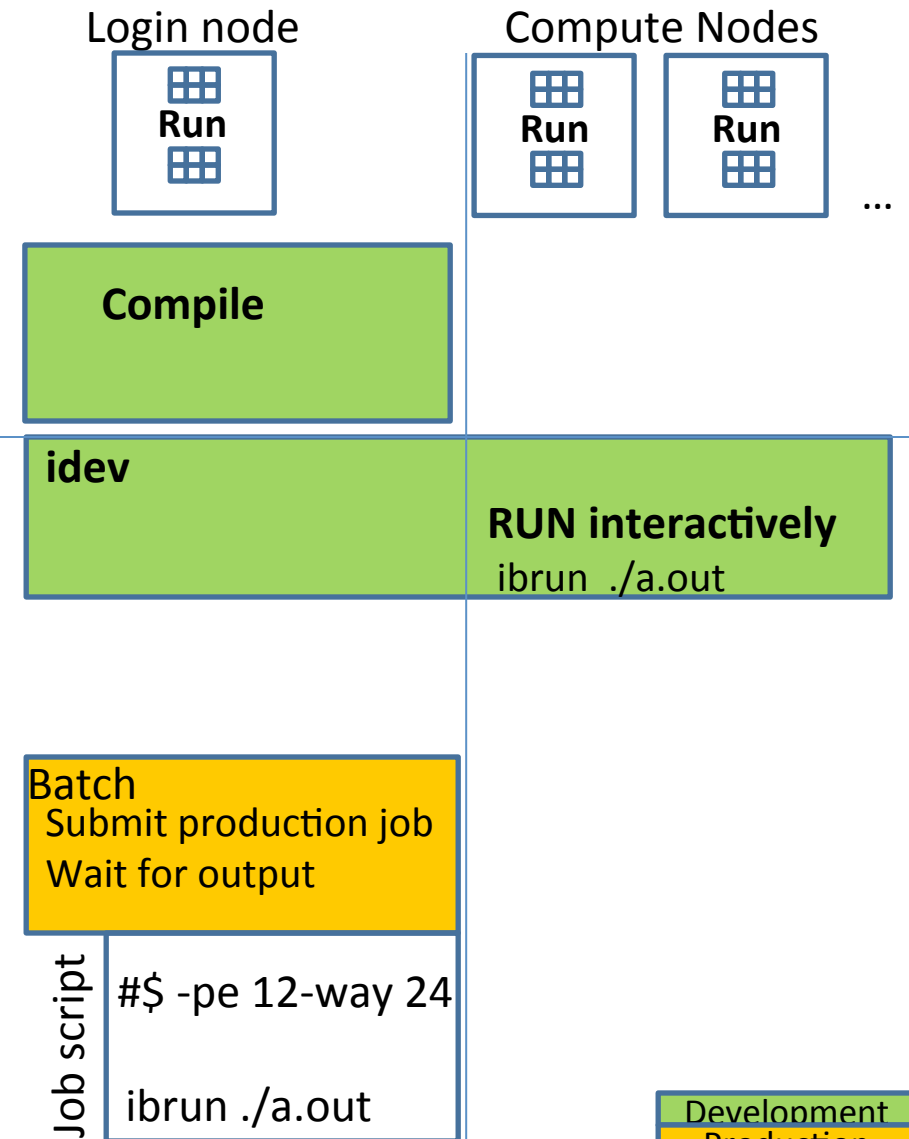
Include files

- The MPI include file
 - C:
`mpi.h`
 - Fortran
`mpif.h`
MPI module → `use MPI`
- Defines many constants used within MPI programs
 - In C, defines the interfaces for the `functions`
 - In C++, the interfaces are different, so be careful
 - In F90, module defines interface for `subroutines`
- Compilers know where to find the include files
 - regular compilers are usually called through `mpif90/`
`mpicc wrapper scripts`

OpenMP World




MPI World



Development
Production

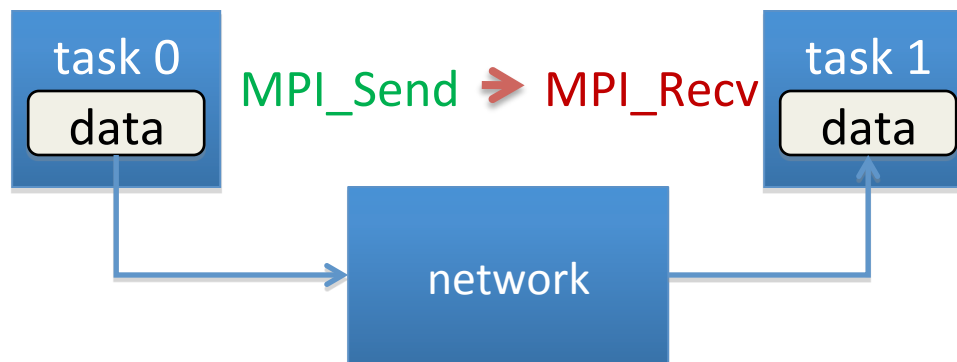
Parallel Code

- The programmer is responsible for determining all parallelism.
 - Data Partitioning
 - Deriving Parallel Algorithms
 - Moving Data between Processes
- Tasks (independent processes executing anywhere) send and receive “messages” to exchange data.
- Data transfer requires cooperative operation to be performed by each process (point to point communications).
- Message Passing Interface (MPI) was released in 1994. (MPI-2 in 1996) Now the MPI is the de facto standard for message passing.
- <http://www-unix.mcs.anl.gov/mpi/>



Point-to-Point Communication

- Sending data from one point (process/task) to another point (process/task)
- One task **sends** while another **receives**



Basic Communications in MPI

- Standard `MPI_Send/MPI_Recv` routines
 - Used for basic messaging

Modes of Operation

- Blocking
 - Call does not return until the `data area is safe to use`
- Non-blocking
 - Initiates send or receive operation, returns immediately
 - Can check or wait for completion of the operation
 - `Data area is not safe to used until completion.`
- Synchronous and Buffered (later)

Data Types (basics)

- Data types (more a mapping than declaration)
 - Specifies the data type and size in MPI routines
 - Predefined MPI types correspond to language types

Representation	MPI Type Fortran	Fortran	MPI Type C	C
32-bit floating point	MPI_REAL	REAL	MPI_FLOAT	float
64-bit floating point	MPI_DOUBLE_PRECISION	DOUBLE_PRECISION	MPI_DOUBLE	double
32-bit integer	MPI_INTEGER	INTEGER	MPI_INT	int

- Methods exists for creating user-defined types
 - Simple (just combinations of normal data types)
 - Advanced (a map of data to be send)