

# Scalability of Collective Operations

Victor Eijkhout

- Cost analysis
- Dense matrix-vector product

## Cost analysis

# Concepts

- $\alpha$ : message latency
- $\beta$ : transfer time per byte
- $\gamma$ : time for floating point operation

Recall definitions of weak/strong scalability

# Dense matrix-vector product

# Parallel matrix-vector product; dense

- Assume a division by block rows
- Every processor  $p$  has a set of row indices  $I_p$

Mvp on processor  $p$ :


$$\forall_i: y_i = \sum_j a_{ij}x_j$$

$$\forall_i: y_i = \sum_q \sum_{j \in I_q} a_{ij}x_j$$

Local and remote parts:

$$\forall_i: y_i = \sum_{j \in I_p} a_{ij} x_j + \sum_{q \neq p} \sum_{j \in I_q} a_{ij} x_j$$

Each processor needs to collect the whole vector: Allgather



# Cost computation 1.

Algorithm:

Step	Cost (lower bound)
Allgather $x_i$ so that $x$ is available on all nodes	
Locally compute $y_i = A_i x$	$\approx 2 \frac{n^2}{P} \gamma$



# Allgather

Assume that data arrives over a binary tree:

- latency  $\alpha \log_2 P$
- transmission time, receiving  $n/P$  elements from  $P - 1$  processors

Algorithm with cost:

Step	Cost (lower bound)
Allgather $x_i$ so that $x$ is available on all nodes	$\lceil \log_2(P) \rceil \alpha + \frac{P-1}{P} n \beta \approx \log_2(P) \alpha + n \beta$
Locally compute $y_i = A_i x$	$\approx 2 \frac{n^2}{P} \gamma$

# Parallel efficiency

$$E_p^{1\text{D-row}}(n) = \frac{S_p^{1\text{D-row}}(n)}{p} = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{p}{2n} \frac{\beta}{\gamma}}.$$

Strong scaling, weak scaling?

# Two-dimensional partitioning

$x_0$ $a_{00}$ $a_{01}$ $a_{02}$ $y_0$ $a_{10}$ $a_{11}$ $a_{12}$ $a_{20}$ $a_{21}$ $a_{22}$ $a_{30}$ $a_{31}$ $a_{32}$	$x_3$ $a_{03}$ $a_{04}$ $a_{05}$ $y_1$ $a_{13}$ $a_{14}$ $a_{15}$ $a_{23}$ $a_{24}$ $a_{25}$ $a_{33}$ $a_{34}$ $a_{35}$	$x_6$ $a_{06}$ $a_{07}$ $a_{08}$ $a_{16}$ $a_{17}$ $a_{18}$ $a_{26}$ $a_{27}$ $a_{28}$ $y_2$ $a_{36}$ $a_{37}$ $a_{38}$	$x_9$ $a_{09}$ $a_{10}$ $a_{11}$ $a_{12}$ $a_{19}$ $a_{20}$ $a_{21}$ $a_{22}$ $a_{29}$ $a_{30}$ $a_{31}$ $a_{32}$ $a_{39}$ $a_{40}$ $a_{41}$ $a_{42}$
$x_1$ $a_{40}$ $a_{41}$ $a_{42}$ $y_4$ $a_{50}$ $a_{51}$ $a_{52}$ $a_{60}$ $a_{61}$ $a_{62}$ $a_{70}$ $a_{71}$ $a_{72}$	$x_4$ $a_{43}$ $a_{44}$ $a_{45}$ $y_5$ $a_{53}$ $a_{54}$ $a_{55}$ $a_{63}$ $a_{64}$ $a_{65}$ $a_{73}$ $a_{74}$ $a_{75}$	$x_7$ $a_{46}$ $a_{47}$ $a_{48}$ $a_{56}$ $a_{57}$ $a_{58}$ $a_{66}$ $a_{67}$ $a_{68}$ $y_6$ $a_{76}$ $a_{77}$ $a_{78}$	$x_{10}$ $a_{49}$ $a_{50}$ $a_{51}$ $a_{52}$ $a_{59}$ $a_{60}$ $a_{61}$ $a_{62}$ $a_{69}$ $a_{70}$ $a_{71}$ $a_{72}$ $a_{79}$ $a_{80}$ $a_{81}$ $a_{82}$
$x_2$ $a_{80}$ $a_{81}$ $a_{82}$ $y_8$ $a_{90}$ $a_{91}$ $a_{92}$ $a_{10,0}$ $a_{10,1}$ $a_{10,2}$ $a_{11,0}$ $a_{11,1}$ $a_{11,2}$	$x_5$ $a_{83}$ $a_{84}$ $a_{85}$ $y_9$ $a_{93}$ $a_{94}$ $a_{95}$ $a_{10,3}$ $a_{10,4}$ $a_{10,5}$ $a_{11,3}$ $a_{11,4}$ $a_{11,5}$	$x_8$ $a_{86}$ $a_{87}$ $a_{88}$ $a_{96}$ $a_{97}$ $a_{98}$ $a_{10,6}$ $a_{10,7}$ $a_{10,8}$ $y_{10}$ $a_{11,6}$ $a_{11,7}$ $a_{11,8}$	$x_{11}$ $a_{89}$ $a_{90}$ $a_{91}$ $a_{92}$ $a_{99}$ $a_{100}$ $a_{101}$ $a_{102}$ $a_{10,9}$ $a_{10,10}$ $a_{10,11}$ $a_{10,12}$ $a_{11,9}$ $a_{11,10}$ $a_{11,11}$ $a_{11,12}$

# Algorithm

- Collecting  $x_j$  on each processor  $p_{ij}$  by an *allgather* inside the processor columns.
- Each processor  $p_{ij}$  then computes  $y_{ij} = A_{ij}x_j$ .
- Gathering together the pieces  $y_{ij}$  in each processor row to form  $y_i$ , distribute this over the processor row: combine to form a *reduce-scatter*.
- Setup for the next  $A$  or  $A^t$  product

# Analysis 1.

Step	Cost (lower bound)
Allgather $x_i$ 's within columns	$\lceil \log_2(r) \rceil \alpha + \frac{r-1}{p} n \beta \approx$
Perform local matrix-vector multiply	$\log_2(r) \alpha + \frac{n}{c} \beta$
Reduce-scatter $y_i$ 's within rows	$\approx 2 \frac{n^2}{p} \gamma$

# Reduce-scatter

	$t = 1$	$t = 2$	$t = 3$
$p_0$	$x_0^{(0)}, x_1^{(0)}, x_2^{(0)} \downarrow, x_3^{(0)} \downarrow$	$x_0^{(0:2:2)}, x_1^{(0:2:2)} \downarrow$	$x_0^{(0:3)}$
$p_1$	$x_0^{(1)}, x_1^{(1)}, x_2^{(1)} \downarrow, x_3^{(1)} \downarrow$	$x_0^{(1:3:2)} \uparrow, x_1^{(1:3:2)}$	$x_1^{(0:3)}$
$p_2$	$x_0^{(2)} \uparrow, x_1^{(2)} \uparrow, x_2^{(2)}, x_3^{(2)}$	$x_2^{(0:2:2)}, x_3^{(0:2:2)} \downarrow$	$x_2^{(0:3)}$
$p_3$	$x_0^{(3)} \uparrow, x_1^{(3)} \uparrow, x_2^{(3)}, x_3^{(3)}$	$x_0^{(1:3:2)} \uparrow, x_1^{(1:3:2)}$	$x_3^{(0:3)}$

Time:

$$\lceil \log_2 p \rceil \alpha + \frac{p-1}{p} n(\beta + \gamma).$$

Step	Cost (lower bound)
Allgather $x_i$ 's within columns	$\lceil \log_2(r) \rceil \alpha + \frac{r-1}{p} n \beta \approx \log_2(r) \alpha + \frac{n}{c} \beta$
Perform local matrix-vector multiply	$\approx 2 \frac{n^2}{p} \gamma$
Reduce-scatter $y_i$ 's within rows	$\lceil \log_2(c) \rceil \alpha + \frac{c-1}{p} n \beta + \frac{c-1}{p} n \gamma \approx \log_2(r) \alpha + \frac{n}{c} \beta + \frac{n}{c} \gamma$



# Efficiency

$$E_p^{\sqrt{p} \times \sqrt{p}}(n) = \frac{1}{1 + \frac{p \log_2(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{\sqrt{p}}{2n} \frac{(2\beta + \gamma)}{\gamma}}$$

Weak scaling:

for  $p \rightarrow \infty$  this is  $\approx 1/\log_2 P$ :

only slowly decreasing.